

AZ-400

Number: 000-000
Passing Score: 800
Time Limit: 120 min
File Version: 1.0

AZ-400

Microsoft Azure DevOps Solutions

TestGuide4U (ExamGuidesForIT)

Check Out Our Site at:

[www.e-junkie.com\TestGuide4u](http://www.e-junkie.com/TestGuide4u)

More Exams Can be Purchased through Credit Cards or Paypal Online Directly.
Download link will be sent to your email immediately after the purchase.

Exam A

QUESTION 1

You use Azure Artifacts to host NuGet packages that you create.

You need to make one of the packages available to anonymous users outside your organization. The solution must minimize the number of publication points.

What should you do?

- A. Change the feed URL of the package
- B. Create a new feed for the package
- C. Promote the package to a release view.
- D. Publish the package to a public NuGet repository.

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Azure Artifacts introduces the concept of multiple feeds that you can use to organize and control access to your packages.

Packages you host in Azure Artifacts are stored in a feed. Setting permissions on the feed allows you to share your packages with as many or as few people as your scenario requires.

Feeds have four levels of access: Owners, Contributors, Collaborators, and Readers.

References: <https://docs.microsoft.com/en-us/azure/devops/artifacts/feeds/feed-permissions?view=vsts&tabs=new-nav>

QUESTION 2

Your company plans to use an agile approach to software development.

You need to recommend an application to provide communication between members of the development team who work in locations around the world. The applications must meet the following requirements:

Provide the ability to isolate the members of different project teams into separate communication channels and to keep a history of the chats within those channels.

Be available on Windows 10, Mac OS, iOS, and Android operating systems.

Provide the ability to add external contractors and suppliers to projects.

Integrate directly with Azure DevOps.

What should you recommend?

- A. Microsoft Project
- B. Bamboo
- C. Microsoft Lync
- D. Microsoft Teams

Answer: D

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Within each team, users can create different channels to organize their communications by topic. Each channel can include a couple of users or scale to thousands of users.

Microsoft Teams works on Android, iOS, Mac and Windows systems and devices. It also works in Chrome, Firefox, Internet Explorer 11 and Microsoft Edge web browsers.

The guest-access feature in Microsoft Teams allows users to invite people outside their organizations to join internal channels for messaging, meetings and file sharing. This capability helps to facilitate business-to-business project management.

Teams integrates with Azure DevOps.

Note: Slack would also be a correct answer, but it is not an option here.

References:

<https://searchunifiedcommunications.techtarget.com/definition/Microsoft-Teams>

QUESTION 3

DRAG DROP

You need to recommend project metrics for dashboards in Azure DevOps.

Which chart widgets should you recommend for each metric? To answer, drag the appropriate chart widgets to the correct metrics. Each chart widget may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Chart Widgets	Answer Area
Burndown	The elapsed time from the creation of work items to their completion: <input type="text" value="0"/>
Cycle Time	The elapsed time to complete work items once they are active: <input type="text" value="0"/>
Lead Time	The remaining work: <input type="text" value="0"/>
Velocity	

A. Answer:

Chart Widgets	Answer Area
The elapsed time from the creation of work items to their completion:	Lead Time
The elapsed time to complete work items once they are active:	Cycle Time
Velocity	The remaining work:
	Burndown

Explanation

Explanation/Reference:

Explanation:

Box 1: Lead time

Lead time measures the total time elapsed from the creation of work items to their completion.

Box 2: Cycle time

Cycle time measures the time it takes for your team to complete work items once they begin actively working on them.

Box 3: Burndown

Burndown charts focus on remaining work within a specific time period.

Incorrect Answers:

Velocity provides a useful metric for these activities:

Support sprint planning

Forecast future sprints and the backlog items that can be completed

A guide for determining how well the team estimates and meets their planned commitments

References:

<https://docs.microsoft.com/en-us/azure/devops/report/dashboards/velocity-guidance?view=vsts>

<https://docs.microsoft.com/en-us/azure/devops/report/dashboards/cycle-time-and-lead-time?view=vsts>

<https://docs.microsoft.com/en-us/azure/devops/report/dashboards/configure-burndown-burnup-widgets?view=vsts>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box 1: Lead time

Lead time measures the total time elapsed from the creation of work items to their completion.

Box 2: Cycle time

Cycle time measures the time it takes for your team to complete work items once they begin actively working on them.

Box 3: Burndown

Burndown charts focus on remaining work within a specific time period.

Incorrect Answers:

Velocity provides a useful metric for these activities:

Support sprint planning

Forecast future sprints and the backlog items that can be completed

A guide for determining how well the team estimates and meets their planned commitments

References:

<https://docs.microsoft.com/en-us/azure/devops/report/dashboards/velocity-guidance?view=vsts>

<https://docs.microsoft.com/en-us/azure/devops/report/dashboards/cycle-time-and-lead-time?view=vsts>

<https://docs.microsoft.com/en-us/azure/devops/report/dashboards/configure-burndown-burnup-widgets?view=vsts>

QUESTION 4

You manage build pipelines and deployment pipelines by using Azure DevOps.

Your company has a team of 500 developers. New members are added continually to the team.

You need to automate the management of users and licenses whenever possible.

Which task must you perform manually?

- A. modifying group memberships
- B. adding users
- C. assigning entitlements
- D. procuring licenses

Answer: D

Explanation/Reference:

Explanation

Explanation/Reference:

Incorrect Answers:

A: You can seamlessly replace existing solutions with group-based licensing to more easily manage licenses in Azure DevOps. You can use Group rules.

C: Member Entitlement Management APIs allow managing Entitlements that include -
License

Extensions

Project/Team memberships

References:

<https://docs.microsoft.com/en-us/azure/devops/organizations/accounts/migrate-to-group-based-resource-management?view=vsts&tabs=new-nav>

<https://docs.microsoft.com/en-us/rest/api/azure/devops/memberentitlementmanagement/?view=azure-devops-rest-5.0>

QUESTION 5

DRAG DROP

You need to increase the security of your team's development process.

Which type of security tool should you recommend for each stage of the development process? To answer, drag the appropriate security tools to the correct stages. Each security tool may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Security Tools	Answer Area
Penetration testing	Pull request: <input type="text"/>
Static code analysis	Continuous integration: <input type="text"/>
Threat modeling	Continuous delivery: <input type="text"/>

A. Answer:

Security Tools	Answer Area
Pull request:	Threat modeling
Continuous integration:	Static code analysis
Continuous delivery:	Penetration testing

Explanation

Explanation/Reference:

Explanation:

Box 1: Threat modeling

Threat modeling's motto should be, "The earlier the better, but not too late and never ignore."

Box 2: Static code analysis

Validation in the CI/CD begins before the developer commits his or her code. Static code analysis tools in the IDE provide the first line of defense to help ensure that security vulnerabilities are not introduced into the CI/CD process.

Box 3: Penetration testing

Once your code quality is verified, and the application is deployed to a lower environment like development or QA, the process should verify that there are not any security vulnerabilities in the running application. This can be accomplished by executing automated penetration test against the running application to scan it for vulnerabilities.

References: <https://docs.microsoft.com/en-us/azure/devops/articles/security-validation-cicd-pipeline?view=vsts>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box 1: Threat modeling

Threat modeling's motto should be, "The earlier the better, but not too late and never ignore."

Box 2: Static code analysis

Validation in the CI/CD begins before the developer commits his or her code. Static code analysis tools in the IDE provide the first line of defense to help ensure that security vulnerabilities are not introduced into the CI/CD process.

Box 3: Penetration testing

Once your code quality is verified, and the application is deployed to a lower environment like development or QA, the process should verify that there are not any security vulnerabilities in the running application. This can be accomplished by executing automated penetration test against the running application to scan it for vulnerabilities.

References: <https://docs.microsoft.com/en-us/azure/devops/articles/security-validation-cicd-pipeline?view=vsts>

QUESTION 6

HOTSPOT

Your company uses Team Foundation Server 2013 (TFS 2013).

You plan to migrate to Azure DevOps.

You need to recommend a migration strategy that meets the following requirements:

Preserves the dates of Team Foundation Version Control changesets

Preserves the changes dates of work items revisions

Minimizes migration effort

Migrates all TFS artifacts

What should you recommend? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area	
On the TFS server:	<ul style="list-style-type: none"><input type="checkbox"/> Install the TFS Java SDK.<input type="checkbox"/> Upgrade TFS to the most recent RTW release.<input type="checkbox"/> Upgrade to the most recent version of PowerShell Core.
To perform the migration:	<ul style="list-style-type: none"><input type="checkbox"/> Copy the assets manually.<input type="checkbox"/> Use public API-based tools.<input type="checkbox"/> Use the TFS Database Import Service.<input type="checkbox"/> Use the TFS Integration Platform.

A. Answer:

Answer Area

On the TFS server:

- Install the TFS Java SDK.
- Upgrade TFS to the most recent RTW release.
- Upgrade to the most recent version of PowerShell Core.

To perform the migration:

- Copy the assets manually.
- Use public API-based tools.
- Use the TFS Database Import Service.
- Use the TFS Integration Platform.

Explanation

Explanation/Reference:

Explanation:

Box 1: Upgrade TFS to the most recent RTM release.

One of the major prerequisites for migrating your Team Foundation Server database is to get your database schema version as close as possible to what is currently deployed in Azure Devops Services.

Box 2: Use the TFS Database Import Service

In Phase 3 of your migration project, you will work on upgrading your Team Foundation Server to one of the supported versions for the Database Import Service in Azure Devops Services.

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box 1: Upgrade TFS to the most recent RTM release.

One of the major prerequisites for migrating your Team Foundation Server database is to get your database schema version as close as possible to what is currently deployed in Azure Devops Services.

Box 2: Use the TFS Database Import Service

In Phase 3 of your migration project, you will work on upgrading your Team Foundation Server to one of the supported versions for the Database Import Service in Azure Devops Services.

QUESTION 7

You are developing a multi-tier application. The application will use Azure App Service web apps as the front end and an Azure SQL database as the back end. The application will use Azure functions to write some data to Azure Storage.

You need to send the Azure DevOps team an email message when the front end fails to return a status code of 200.

Which feature should you use?

- A. Service Map in Azure Log Analytics
- B. Availability tests in Azure Application Insights

- C. Profiler in Azure Application Insights
- D. Application Map in Azure Application Insights

Answer: D

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Application Map helps you spot performance bottlenecks or failure hotspots across all components of your distributed application. Each node on the map represents an application component or its dependencies; and has health KPI and alerts status.

Incorrect Answers:

A: Service Map automatically discovers application components on Windows and Linux systems and maps the communication between services. You can use it to view your servers as you think of them--interconnected systems that deliver critical services.

Service Map shows connections between servers, processes, and ports across any TCP-connected architecture with no configuration required, other than installation of an agent.

References: <https://docs.microsoft.com/en-us/azure/azure-monitor/app/app-map>

QUESTION 8

During a code review, you discover many quality issues. Many modules contain unused variables and empty catch blocks.

You need to recommend a solution to improve the quality of the code.

What should you recommend?

- A. In a Grunt build task, select Enabled from Control Options.
- B. In a Maven build task, select Run PMD.
- C. In a Xcode build task, select Use xcpretty from Advanced.
- D. In a Gradle build task, select Run Checkstyle.

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

PMD is a source code analyzer. It finds common programming flaws like unused variables, empty catch blocks, unnecessary object creation, and so forth.

There is an Apache Maven PMD Plugin which allows you to automatically run the PMD code analysis tool on your project's source code and generate a site report with its results.

Incorrect Answers:

C: xcpretty is a fast and flexible formatter for xcodebuild.

References: <https://pmd.github.io/>

QUESTION 9

Your company has an on-premises Bitbucket Server that is used for Git-based source control. The server is protected by a firewall that blocks inbound Internet traffic. You plan to use Azure DevOps to manage the build and release processes.

Which two components are required to integrate Azure DevOps and Bitbucket? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. a deployment group
- B. a Microsoft-hosted agent
- C. service hooks
- D. a self-hosted agent
- E. an External Git service connection

Answer: DE

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

E: GitLab CI/CD can be used with GitHub or any other Git server such as BitBucket. Instead of moving your entire project to GitLab, you can connect your external repository to get the benefits of GitLab CI/CD.

Note: When a pipeline uses a remote, 3rd-party repository host such as Bitbucket Cloud, the repository is configured with webhooks that notify Azure Pipelines Server or TFS when code has changed and a build should be triggered. Since on-premises

installations are normally protected behind a firewall, 3rd-party webhooks are unable to reach the on-premises server. As a workaround, you can use the External Git repository type which uses polling instead of webhooks to trigger a build when code has changed.

References: <https://docs.microsoft.com/en-us/azure/devops/pipelines/repos/pipeline-options-for-git>

QUESTION 10

Your company plans to use an agile approach to software development.

You need to recommend an application to provide communication between members of the development team who work in locations around the world. The applications must meet the following requirements:

Provide the ability to isolate the members of different project teams into separate communication channels and to keep a history of the chats within those channels.

Be available on Windows 10, Mac OS, iOS, and Android operating systems.

Provide the ability to add external contractors and suppliers to projects.

Integrate directly with Azure DevOps.

What should you recommend?

- A. Microsoft Project
- B. Bamboo
- C. Octopus
- D. Slack

Answer: D

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Slack is a popular team collaboration service that helps teams be more productive by keeping all communications in one place and easily searchable from virtually anywhere. All your messages, your files, and everything from Twitter, Dropbox, Google Docs,

Azure DevOps, and more all together. Slack also has fully native apps for iOS and Android to give you the full functionality of Slack wherever you go.

Integrated with Azure DevOps

This integration keeps your team informed of activity happening in its Azure DevOps projects. With this integration, code check-ins, pull requests, work item updates, and build events show up directly in your team's Slack channel.

Note: Microsoft Teams would also be a correct answer, but it is not an option here.

References:

<https://marketplace.visualstudio.com/items?itemName=ms-vsts.vss-services-slack>

QUESTION 11

DRAG DROP

You are planning projects for three customers. Each customer's preferred process for work items is shown in the following table.

Customer name	Preferred process
Litware, Inc.	Track product backlog items (PBIs) and bugs on the Kanban board. Break the PBIs down into tasks on the task board.
Contoso, Ltd.	Track user stories and bugs on the Kanban board. Track the bugs and tasks on the task board.
A. Datum Corporation	Track requirements, change requests, risks, and reviews.

The customers all plan to use Azure DevOps for work item management.

Which work item process should you use for each customer? To answer, drag the appropriate work item processes to the correct customers. Each work item process may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Processes	Answer Area
Agile	Litware
CMMI	Contoso:
Scrum	A. Datum:
XP	

A. Answer:

Processes	Answer Area
	Litware
	Scrum
	Contoso:
	Agile
	A. Datum:
	CMMI
XP	

Explanation

Explanation/Reference:

Explanation:

Box 1: Scrum

Choose Scrum when your team practices Scrum. This process works great if you want to track product backlog items (PBIs) and bugs on the Kanban board, or break PBIs and bugs down into tasks on the taskboard.

Box 2: Agile

Choose Agile when your team uses Agile planning methods, including Scrum, and tracks development and test activities separately. This process works great if you want to track user stories and (optionally) bugs on the Kanban board, or track bugs and tasks on the taskboard.

Box 3: CMMI

Choose CMMI when your team follows more formal project methods that require a framework for process improvement and an auditable record of decisions. With this process, you can track requirements, change requests, risks, and reviews.

Incorrect Answers:

XP:

The work tracking objects contained within the default DevOps processes and DevOps process templates are Basic, Agile, CMMI, and Scrum

XP (Extreme Programming) and DevOps are different things. They don't contradict with each other, they can be used together, but they have different base concepts inside them.

References:

<https://docs.microsoft.com/en-us/azure/devops/boards/work-items/guidance/choose-process?view=azure-devops>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box 1: Scrum

Choose Scrum when your team practices Scrum. This process works great if you want to track product backlog items (PBIs) and bugs on the Kanban board, or break PBIs and bugs down into tasks on the taskboard.

Box 2: Agile

Choose Agile when your team uses Agile planning methods, including Scrum, and tracks development and test activities separately. This process works great if you want to track user stories and (optionally) bugs on the Kanban board, or track bugs and tasks on the taskboard.

Box 3: CMMI

Choose CMMI when your team follows more formal project methods that require a framework for process improvement and an auditable record of decisions. With this process, you can track requirements, change requests, risks, and reviews.

Incorrect Answers:

XP:

The work tracking objects contained within the default DevOps processes and DevOps process templates are Basic, Agile, CMMI, and Scrum

XP (Extreme Programming) and DevOps are different things. They don't contradict with each other, they can be used together, but they have different base concepts inside them.

References:

<https://docs.microsoft.com/en-us/azure/devops/boards/work-items/guidance/choose-process?view=azure-devops>

QUESTION 12

Your development team is building a new web solution by using the Microsoft Visual Studio integrated development environment (IDE).

You need to make a custom package available to all the developers. The package must be managed centrally, and the latest version must be available for consumption in Visual Studio automatically.

Which three actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Publish the package to a feed.
- B. Create a new feed in Azure Artifacts.
- C. Upload a package to a Git repository.
- D. Add the package URL to the Environment settings in Visual Studio.
- E. Add the package URL to the NuGet Package Manager settings in Visual Studio.
- F. Create a Git repository in Azure Repos.

Answer: ABE

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

B: By using your custom NuGet package feed within your Azure DevOps (previously VSTS) instance, you'll be able to distribute your packages within your organization with ease.

Start by creating a new feed.

A: We can publish, pack and push the built project to our NuGet feed.

E: Consume your private NuGet Feed

Go back to the Packages area in Azure DevOps, select your feed and hit "Connect to feed". You'll see some instructions for your feed, but it's fairly simple to set up.

Just copy your package source URL, go to Visual Studio, open the NuGet Package Manager, go to its settings and add a new source. Choose a fancy name, insert the source URL. Done.

Search for your package in the NuGet Package Manager and it should appear there, ready for installation. Make sure to select the appropriate feed (or just all feeds) from the top right select box.

References:

<https://medium.com/medialesson/get-started-with-private-nuget-feeds-in-azure-devops-8c7b5f022a68>

QUESTION 13

You have a GitHub repository.

You create a new repository in Azure DevOps.

You need to recommend a procedure to clone the repository from GitHub to Azure DevOps.

What should you recommend?

- A. Create a pull request.
- B. Create a webhook.
- C. Create a service connection for GitHub.
- D. From Import a Git repository, click Import.
- E. Create a personal access token in Azure DevOps.

Answer: D

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

You can import an existing Git repo from GitHub, Bitbucket, GitLab, or other location into a new or empty existing repo in your project in Azure DevOps.

Import into a new repo

1. Select Repos, Files.
2. From the repo drop-down, select Import repository.
3. If the source repo is publicly available, just enter the clone URL of the source repository and a name for your new Git repository.

References:

<https://docs.microsoft.com/en-us/azure/devops/repos/git/import-git-repository?view=azure-devops>

QUESTION 14

Your company is concerned that when developers introduce open source libraries, it creates licensing compliance issues.

You need to add an automated process to the build pipeline to detect when common open source libraries are added to the code base.

What should you use?

- A. SourceGear Vault
- B. Jenkins
- C. Microsoft Visual SourceSafe
- D. WhiteSource Bolt

Answer: D

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

WhiteSource provides WhiteSource Bolt, a lightweight open source security and management solution developed specifically for integration with Azure DevOps and Azure DevOps Server.

Note: WhiteSource is the leader in continuous open source software security and compliance management. WhiteSource integrates into your build process, irrespective of your programming languages, build tools, or development environments. It works automatically, continuously, and silently in the background, checking the security, licensing, and quality of your open source components against WhiteSource constantly-updated definitive database of open source repositories.

Reference:

<https://www.azuredevopslabs.com/labs/vstsextend/whitesource/>

QUESTION 15

You plan to use a NuGet package in a project in Azure DevOps. The NuGet package is in a feed that requires authentication.

You need to ensure that the project can restore the NuGet package automatically.

What should the project use to automate the authentication?

- A. an Azure Automation account
- B. an Azure Artifacts Credential Provider
- C. an Azure Active Directory (Azure AD) account that has multi-factor authentication (MFA) enabled
- D. an Azure Active Directory (Azure AD) service principal

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

The Azure Artifacts Credential Provider automates the acquisition of credentials needed to restore NuGet packages as part of your .NET development workflow. It integrates with MSBuild, dotnet, and NuGet(.exe) and works on Windows, Mac, and Linux. Any

time you want to use packages from an Azure Artifacts feed, the Credential Provider will automatically acquire and securely store a token on behalf of the NuGet client you're using.

Reference:

<https://github.com/Microsoft/artifacts-credprovider>

QUESTION 16

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

The lead developer at your company reports that adding new application features takes longer than expected due to a large accumulated technical debt.

You need to recommend changes to reduce the accumulated technical debt.

Solution: You recommend increasing the code duplication.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Instead reduce the code complexity.

Reference:

<https://dzone.com/articles/fight-through-the-pain-how-to-deal-with-technical>

QUESTION 17

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

The lead developer at your company reports that adding new application features takes longer than expected due to a large accumulated technical debt.

You need to recommend changes to reduce the accumulated technical debt.

Solution: You recommend increasing the test coverage.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Instead reduce the code complexity.

Reference:

<https://dzone.com/articles/fight-through-the-pain-how-to-deal-with-technical>

QUESTION 18

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

The lead developer at your company reports that adding new application features takes longer than expected due to a large accumulated technical debt.

You need to recommend changes to reduce the accumulated technical debt.

Solution: You recommend reducing the code complexity.

Does this meet the goal?

- A. Yes
- B. No

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Reference:
<https://dzone.com/articles/fight-through-the-pain-how-to-deal-with-technical>

QUESTION 19

Your company has 60 developers who are assigned to four teams. Each team has 15 members.

The company uses an agile development methodology.

You need to structure the work of the development teams so that each team owns their respective work while working together to reach a common goal.

Which parts of the taxonomy should you enable the team to perform autonomously?

- A. Features and Tasks
- B. Initiatives and Epics
- C. Epics and Features
- D. Stories and Tasks

Answer: A

Explanation/Reference:
Explanation

Explanation/Reference:

Explanation:

A feature typically represents a shippable component of software.

Features, examples:

Add view options to the new work hub

Add mobile shopping cart

Support text alerts

Refresh the web portal with new look and feel

User Stories and Tasks are used to track work. Teams can choose how they track bugs, either as requirements or as tasks

Incorrect Answers:

B, C: An epic represents a business initiative to be accomplished.

Epics, examples:

Increase customer engagement

Improve and simplify the user experience

Implement new architecture to improve performance

Engineer the application to support future growth

Support integration with external services

Support mobile apps

Reference:

<https://docs.microsoft.com/en-us/azure/devops/boards/backlogs/define-features-epics>

<https://docs.microsoft.com/en-us/azure/devops/boards/work-items/about-work-items>

QUESTION 20

HOTSPOT

Your company uses Git as a source code control system for a complex app named App1.

You plan to add a new functionality to App1.

You need to design a branching model for the new functionality.

Which branch lifetime and branch type should you use in the branching model? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Branch lifetime:

- Long-lived
- Short-lived

Branch type:

- Master
- Feature
- Integration

A. Answer:

Answer Area

Branch lifetime:	<input type="button" value="▼"/>
Long-lived	
Short-lived	
Branch type:	<input type="button" value="▼"/>
Master	
Feature	
Integration	

Explanation

Explanation/Reference:

Explanation:

Branch lifetime: Short-lived

Branch type: Feature

Feature branches are used when developing a new feature or enhancement which has the potential of a development lifespan longer than a single deployment. When starting development, the deployment in which this feature will be released may not be known. No matter when the feature branch will be finished, it will always be merged back into the master branch.

References:

<https://gist.github.com/digitaljhelms/4287848>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Branch lifetime: Short-lived

Branch type: Feature

Feature branches are used when developing a new feature or enhancement which has the potential of a development lifespan longer than a single deployment. When starting development, the deployment in which this feature will be released may not be known. No matter when the feature branch will be finished, it will always be merged back into the master branch.

References:

<https://gist.github.com/digitaljhelms/4287848>

QUESTION 21

You store source code in a Git repository in Azure repos. You use a third-party continuous integration (CI) tool to control builds.

What will Azure DevOps use to authenticate with the tool?

- A. certificate authentication
- B. a personal access token (PAT)
- C. a Shared Access Signature (SAS) token
- D. NTLM authentication

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Personal access tokens (PATs) give you access to Azure DevOps and Team Foundation Server (TFS), without using your username and password directly.

Reference:

<https://docs.microsoft.com/en-us/azure/devops/repos/git/auth-overview>

QUESTION 22

During a code review, you discover quality issues in a Java application.

You need to recommend a solution to detect quality issues including unused variables and empty catch blocks.

What should you recommend?

- A. In a Maven build task, select Run PMD.
- B. In an Xcode build task, select Use xcpretty from Advanced.
- C. In a Gulp build task, specify a custom condition expression.
- D. In a Grunt build task, select Enabled from Control Options.

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

PMD is a source code analyzer. It finds common programming flaws like unused variables, empty catch blocks, unnecessary object creation, and so forth.

There is an Apache Maven PMD Plugin which allows you to automatically run the PMD code analysis tool on your project's source code and generate a site report with its results.

Incorrect Answers:

B: xcpretty is a fast and flexible formatter for xcodebuild.

Reference:

<https://pmd.github.io/>

QUESTION 23

Your company creates a new Azure DevOps team.

You plan to use Azure DevOps for sprint planning.

You need to visualize the flow of your work by using an agile methodology.

Which Azure DevOps component should you use?

- A. Kanban boards
- B. sprint planning
- C. delivery plans
- D. portfolio backlogs

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Customizing Kanban boards

To maximize a team's ability to consistently deliver high quality software, Kanban emphasize two main practices. The first, visualize the flow of work, requires you to map your team's workflow stages and configure your Kanban board to match. Your Kanban

board turns your backlog into an interactive signboard, providing a visual flow of work.

Reference:

<https://azuredovplabs.com/labs/azuredevops/agile/>

QUESTION 24

Testlet 1

Case Study

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other question in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview

Litware, Inc. is an independent software vendor (ISV). Litware has a main office and five branch offices.

Existing Environment

Application Architecture

The company's primary application is a single monolithic retirement fund management system based on ASP.NET web forms that use logic written in VB.NET. Some new sections of the application are written in C#.

Variations of the application are created for individual customers. Currently, there are more than 80 live code branches in the application's code base.

The application was developed by using Microsoft Visual Studio. Source code is stored in Team Foundation Server (TFS) in the main office. The branch offices access the source code by using TFS proxy servers.

Architectural Issues

Litware focuses on writing new code for customers. No resources are provided to refactor or remove existing code. Changes to the code base take a long time, as dependencies are not obvious to individual developers.

Merge operations of the code often take months and involve many developers. Code merging frequently introduces bugs that are difficult to locate and resolve.

Customers report that ownership costs of the retirement fund management system increase continually. The need to merge unrelated code makes even minor code changes expensive.

Customers report that bug reporting is overly complex.

Requirements

Planned changes

Litware plans to develop a new suite of applications for investment planning. The investment planning applications will require only minor integration with the existing retirement fund management system.

The investment planning applications suite will include one multi-tier web application and two iOS mobile applications. One mobile application will be used by employees; the other will be used by customers.

Litware plans to move to a more agile development methodology. Shared code will be extracted into a series of packages.

Litware has started an internal cloud transformation process and plans to use cloud-based services whenever suitable.

Litware wants to become proactive in detecting failures, rather than always waiting for customer bug reports.

Technical requirements

The company's investment planning applications suite must meet the following requirements:

New incoming connections through the firewall must be minimized.

Members of a group named Developers must be able to install packages.

The principle of least privilege must be used for all permission assignments.

A branching strategy that supports developing new functionality in isolation must be used.

Members of a group named Team Leaders must be able to create new packages and edit the permissions of package feeds.

Visual Studio App Center must be used to centralize the reporting of mobile application crashes and device types in use.

By default, all releases must remain available for 30 days, except for production releases, which must be kept for 60 days.

Code quality and release quality are critical. During release, deployments must not proceed between stages if any active bugs are logged against the release.

The mobile applications must be able to call the share pricing service of the existing retirement fund management system. Until the system is upgraded, the service will only support basic authentication over HTTPS.

The required operating system configuration for the test servers changes weekly. Azure Automation State Configuration must be used to ensure that the operating system on each test server is configured the same way when the servers are created and

checked periodically.

Current Technical Issue

The test servers are configured correctly when first deployed, but they experience configuration drift over time. Azure Automation State Configuration fails to correct the configurations.

Azure Automation State Configuration nodes are registered by using the following command.

```
Register-AzureRmAutomationDscNode
-ResourceGroupName 'TestResourceGroup'
-AutomationAccountName 'LitwareAutomationAccount'
-AzureVMName $vname
-ConfigurationMode 'ApplyOnly'
```

To resolve the current technical issue, what should you do to the Register-AzureRmAutomationDscNode command?

- A. Change the value of the ConfigurationMode parameter.
- B. Replace the Register-AzureRmAutomationDscNode cmdlet with Register-AzureRmAutomationScheduledRunbook
- C. Add the AllowModuleOverwrite parameter.
- D. Add the DefaultProfile parameter.

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Change the ConfigurationMode parameter from ApplyOnly to ApplyAndAutocorrect.

The Register-AzureRmAutomationDscNode cmdlet registers an Azure virtual machine as an APS Desired State Configuration (DSC) node in an Azure Automation account.

Scenario: Current Technical Issue

The test servers are configured correctly when first deployed, but they experience configuration drift over time. Azure Automation State Configuration fails to correct the configurations.

Azure Automation State Configuration nodes are registered by using the following command.

```
Register-AzureRmAutomationDscNode  
-ResourceGroupName 'TestResourceGroup'  
-AutomationAccountName 'LitwareAutomationAccount'  
-AzureVMName $vmaname  
-ConfigurationMode 'ApplyOnly'
```

References:

<https://docs.microsoft.com/en-us/powershell/module/azurerm.automation/register-azurermautomationdscnode?view=azurermps-6.13.0>

QUESTION 25

HOTSPOT

You need to configure a cloud service to store the secrets required by the mobile applications to call the share pricing service. What should you include in the solution? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Required secrets:

- Certificate
- Personal access token
- Shared Access Authorization token
- Username and password

Storage location:

- Azure Data Lake
- Azure Key Vault
- Azure Storage with HTTP access
- Azure Storage with HTTPS access

A. Answer:

Answer Area

Required secrets: Certificate
 Personal access token
 Shared Access Authorization token
 Username and password

Storage location: Azure Data Lake
 Azure Key Vault
 Azure Storage with HTTP access
 Azure Storage with HTTPS access

Explanation

Explanation/Reference:

Explanation:

Every request made against a storage service must be authorized, unless the request is for a blob or container resource that has been made available for public or signed access. One option for authorizing a request is by using Shared Key.

Scenario: The mobile applications must be able to call the share pricing service of the existing retirement fund management system. Until the system is upgraded, the service will only support basic authentication over HTTPS.

The investment planning applications suite will include one multi-tier web application and two iOS mobile application. One mobile application will be used by employees; the other will be used by customers.

References: <https://docs.microsoft.com/en-us/rest/api/storageservices/authorize-with-shared-key>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Every request made against a storage service must be authorized, unless the request is for a blob or container resource that has been made available for public or signed access. One option for authorizing a request is by using Shared Key.

Scenario: The mobile applications must be able to call the share pricing service of the existing retirement fund management system. Until the system is upgraded, the service will only support basic authentication over HTTPS.

The investment planning applications suite will include one multi-tier web application and two iOS mobile application. One mobile application will be used by employees; the other will be used by customers.

References: <https://docs.microsoft.com/en-us/rest/api/storageservices/authorize-with-shared-key>

QUESTION 26

HOTSPOT

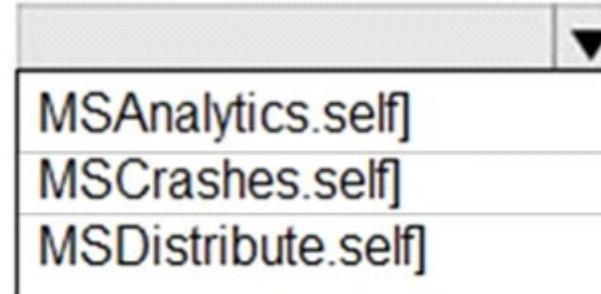
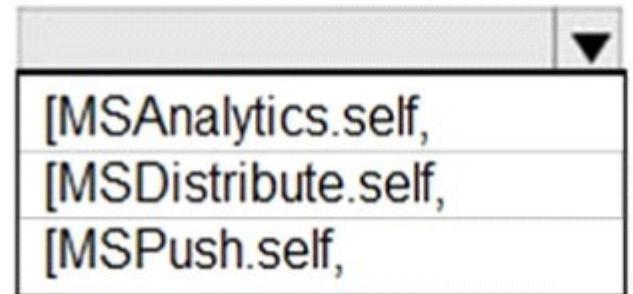
How should you complete the code to initialize App Center in the mobile application? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

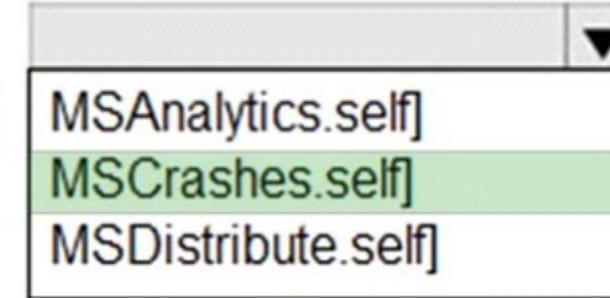
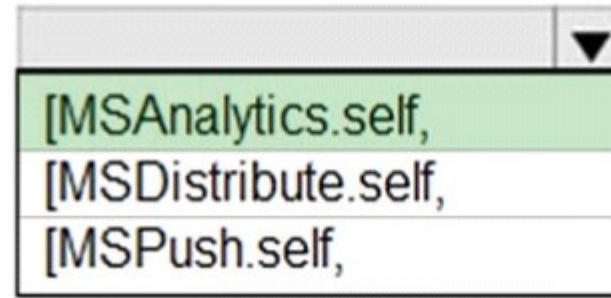
```
MSAppCenter.start  
( "{Your App Secret}",  
  withServices:  
)
```



A. Answer:

Answer Area

```
MSAppCenter.start  
( "{Your App Secret}",  
  withServices:  
)
```



Explanation

Explanation/Reference:

Scenario: Visual Studio App Center must be used to centralize the reporting of mobile application crashes and device types in use.

In order to use App Center, you need to opt in to the service(s) that you want to use, meaning by default no services are started and you will have to explicitly call each of them when starting the SDK.

Insert the following line to start the SDK in your app's AppDelegate class in the didFinishLaunchingWithOptions method.

```
MSAppCenter.start("{Your App Secret}", withServices: [MSAnalytics.self, MSCrashes.self])
```

References:

<https://docs.microsoft.com/en-us/appcenter/sdk/getting-started/ios>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Scenario: Visual Studio App Center must be used to centralize the reporting of mobile application crashes and device types in use.

In order to use App Center, you need to opt in to the service(s) that you want to use, meaning by default no services are started and you will have to explicitly call each of them when starting the SDK.

Insert the following line to start the SDK in your app's AppDelegate class in the didFinishLaunchingWithOptions method.

```
MSAppCenter.start("{Your App Secret}", withServices: [MSAnalytics.self, MSCrashes.self])
```

References:

<https://docs.microsoft.com/en-us/appcenter/sdk/getting-started/ios>

QUESTION 27

Which branching strategy should you recommend for the investment planning applications suite?

- A. release isolation
- B. main only
- C. development isolation
- D. feature isolation

Answer: D

Explanation/Reference:

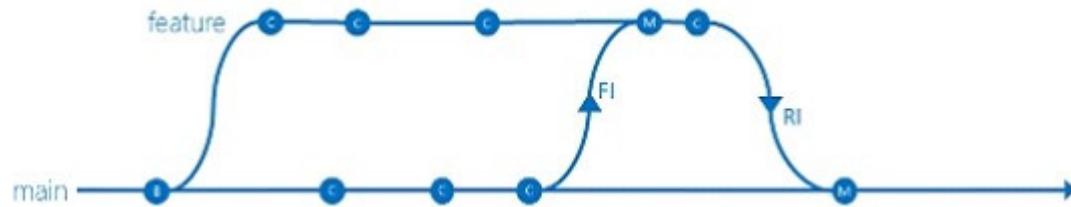
Explanation

Explanation/Reference:

Explanation:

Scenario: A branching strategy that supports developing new functionality in isolation must be used.

Feature isolation is a special derivation of the development isolation, allowing you to branch one or more feature branches from main, as shown, or from your dev branches.



When you need to work on a particular feature, it might be a good idea to create a feature branch.

Incorrect Answers:

A: Release isolation introduces one or more release branches from main. The strategy allows concurrent release management, multiple and parallel releases, and codebase snapshots at release time.

B: The Main Only strategy can be folder-based or with the main folder converted to a Branch, to enable additional visibility features. You commit your changes to the main branch and optionally indicate development and release milestones with labels.

C: Development isolation: When you need to maintain and protect a stable main branch, you can branch one or more dev branches from main. It enables isolation and concurrent development. Work can be isolated in development branches by feature, organization, or temporary collaboration.

References:

<https://docs.microsoft.com/en-us/azure/devops/repos/tfvc/branching-strategies-with-tfvc?view=azure-devops>

QUESTION 28

Testlet 2

Case Study

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case

study. Each question is independent of the other question in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview

Contoso, Ltd. is a manufacturing company that has a main office in Chicago.

Existing Environment

Contoso plans to improve its IT development and operations processes by implementing Azure DevOps principles. Contoso has an Azure subscription and creates an Azure DevOps organization.

The Azure DevOps organization includes:

The Docker extension

A deployment pool named Pool7 that contains 10 Azure virtual machines that run Windows Server 2016

The Azure subscription contains an Azure Automation account.

Requirements

Planned changes

Contoso plans to create projects in Azure DevOps as shown in the following table.

Project name	Project details
Project 1	Project1 will provide support for incremental builds and third-party SDK components
Project 2	Project2 will use an automatic build policy. A small team of developers named Team2 will work independently on changes to the project. The Team2 members will not have permissions to Project2.
Project 3	Project3 will be integrated with SonarQube
Project 4	Project4 will provide support for a build pipeline that creates a Docker image and pushes the image to the Azure Container Registry. Project4 will use an existing Dockerfile.
Project 5	Project5 will contain a Git repository in Azure Repos and a continuous integration trigger that will initiate a build in response to any change except for changes within /folder1 of the repository.
Project 6	Project6 will provide support for build and deployment pipelines. Deployment will be allowed only if the number of current work items representing active software bugs is 0.
Project 7	Project7 will contain a target deployment group named Group7 that maps to Pool7. Project7 will use Azure Automation State Configuration to maintain the desired state of the computers in Group7.

Technical requirements

Contoso identifies the following technical requirements:

Implement build agents for Project1.

Whenever possible, use Azure resources.

Avoid using deprecated technologies.

Implement a code flow strategy for Project2 that will:

- Enable Team2 to submit pull requests for Project2.
 - Enable Team2 to work independently on changes to a copy of Project2.
 - Ensure that any intermediary changes performed by Team2 on a copy of Project2 will be subject to the same restrictions as the ones defined in the build policy of Project2.
- Whenever possible, implement automation and minimize administrative effort.

Implement Project3, Project5, Project6, and Project7 based on the planned changes.

Implement Project4 and configure the project to push Docker images to Azure Container Registry.

You add the virtual machines as managed nodes in Azure Automation State Configuration.

You need to configure the managed computers in Pool7.

What should you do next?

- A. Modify the RefreshMode property of the Local Configuration Manager (LCM).
- B. Run the Register-AzureRmAutomationDscNode Azure Powershell cmdlet.
- C. Modify the ConfigurationMode property of the Local Configuration Manager (LCM).
- D. Install PowerShell Core.

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

The Register-AzureRmAutomationDscNode cmdlet registers an Azure virtual machine as an APS Desired State Configuration (DSC) node in an Azure Automation account.

Scenario: The Azure DevOps organization includes:

The Docker extension

A deployment pool named Pool7 that contains 10 Azure virtual machines that run Windows Server 2016

Project 7	Project7 will contain a target deployment group named Group7 that maps to Pool7. Project7 will use Azure Automation State Configuration to maintain the desired state of the computers in Group7.
-----------	---

References: <https://docs.microsoft.com/en-us/powershell/module/azurerm.automation/register-azurermautomationdscnode>

QUESTION 29

DRAG DROP

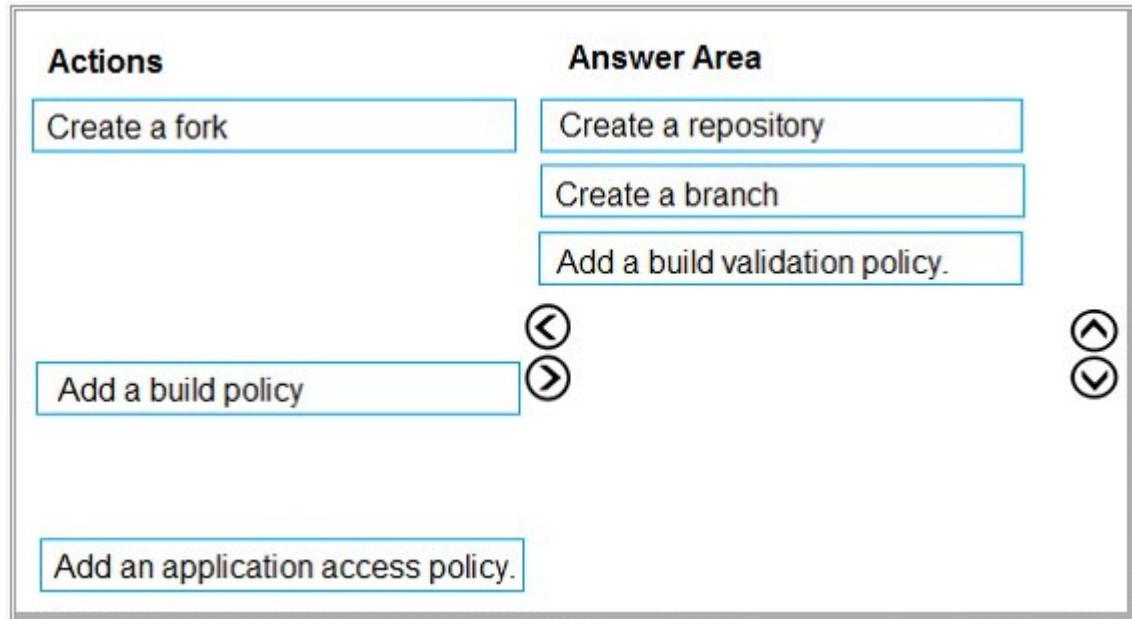
You need to implement the code flow strategy for Project2 in Azure DevOps.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions	Answer Area
Create a fork	0
Create a branch	0
Add a build validation policy.	0
Add a build policy	< >
Create a repository	
Add an application access policy.	

A. Answer:



Explanation

Explanation/Reference:

Explanation:

Step 1: Create a repository

A Git repository, or repo, is a folder that you've told Git to help you track file changes in. You can have any number of repos on your computer, each stored in their own folder.

Step 2: Create a branch

Branch policies help teams protect their important branches of development. Policies enforce your team's code quality and change management standards.

Step 3: Add a build validation policy

When a build validation policy is enabled, a new build is queued when a new pull request is created or when changes are pushed to an existing pull request targeting this branch. The build policy then evaluates the results of the build to determine whether the

pull request can be completed.

Note: Each branch must have a defined policy about how to integrate code into this branch.

Scenario:

Implement a code flow strategy for Project2 that will:

Enable Team2 to submit pull requests for Project2.

Enable Team2 to work independently on changes to a copy of Project2.

Ensure that any intermediary changes performed by Team2 on a copy of Project2 will be subject to the same restrictions as the ones defined in the build policy of Project2.

Project2 will use an automatic build policy. A small team of developers named Team2 will work independently on changes to the project. The Team2 members will not have permissions to Project2.

References:

<https://docs.microsoft.com/en-us/azure/devops/repos/git/manage-your-branches>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation

Step 1: Create a repository

A Git repository, or repo, is a folder that you've told Git to help you track file changes in. You can have any number of repos on your computer, each stored in their own folder.

Step 2: Create a branch

Branch policies help teams protect their important branches of development. Policies enforce your team's code quality and change management standards.

Step 3: Add a build validation policy

When a build validation policy is enabled, a new build is queued when a new pull request is created or when changes are pushed to an existing pull request targeting this branch. The build policy then evaluates the results of the build to determine whether the pull request can be completed.

Note: Each branch must have a defined policy about how to integrate code into this branch.

Scenario:

Implement a code flow strategy for Project2 that will:

Enable Team2 to submit pull requests for Project2.

Enable Team2 to work independently on changes to a copy of Project2.

Ensure that any intermediary changes performed by Team2 on a copy of Project2 will be subject to the same restrictions as the ones defined in the build policy of Project2.

Project2 will use an automatic build policy. A small team of developers named Team2 will work independently on changes to the project. The Team2 members will not have permissions to Project2.

References:

<https://docs.microsoft.com/en-us/azure/devops/repos/git/manage-your-branches>

QUESTION 30

DRAG DROP

You need to configure Azure Automation for the computers in Pool7.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions**Answer Area**

Run the `Import-AzureRmAutomationDscConfiguration` Azure PowerShell cmdlet.

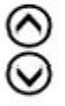
Create a Desired State Configuration (DSC) configuration file that has an extension of `.ps1`.

Run the `New-AzureRmResourceGroupDeployment` Azure PowerShell cmdlet.



Run the `Start-AzureRmAutomationDscCompilationJob` Azure PowerShell cmdlet.

Create an Azure Resource Manager template file that has an extension of `.json`.



A. Answer:

Actions	Answer Area
	Create a Desired State Configuration (DSC) configuration file that has an extension of .ps1.
Run the <code>New-AzureRmResourceGroupDeployment</code> Azure PowerShell cmdlet.	Run the <code>Import-AzureRmAutomationDscConfiguration</code> Azure PowerShell cmdlet.
Create an Azure Resource Manager template file that has an extension of .json.	Run the <code>Start-AzureRmAutomationDscCompilationJob</code> Azure PowerShell cmdlet.

Explanation

Explanation/Reference:
Explanation:

Step 1: Create a Desired State Configuration (DSC) configuration file that has an extension of .ps1.

Step 2: Run the `Import-AzureRmAutomationDscConfiguration` Azure Powershell cmdlet

The `Import-AzureRmAutomationDscConfiguration` cmdlet imports an APS Desired State Configuration (DSC) configuration into Azure Automation. Specify the path of an APS script that contains a single DSC configuration.

Example:

```
PS C:\>Import-AzureRmAutomationDscConfiguration -AutomationAccountName "Contoso17"-ResourceGroupName "ResourceGroup01" -SourcePath "C:\DSC\client.ps1" -Force
```

This command imports the DSC configuration in the file named client.ps1 into the Automation account named Contoso17. The command specifies the Force parameter. If there is an existing DSC configuration, this command replaces it.

Step 3: Run the `Start-AzureRmAutomationDscCompilationJob` Azure Powershell cmdlet

The `Start-AzureRmAutomationDscCompilationJob` cmdlet compiles an APS Desired State Configuration (DSC) configuration in Azure Automation.

References:

<https://docs.microsoft.com/en-us/powershell/module/azurerm.automation/import-azurermautomationdscconfiguration>

<https://docs.microsoft.com/en-us/powershell/module/azurerm.automation/start-azurermautomationdsccompilationjob>
Question Set 3

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Step 1: Create a Desired State Configuration (DSC) configuration file that has an extension of .ps1.

Step 2: Run the Import-AzureRmAutomationDscConfiguration Azure Powershell cmdlet

The Import-AzureRmAutomationDscConfiguration cmdlet imports an APS Desired State Configuration (DSC) configuration into Azure Automation. Specify the path of an APS script that contains a single DSC configuration.

Example:

```
PS C:\>Import-AzureRmAutomationDscConfiguration -AutomationAccountName "Contoso17"-ResourceGroupName "ResourceGroup01" -SourcePath "C:\DSC\client.ps1" -Force
```

This command imports the DSC configuration in the file named client.ps1 into the Automation account named Contoso17. The command specifies the Force parameter. If there is an existing DSC configuration, this command replaces it.

Step 3: Run the Start-AzureRmAutomationDscCompilationJob Azure Powershell cmdlet

The Start-AzureRmAutomationDscCompilationJob cmdlet compiles an APS Desired State Configuration (DSC) configuration in Azure Automation.

References:

<https://docs.microsoft.com/en-us/powershell/module/azurerm.automation/import-azurermautomationdscconfiguration>

<https://docs.microsoft.com/en-us/powershell/module/azurerm.automation/start-azurermautomationdsccompilationjob>

Question Set 3

QUESTION 31

HOTSPOT

Your company is creating a suite of three mobile applications.

You need to control access to the application builds. The solution must be managed at the organization level.

What should you use? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

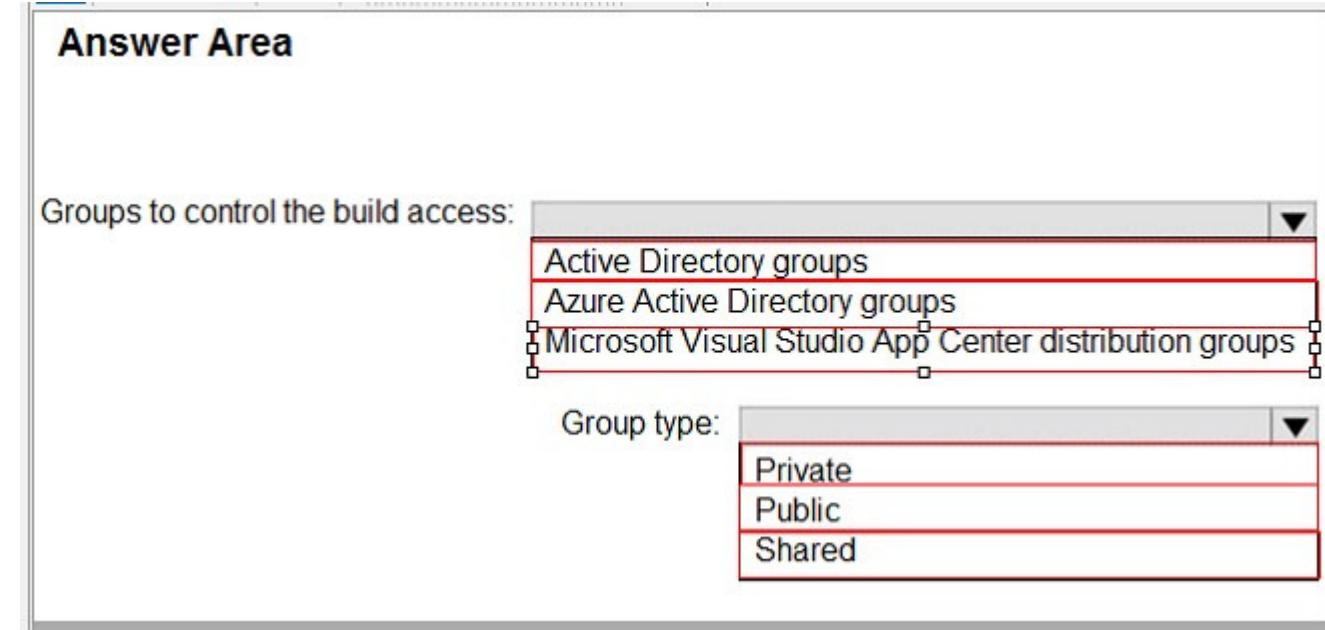
Answer Area

Groups to control the build access:

- Active Directory groups
- Azure Active Directory groups
- Microsoft Visual Studio App Center distribution groups

Group type:

- Private
- Public
- Shared



A. Answer:

Answer Area

Groups to control the build access:

Active Directory groups
Azure Active Directory groups
Microsoft Visual Studio App Center distribution groups

Group type:

Private
Public
Shared

Explanation

Explanation/Reference:

Explanation:

Box 1: Microsoft Visual Studio App Center distribution Groups

Distribution Groups are used to control access to releases. A Distribution Group represents a set of users that can be managed jointly and can have common access to releases. Example of Distribution Groups can be teams of users, like the QA Team or External Beta Testers or can represent stages or rings of releases, such as Staging.

Box 2: Shared

Shared distribution groups are private or public distribution groups that are shared across multiple apps in a single organization. Shared distribution groups eliminate the need to replicate distribution groups across multiple apps.

Note: With the Deploy with App Center Task in Visual Studio Team Services, you can deploy your apps from Azure DevOps (formerly known as VSTS) to App Center. By deploying to App Center, you will be able to distribute your builds to your users.

References: <https://docs.microsoft.com/en-us/appcenter/distribution/groups>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box 1: Microsoft Visual Studio App Center distribution Groups

Distribution Groups are used to control access to releases. A Distribution Group represents a set of users that can be managed jointly and can have common access to releases. Example of Distribution Groups can be teams of users, like the QA Team or External Beta Testers or can represent stages or rings of releases, such as Staging.

Box 2: Shared

Shared distribution groups are private or public distribution groups that are shared across multiple apps in a single organization. Shared distribution groups eliminate the need to replicate distribution groups across multiple apps.

Note: With the Deploy with App Center Task in Visual Studio Team Services, you can deploy your apps from Azure DevOps (formerly known as VSTS) to App Center. By deploying to App Center, you will be able to distribute your builds to your users.

References: <https://docs.microsoft.com/en-us/appcenter/distribution/groups>

QUESTION 32

DRAG DROP

You are configuring Azure DevOps build pipelines.

You plan to use hosted build agents.

Which build agent pool should you use to compile each application type? To answer, drag the appropriate build agent pools to the correct application types. Each build agent pool may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Answer Area

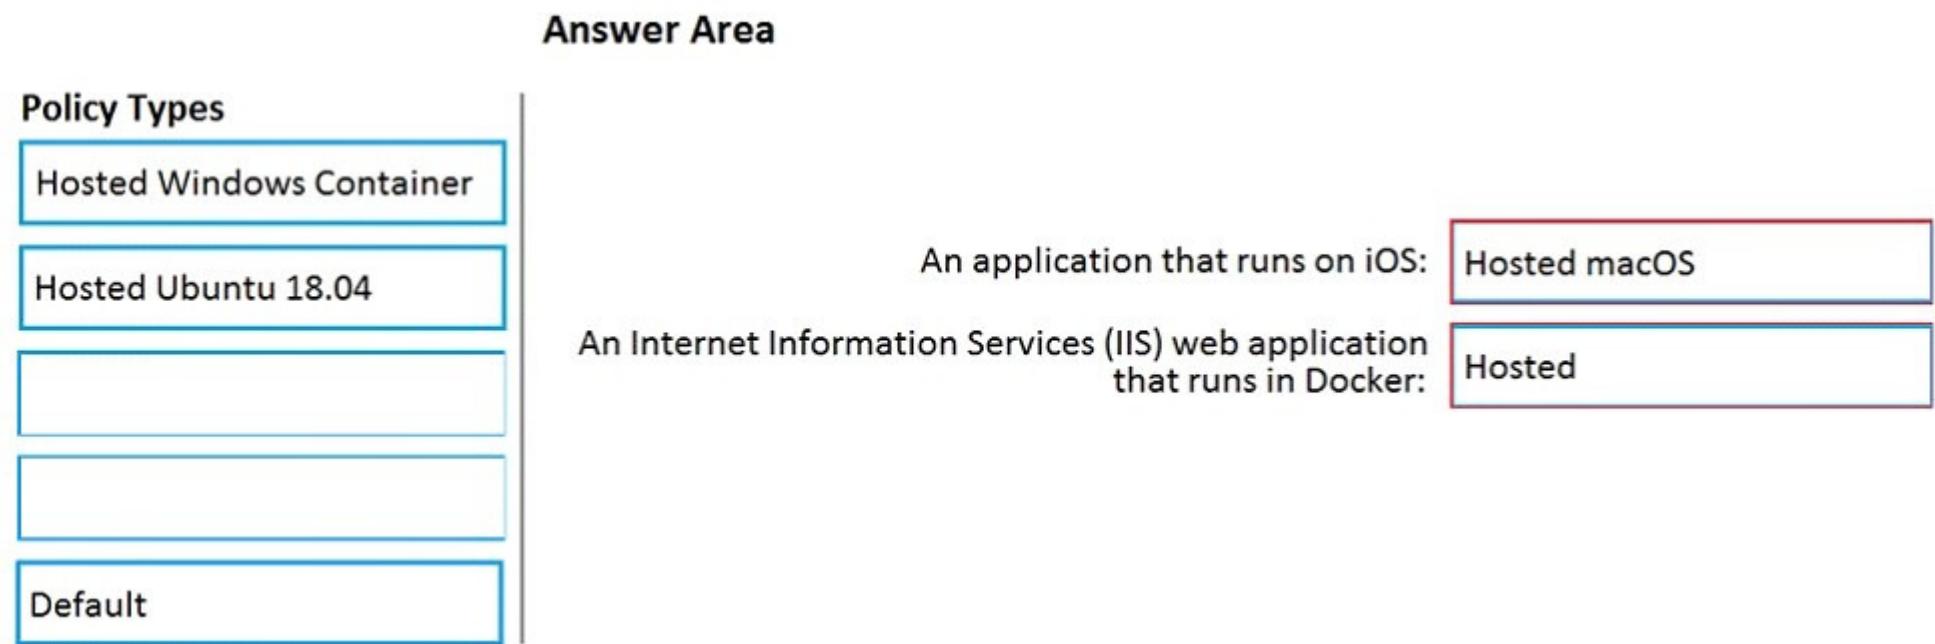
Policy Types

- Hosted Windows Container
- Hosted Ubuntu 18.04
- Hosted macOS
- Hosted
- Default

An application that runs on iOS:

An Internet Information Services (IIS) web application
that runs in Docker:

A. Answer:



Explanation

Explanation/Reference:

Explanation:

Box 1: Hosted macOS

Hosted macOS pool (Azure Pipelines only): Enables you to build and release on macOS without having to configure a self-hosted macOS agent. This option affects where your data is stored.

Box 2: Hosted

Hosted pool (Azure Pipelines only): The Hosted pool is the built-in pool that is a collection of Microsoft-hosted agents.

Incorrect Answers:

Default pool: Use it to register self-hosted agents that you've set up.

Hosted Windows Container pool (Azure Pipelines only): Enabled you to build and release inside Windows containers. Unless you're building using containers, Windows builds should run in the Hosted VS2017 or Hosted pools.

Reference:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/agents/v2-osx>

<https://docs.microsoft.com/en-us/azure/devops/pipelines/agents>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box 1: Hosted macOS

Hosted macOS pool (Azure Pipelines only): Enables you to build and release on macOS without having to configure a self-hosted macOS agent. This option affects where your data is stored.

Box 2: Hosted

Hosted pool (Azure Pipelines only): The Hosted pool is the built-in pool that is a collection of Microsoft-hosted agents.

Incorrect Answers:

Default pool: Use it to register self-hosted agents that you've set up.

Hosted Windows Container pool (Azure Pipelines only): Enabled you to build and release inside Windows containers. Unless you're building using containers, Windows builds should run in the Hosted VS2017 or Hosted pools.

Reference:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/agents/v2-osx>

<https://docs.microsoft.com/en-us/azure/devops/pipelines/agents>

QUESTION 33

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You integrate a cloud-hosted Jenkins server and a new Azure DevOps deployment.

You need Azure DevOps to send a notification to Jenkins when a developer commits changes to a branch in Azure Repos.

Solution: You create an email subscription to an Azure DevOps notification.

Does this meet the goal?

A. Yes

B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

You can create a service hook for Azure DevOps Services and TFS with Jenkins.

References:

<https://docs.microsoft.com/en-us/azure/devops/service-hooks/services/jenkins>

QUESTION 34

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You integrate a cloud-hosted Jenkins server and a new Azure DevOps deployment.

You need Azure DevOps to send a notification to Jenkins when a developer commits changes to a branch in Azure Repos.

Solution: You create a service hook subscription that uses the code pushed event.

Does this meet the goal?

- A. Yes
- B. No

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

You can create a service hook for Azure DevOps Services and TFS with Jenkins.
The code push event is triggered when the code is pushed to a Git repository.

References:

<https://docs.microsoft.com/en-us/azure/devops/service-hooks/services/jenkins>

<https://docs.microsoft.com/en-us/azure/devops/service-hooks/events>

QUESTION 35

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You integrate a cloud-hosted Jenkins server and a new Azure DevOps deployment.

You need Azure DevOps to send a notification to Jenkins when a developer commits changes to a branch in Azure Repos.

Solution: You add a trigger to the build pipeline.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

You can create a service hook for Azure DevOps Services and TFS with Jenkins.

References:

<https://docs.microsoft.com/en-us/azure/devops/service-hooks/services/jenkins>

QUESTION 36

You are automating the build process for a Java-based application by using Azure DevOps.

You need to add code coverage testing and publish the outcomes to the pipeline.

What should you use?

- A. Cobertura
- B. Bullseye Coverage
- C. MSTest
- D. Coverlet
- E. NUnit
- F. Coverage.py

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Use Publish Code Coverage Results task in a build pipeline to publish code coverage results to Azure Pipelines or TFS, which were produced by a build in Cobertura or JaCoCo format.

Incorrect:

Not B: Bullseye Coverage is used for C++ code, and not for Java.

Not D: If you're building on Linux or macOS, you can use Coverlet or a similar tool to collect code coverage metrics. Code coverage results can be published to the server by using the Publish Code Coverage Results task. To leverage this functionality, the coverage tool must be configured to generate results in Cobertura or JaCoCo coverage format.

Not F: Coverage.py is used for Python, not for Java.

Reference:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/tasks/test/publish-code-coverage-results>

QUESTION 37

Your company uses Azure DevOps.

Only users who have accounts in Azure Active Directory can access the Azure DevOps environment.

You need to ensure that only devices that are connected to the on-premises network can access the Azure DevOps environment.

What should you do?

- A. Assign the Stakeholder access level to all users.
- B. In Azure Active Directory, configure risky sign-ins.
- C. In Azure DevOps, configure Security in Project Settings.
- D. In Azure Active Directory, configure conditional access.

Answer: D

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Conditional Access is a capability of Azure Active Directory. With Conditional Access, you can implement automated access control decisions for accessing your cloud apps that are based on conditions.

Conditional Access policies are enforced after the first-factor authentication has been completed.

References:

<https://docs.microsoft.com/en-us/azure/active-directory/conditional-access/overview>

QUESTION 38

You are automating the testing process for your company.

You need to automate UI testing of a web application.

Which framework should you use?

- A. JaCoco
- B. Selenium
- C. Xamarin.UITest
- D. Microsoft.CodeAnalysis

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Performing user interface (UI) testing as part of the release pipeline is a great way of detecting unexpected changes, and need not be difficult. Selenium can be used to test your website during a continuous deployment release and test automation.

References:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/test/continuous-test-selenium?view=azure-devops>

QUESTION 39

You have an Azure DevOps organization named Contoso, an Azure DevOps project named Project1, an Azure subscription named Sub1, and an Azure key vault named vault1.

You need to ensure that you can reference the values of the secrets stored in vault1 in all the pipelines of Project1. The solution must prevent the values from being stored in the pipelines.

What should you do?

- A. Create a variable group in Project1.
- B. Add a secure file to Project1.
- C. Modify the security settings of the pipelines.
- D. Configure the security policy of Contoso.

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Use a variable group to store values that you want to control and make available across multiple pipelines.

References:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/library/variable-groups>

QUESTION 40

DRAG DROP

You are configuring Azure Pipelines for three projects in Azure DevOps as shown in the following table.

Project name	Project Details
Project1	The project team provides preconfigured YAML files that it wants to use to manage future pipeline configuration changes.
Project2	The sensitivity of the project requires that the source code be hosted on the managed Windows server on your company's network.
Project3	The project team requires a centralized version control system to ensure that developers work with the most recent version.

Which version control system should you recommend for each project? To answer, drag the appropriate version control systems to the correct projects. Each version control system may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Version Control Systems	Answer Area
Assembla Subversion	Project1: <input type="text"/>
Bitbucket Cloud	Project2: <input type="text"/>
Git in Azure Repos	Project3: <input type="text"/>
GitHub Enterprise	

A. Answer:

Version Control Systems	Answer Area
Assembla Subversion	Project1: Git in Azure Repos
Bitbucket Cloud	Project2: GitHub Enterprise
Git in Azure Repos	Project3: Bitbucket Cloud
GitHub Enterprise	

Explanation

Explanation/Reference:

Explanation:

Project1: Git in Azure Repos

Project2: Github Enterprise

GitHub Enterprise is the on-premises version of GitHub.com. GitHub Enterprise includes the same great set of features as GitHub.com but packaged for running on your organization's local network. All repository data is stored on machines that you control, and access is integrated with your organization's authentication system (LDAP, SAML, or CAS).

Project3: Bitbucket cloud

One downside, however, is that Bitbucket does not include support for SVN but this can be easily amended migrating the SVN repos to Git with tools such as SVN Mirror for Bitbucket .

Note: SVN is a centralized version control system.

Incorrect Answers:

Bitbucket:

Bitbucket comes as a distributed version control system based on Git.

Note: A source control system, also called a version control system, allows developers to collaborate on code and track changes. Source control is an essential tool for multi-developer projects.

Our systems support two types of source control: Git (distributed) and Team Foundation Version Control (TFVC). TFVC is a centralized, client-server system. In both Git and TFVC, you can check in files and organize files in folders, branches, and repositories.

References:

<https://www.azuredevopslabs.com/labs/azuredevops/yaml/>

<https://enterprise.github.com/faq>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Project1:Git in Azure Repos

Project2: Github Enterprise

GitHub Enterprise is the on-premises version of GitHub.com. GitHub Enterprise includes the same great set of features as GitHub.com but packaged for running on your organization's local network. All repository data is stored on machines that you control, and access is integrated with your organization's authentication system (LDAP, SAML, or CAS).

Project3: Bitbucket cloud

One downside, however, is that Bitbucket does not include support for SVN but this can be easily amended migrating the SVN repos to Git with tools such as SVN Mirror for Bitbucket .

Note: SVN is a centralized version control system.

Incorrect Answers:

Bitbucket:

Bitbucket comes as a distributed version control system based on Git.

Note: A source control system, also called a version control system, allows developers to collaborate on code and track changes. Source control is an essential tool for multi-developer projects.

Our systems support two types of source control: Git (distributed) and Team Foundation Version Control (TFVC). TFVC is a centralized, client-server system. In both Git and TFVC, you can check in files and organize files in folders, branches, and repositories.

References:

<https://www.azuredevopslabs.com/labs/azuredevops/yaml/>

<https://enterprise.github.com/faq>

QUESTION 41

Your team uses an agile development approach.

You need to recommend a branching strategy for the team's Git repository. The strategy must meet the following requirements.

Provide the ability to work on multiple independent tasks in parallel.

Ensure that checked-in code remains in a releasable state always.

Ensure that new features can be abandoned at any time.

Encourage experimentation.

What should you recommend?

- A. a single long-running branch without forking
- B. multiple long-running branches
- C. a single fork per team member
- D. a single long-running branch with multiple short-lived topic branches

Answer: D

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Topic branches, however, are useful in projects of any size. A topic branch is a short-lived branch that you create and use for a single particular feature or related work. This is something you've likely never done with a VCS before because it's generally too

expensive to create and merge branches. But in Git it's common to create, work on, merge, and delete branches several times a day.

Reference:

<https://git-scm.com/book/en/v2/Git-Branching-Branching-Workflows>

QUESTION 42

Your company has a project in Azure DevOps for a new web application.

The company identifies security as one of the highest priorities.

You need to recommend a solution to minimize the likelihood that infrastructure credentials will be leaked.

What should you recommend?

- A. Add a Run Inline Azure PowerShell task to the pipeline.
- B. Add a PowerShell task to the pipeline and run Set-AzureKeyVaultSecret.
- C. Add a Azure Key Vault task to the pipeline.
- D. Add Azure Key Vault references to Azure Resource Manger templates.

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Azure Key Vault provides a way to securely store credentials and other keys and secrets.

The Set-AzureKeyVaultSecret cmdlet creates or updates a secret in a key vault in Azure Key Vault.

References:

<https://docs.microsoft.com/en-us/powershell/module/azurerm.keyvault/set-azurekeyvaultsecret>

QUESTION 43

DRAG DROP

You provision an Azure Kubernetes Service (AKS) cluster that has RBAC enabled. You have a Helm chart for a client application.

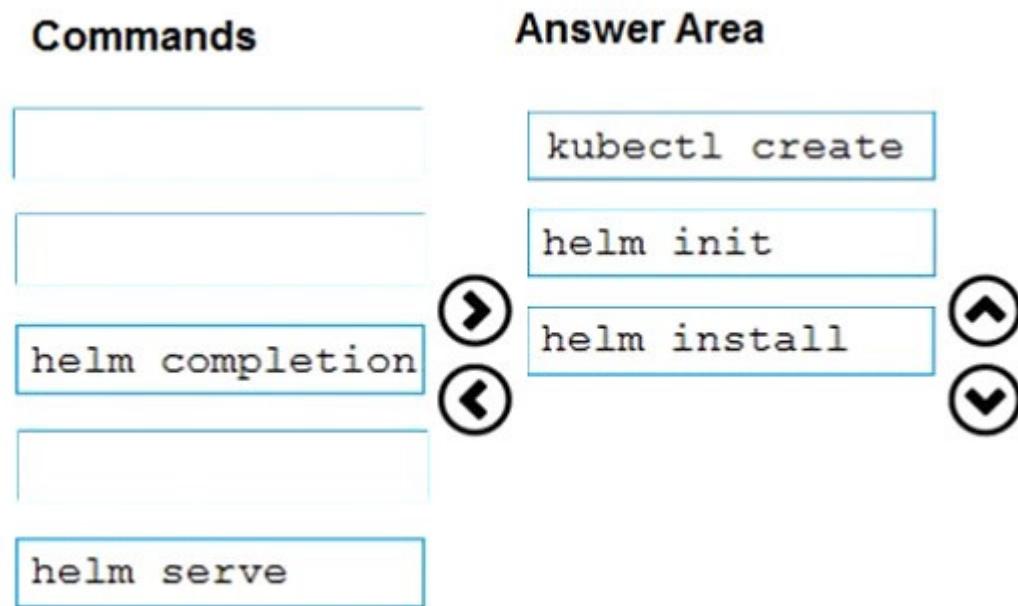
You need to configure Helm and Tiller on the cluster and install the chart.

Which three commands should you recommend be run in sequence? To answer, move the appropriate commands from the list of commands to the answer area and arrange them in the correct order.

Select and Place:

Commands**Answer Area**`helm install``kubectl create``helm completion``helm init``helm serve`

A. Answer:



Explanation

Explanation/Reference:

Explanation:

Step 1: Kubectl create

You can add a service account to Tiller using the --service-account <NAME> flag while you're configuring Helm (step 2 below). As a prerequisite, you'll have to create a role binding which specifies a role and a service account name that have been set up in advance.

Example: Service account with cluster-admin role

```
$ kubectl create -f rbac-config.yaml
serviceaccount "tiller" created
clusterrolebinding "tiller" created
$ helm init --service-account tiller
```

Step 2: helm init

To deploy a basic Tiller into an AKS cluster, use the helm init command.

Step 3: helm install

To install charts with Helm, use the helm install command and specify the name of the chart to install.

References:

<https://docs.microsoft.com/en-us/azure/aks/kubernetes-helm>

https://docs.helm.sh/using_helm/#tiller-namespaces-and-rbac

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Step 1: Kubectl create

You can add a service account to Tiller using the --service-account <NAME> flag while you're configuring Helm (step 2 below). As a prerequisite, you'll have to create a role binding which specifies a role and a service account name that have been set up in advance.

Example: Service account with cluster-admin role

```
$ kubectl create -f rbac-config.yaml
serviceaccount "tiller" created
clusterrolebinding "tiller" created
$ helm init --service-account tiller
```

Step 2: helm init

To deploy a basic Tiller into an AKS cluster, use the helm init command.

Step 3: helm install

To install charts with Helm, use the helm install command and specify the name of the chart to install.

References:

<https://docs.microsoft.com/en-us/azure/aks/kubernetes-helm>

https://docs.helm.sh/using_helm/#tiller-namespaces-and-rbac

QUESTION 44

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You need to recommend an integration strategy for the build process of a Java application. The solution must meet the following requirements:

The builds must access an on-premises dependency management system.

The build outputs must be stored as Server artifacts in Azure DevOps.

The source code must be stored in a Git repository in Azure DevOps.

Solution: Configure an Octopus Tentacle on an on-premises machine. Use the Package Application task in the build pipeline.

Does this meet the goal?

A. Yes

B. No

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Octopus Deploy is an automated deployment server that makes it easy to automate deployment of ASP.NET web applications, Java applications, NodeJS application and custom scripts to multiple environments.

Octopus can be installed on various platforms including Windows, Mac and Linux. It can also be integrated with most version control tools including VSTS and GIT.

When you deploy software to Windows servers, you need to install Tentacle, a lightweight agent service, on your Windows servers so they can communicate with the Octopus server.

When defining your deployment process, the most common step type will be a package step. This step deploys your packaged application onto one or more deployment targets.

When deploying a package you will need to select the machine role that the package will be deployed to.

References:

<https://octopus.com/docs/deployment-examples/package-deployments>

<https://explore.emtecinc.com/blog/octopus-for-automated-deployment-in-devops-models>

QUESTION 45

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You need to recommend an integration strategy for the build process of a Java application. The solution must meet the following requirements:

The builds must access an on-premises dependency management system.

The build outputs must be stored as Server artifacts in Azure DevOps.

The source code must be stored in a Git repository in Azure DevOps.

Solution: Install and configure a self-hosted build agent on an on-premises machine. Configure the build pipeline to use the Default agent pool. Include the Java Tool Installer task in the build pipeline.

Does this meet the goal?

A. Yes

B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Instead use Octopus Tentacle.

References:

<https://explore.emtecinc.com/blog/octopus-for-automated-deployment-in-devops-models>

QUESTION 46

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution,

while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You need to recommend an integration strategy for the build process of a Java application. The solution must meet the following requirements:
The builds must access an on-premises dependency management system.

The build outputs must be stored as Server artifacts in Azure DevOps.

The source code must be stored in a Git repository in Azure DevOps.

Solution: Configure the build pipeline to use a Hosted VS 2019 agent pool. Include the Java Tool Installer task in the build pipeline.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Instead use Octopus Tentacle.

References:

<https://explore.emtecinc.com/blog/octopus-for-automated-deployment-in-devops-models>

QUESTION 47

You are designing the development process for your company.

You need to recommend a solution for continuous inspection of the company's code base to locate common code patterns that are known to be problematic.

What should you include in the recommendation?

- A. Microsoft Visual Studio test plans
- B. Gradle wrapper scripts
- C. SonarCloud analysis
- D. the JavaScript task runner

Answer: C

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

SonarCloud is a cloud service offered by SonarSource and based on SonarQube. SonarQube is a widely adopted open source platform to inspect continuously the quality of source code and detect bugs, vulnerabilities and code smells in more than 20 different

languages.

Note: The SonarCloud Azure DevOps extension brings everything you need to have your projects analyzed on SonarCloud very quickly.

Incorrect Answers:

A: Test plans are used to group together test suites and individual test cases. This includes static test suites, requirement-based suites, and query-based suites.

References:

<https://docs.travis-ci.com/user/sonarcloud/>

<https://sonarcloud.io/documentation/integrations/vsts/>

QUESTION 48

SIMULATION

You need to ensure that an Azure web app named az400-9940427-main can retrieve secrets from an Azure key vault named az400-9940427-kv1 by using a system managed identity.

The solution must use the principle of least privilege.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

1. In Azure portal navigate to the az400-9940427-main app.
2. Scroll down to the Settings group in the left navigation.
3. Select Managed identity.
4. Within the System assigned tab, switch Status to On. Click Save.

The screenshot shows the Azure portal interface for managing identities. The top navigation bar includes 'Dashboard' and the current resource name 'systemassigned-linux - Identity'. The left sidebar lists various settings: Application settings, Authentication / Authorization, Application Insights, Identity (which is selected and highlighted in blue), Backups, Custom domains, SSL settings, Networking, Scale up (App Service plan), and Scale out (App Service plan). The main content area is titled 'systemassigned-linux - Identity' and describes system assigned managed identities. It shows the 'System assigned' tab is selected, with a 'User assigned (preview)' option available. Below this, there is a 'Status' button with 'Off' and 'On' options, where 'On' is currently selected. An 'Object ID' field contains the value '7283a4ee-ac06-4f67-b8e7-513d24f010d1'. A note at the bottom states: 'This resource is registered with Azure Active Directory. You can control its access to services like A...'.

References:

<https://docs.microsoft.com/en-us/azure/app-service/overview-managed-identity>

Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

1. In Azure portal navigate to the az400-9940427-main app.
2. Scroll down to the Settings group in the left navigation.
3. Select Managed identity.
4. Within the System assigned tab, switch Status to On. Click Save.

A system assigned managed identity enables Azure resources to authenticate to cloud services (e.g. Azure Key Vault). The lifecycle of this type of managed identity is tied to the lifecycle of this resource. Additionally, each resource (e.g. App Service) has its own managed identity.

System assigned **User assigned (preview)**

Status **Off** **On**

Object ID 7283a4ee-ac06-4f67-b8e7-513d24f010d1

This resource is registered with Azure Active Directory. You can control its access to services like Azure Key Vault, Azure Storage, and more.

References:
<https://docs.microsoft.com/en-us/azure/app-service/overview-managed-identity>

QUESTION 49

Your company builds a multi-tier web application.

You use Azure DevOps and host the production application on Azure virtual machines.

Your team prepares an Azure Resource Manager template of the virtual machine that you will use to test new features.

You need to create a staging environment in Azure that meets the following requirements:

Minimizes the cost of Azure hosting

Provisions the virtual machines automatically

Uses the custom Azure Resource Manager template to provision the virtual machines

What should you do?

- A. In Azure Cloud Shell, run Azure CLI commands to create and delete the new virtual machines in a staging resource group.
- B. In Azure DevOps, configure new tasks in the release pipeline to deploy to Azure Cloud Services.
- C. From Azure Cloud Shell, run Azure PowerShell commands to create and delete the new virtual machines in a staging resource group.
- D. In Azure DevOps, configure new tasks in the release pipeline to create and delete the virtual machines in Azure DevTest Labs.

Answer: D

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

You can use the Azure DevTest Labs Tasks extension that's installed in Azure DevOps to easily integrate your CI/CD build-and-release pipeline with Azure DevTest Labs. The extension installs three tasks:

Create a VM

Create a custom image from a VM

Delete a VM

The process makes it easy to, for example, quickly deploy a "golden image" for a specific test task and then delete it when the test is finished.

References: <https://docs.microsoft.com/en-us/azure/lab-services/devtest-lab-integrate-ci-cd-vsts>

QUESTION 50

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You need to recommend an integration strategy for the build process of a Java application. The solution must meet the following requirements:

The builds must access an on-premises dependency management system.

The build outputs must be stored as Server artifacts in Azure DevOps.

The source code must be stored in a Git repository in Azure DevOps.

Solution: Configure the build pipeline to use a Hosted Ubuntu agent pool. Include the Java Tool Installer task in the build pipeline.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Instead use Octopus Tentacle.

Reference:

<https://explore.emtecinc.com/blog/octopus-for-automated-deployment-in-devops-models>

QUESTION 51

DRAG DROP

You are implementing an Azure DevOps strategy for mobile devices using App Center.

You plan to use distribution groups to control access to releases.

You need to create the distribution groups shown in the following table.

Name	Use
Group1	Application testers who are invited by email
Group2	Early release users who use unauthenticated public links
Group3	Application testers for all the apps of your company

Which type of distribution group should you use for each group? To answer, drag the appropriate group types to the correct locations. Each group type may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Answer Area

Private



Group1:

Public



Group2:

Shared



Group3:

A. Answer:

Answer Area

	Group1: Private
	Group2: Public
	Group3: Shared

Explanation

Explanation/Reference:

Explanation:

Box1: Private

In App Center, distribution groups are private by default. Only testers invited via email can access the releases available to this group.

Box 2: Public

Distribution groups must be public to enable unauthenticated installs from public links.

Box 3: Shared

Shared distribution groups are private or public distribution groups that are shared across multiple apps in a single organization.

Reference:

<https://docs.microsoft.com/en-us/appcenter/distribution/groups>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box1: Private

In App Center, distribution groups are private by default. Only testers invited via email can access the releases available to this group.

Box 2: Public

Distribution groups must be public to enable unauthenticated installs from public links.

Box 3: Shared

Shared distribution groups are private or public distribution groups that are shared across multiple apps in a single organization.

Reference:

<https://docs.microsoft.com/en-us/appcenter/distribution/groups>

QUESTION 52**SIMULATION**

You need to ensure that the <https://contoso.com/statushook> webhook is called every time a repository named az40010480345acr1 receives a new version of an image named dotnetapp.
To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

1. Sign in to the Azure portal.
2. Navigate to the container registry az40010480345acr1.
3. Under Services, select Webhooks.
4. Select the existing webhook https://contoso.com/statushook, and double-click on it to get its properties.
5. For Trigger actions select image push

Example web hook:

The screenshot shows the Azure Container Registry interface for 'myregistry'. On the left, the sidebar includes 'Overview', 'Activity log', 'Access control (IAM)', 'Tags', 'Quick start', 'SETTINGS' (with 'Access keys', 'Locks', 'Automation script'), 'SERVICES' (with 'Repositories' and 'Webhooks' selected), and 'Replications (Preview)'. The main area displays a 'Webhooks' blade with a search bar and a table showing 'No result'. A 'Create webhook' dialog is open on the right, containing fields for 'Webhook name' (set to 'myacnwebhook'), 'Location' (set to 'East US'), 'Service URI' (set to 'https://contoso.com/acreventendpoint'), 'Custom headers' (set to 'Content-Type: application/json'), and 'Actions' (set to 'Push'). The 'Actions' field is highlighted with a red border. Below the dialog are 'Status' (set to 'On') and 'Scope' (empty input field). At the bottom are 'Pin to dashboard' and a 'Create' button.

Reference:

<https://docs.microsoft.com/en-us/azure/container-registry/container-registry-webhook>

Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation

1. Sign in to the Azure portal.

2. Navigate to the container registry az4001048034acr1.

3. Under Services, select Webhooks.

4. Select the existing webhook https://contoso.com/statushook, and double-click on it to get its properties.

5. For Trigger actions select image push

Example web hook:

The screenshot shows the Azure Container Registry interface for 'myregistry'. On the left, the navigation menu includes 'Overview', 'Activity log', 'Access control (IAM)', 'Tags', 'Quick start', 'SETTINGS' (with 'Access keys', 'Locks', and 'Automation script'), 'SERVICES' (with 'Repositories' and 'Webhooks' selected), and 'Replications (Preview)'. Below these are 'SUPPORT + TROUBLESHOOTING' and a 'More' section. The main area displays the 'Webhooks' blade with a search bar and a table titled 'No result'. A 'Create webhook' dialog is open on the right, prompting for 'Webhook name' (set to 'myacrwebhook'), 'Location' (set to 'East US'), 'Service URI' (set to 'https://contoso.com/acreventoendpoint'), 'Custom headers' (set to 'Content-Type: application/json'), 'Actions' (set to 'Push' which is highlighted with a red box), 'Status' (set to 'On'), and 'Scope' (set to 'Enter the webhook scope'). There is also a 'Pin to dashboard' checkbox and a 'Create' button at the bottom.

Reference:
<https://docs.microsoft.com/en-us/azure/container-registry/container-registry-webhook>

QUESTION 53

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You integrate a cloud-hosted Jenkins server and a new Azure DevOps deployment.

You need Azure DevOps to send a notification to Jenkins when a developer commits changes to a branch in Azure Repos.

Solution: You create a service hook subscription that uses the build completed event.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:
Explanation

Explanation/Reference:
Explanation:

You can create a service hook for Azure DevOps Services and TFS with Jenkins.
However, the service subscription event should use the code pushed event, is triggered when the code is pushed to a Git repository.

QUESTION 54

HOTSPOT

You need to create deployment files for an Azure Kubernetes Service (AKS) cluster. The deployments must meet the provisioning storage requirements shown in the following table.

Deployment	Requirement
Deployment 1	Use files stored on an SMB-based share from the container's file system.
Deployment 2	Use files on a managed disk from the container's file system.
Deployment 3	Securely access X.509 certificates from the container's file system.

Which resource type should you use for each deployment? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Deployment 1:

azurekeyvault-flexvolume
blobfuse-flexvol
kubernetes.io/azure-disk
kubernetes.io/azure-file
volume.beta.kubernetes.io/storage-provisioner

Deployment 2:

azurekeyvault-flexvolume
blobfuse-flexvol
kubernetes.io/azure-disk
kubernetes.io/azure-file
volume.beta.kubernetes.io/storage-provisioner

Deployment 3:

azurekeyvault-flexvolume
blobfuse-flexvol
kubernetes.io/azure-disk
kubernetes.io/azure-file
volume.beta.kubernetes.io/storage-provisioner

A. Answer:

Answer Area

Deployment 1:

azurekeyvault-flexvolume
blobfuse-flexvol
kubernetes.io/azure-disk
kubernetes.io/azure-file
volume.beta.kubernetes.io/storage-provisioner

Deployment 2:

azurekeyvault-flexvolume
blobfuse-flexvol
kubernetes.io/azure-disk
kubernetes.io/azure-file
volume.beta.kubernetes.io/storage-provisioner

Deployment 3:

azurekeyvault-flexvolume
blobfuse-flexvol
kubernetes.io/azure-disk
kubernetes.io/azure-file
volume.beta.kubernetes.io/storage-provisioner

Explanation

Explanation/Reference:
Explanation:

Deployment 1: Kubernetes.io/azure-file
You can use Azure Files to connect using the Server Message Block (SMB) protocol.

Deployment 2: Kubernetes.io/azure-disk

Deployment 3: azurekeyvault-flexvolume
azurekeyvault-flexvolume: Key Vault FlexVolume: Seamlessly integrate your key management systems with Kubernetes.
Secrets, keys, and certificates in a key management system become a volume accessible to pods. Once the volume is mounted, its data is available directly in the container filesystem for your application.

Incorrect Answers:
blobfuse-flexvolume: This driver allows Kubernetes to access virtual filesystem backed by the Azure Blob storage.

References:
<https://docs.microsoft.com/bs-cyrillic/azure/aks/azure-files-dynamic-pv>

<https://docs.microsoft.com/en-us/azure/aks/azure-disks-dynamic-pv>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation

Deployment 1: Kubernetes.io/azure-file

You can use Azure Files to connect using the Server Message Block (SMB) protocol.

Deployment 2: Kubernetes.io/azure-disk

Deployment 3: azurekeyvault-flexvolume

azurekeyvault-flexvolume: Key Vault FlexVolume: Seamlessly integrate your key management systems with Kubernetes.

Secrets, keys, and certificates in a key management system become a volume accessible to pods. Once the volume is mounted, its data is available directly in the container filesystem for your application.

Incorrect Answers:

blobfuse-flexvolume: This driver allows Kubernetes to access virtual filesystem backed by the Azure Blob storage.

References:

<https://docs.microsoft.com/bs-cyrillic/azure/aks/azure-files-dynamic-pv>

<https://docs.microsoft.com/en-us/azure/aks/azure-disks-dynamic-pv>

QUESTION 55

You create a Microsoft ASP.NET Core application.

You plan to use Azure Key Vault to provide secrets to the application as configuration data.

You need to create a Key Vault access policy to assign secret permissions to the application. The solution must use the principle of least privilege.

Which secret permissions should you use?

- A. List only
- B. Get only
- C. Get and List

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation

Application data plane permissions:

Keys: sign

Secrets: get

Reference:

<https://docs.microsoft.com/en-us/azure/key-vault/key-vault-secure-your-key-vault>

QUESTION 56

DRAG DROP

You need to recommend a solution for deploying charts by using Helm and Tiller to Azure Kubernetes Service (AKS) in an RBAC-enabled cluster.

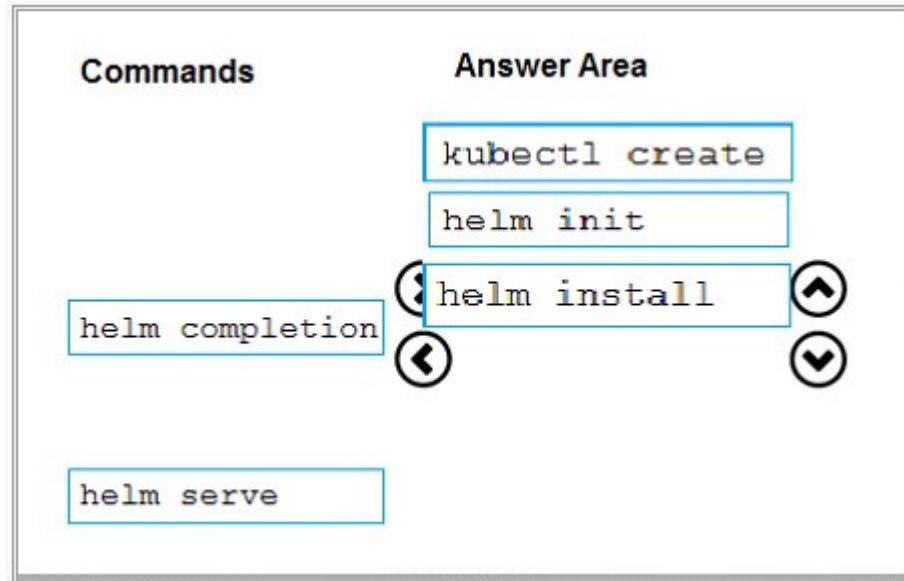
Which three commands should you recommend be run in sequence? To answer, move the appropriate commands from the list of commands to the answer area and arrange them in the correct order.

Select and Place:

Commands	Answer Area
helm install	0
kubectl create	0
helm completion	0
helm init	
helm serve	

The 'Commands' column contains five items: 'helm install', 'kubectl create', 'helm completion', 'helm init', and 'helm serve'. The 'Answer Area' column has five slots, each with a red border and a small red number '0' in the top right corner. The 'helm completion' item in the 'Commands' column has a circular selection icon with a red outline and a black arrow pointing to the right, indicating it is selected for dragging. The 'helm init' item in the 'Commands' column also has a similar circular selection icon with a black arrow pointing to the left, indicating it is selected for dragging. The other three items ('helm install', 'kubectl create', and 'helm serve') do not have these selection icons.

A. Answer:



Explanation

Explanation/Reference:

Explanation:

Step 1: Kubectl create

You can add a service account to Tiller using the --service-account <NAME> flag while you're configuring Helm (step 2 below). As a prerequisite, you'll have to create a role binding which specifies a role and a service account name that have been set up in advance.

Example: Service account with cluster-admin role

```
$ kubectl create -f rbac-config.yaml
serviceaccount "tiller" created
clusterrolebinding "tiller" created
$ helm init --service-account tiller
```

Step 2: helm init

To deploy a basic Tiller into an AKS cluster, use the helm init command.

Step 3: helm install

To install charts with Helm, use the helm install command and specify the name of the chart to install.

References:

<https://docs.microsoft.com/en-us/azure/aks/kubernetes-helm>

https://docs.helm.sh/using_helm/#tiller-namespaces-and-rbac

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Step 1: Kubectl create

You can add a service account to Tiller using the --service-account <NAME> flag while you're configuring Helm (step 2 below). As a prerequisite, you'll have to create a role binding which specifies a role and a service account name that have been set up in advance.

Example: Service account with cluster-admin role
\$ kubectl create -f rbac-config.yaml
serviceaccount "tiller" created
clusterrolebinding "tiller" created
\$ helm init --service-account tiller

Step 2: helm init
To deploy a basic Tiller into an AKS cluster, use the helm init command.

Step 3: helm install
To install charts with Helm, use the helm install command and specify the name of the chart to install.

References:
<https://docs.microsoft.com/en-us/azure/aks/kubernetes-helm>
https://docs.helm.sh/using_helm/#tiller-namespaces-and-rbac

QUESTION 57

DRAG DROP

Your company has a project in Azure DevOps.

You plan to create a release pipeline that will deploy resources by using Azure Resource Manager templates. The templates will reference secrets stored in Azure Key Vault.

You need to recommend a solution for accessing the secrets stored in the key vault during deployments. The solution must use the principle of least privilege.

What should you include in the recommendation? To answer, drag the appropriate configurations to the correct targets. Each configuration may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Configurations	Answer Area
A Key Vault access policy	Enable key vaults for template deployment by using: <input type="text"/>
A Key Vault advanced access policy	Restrict access to the secrets in Key Vault by using: <input type="text"/>
RBAC	

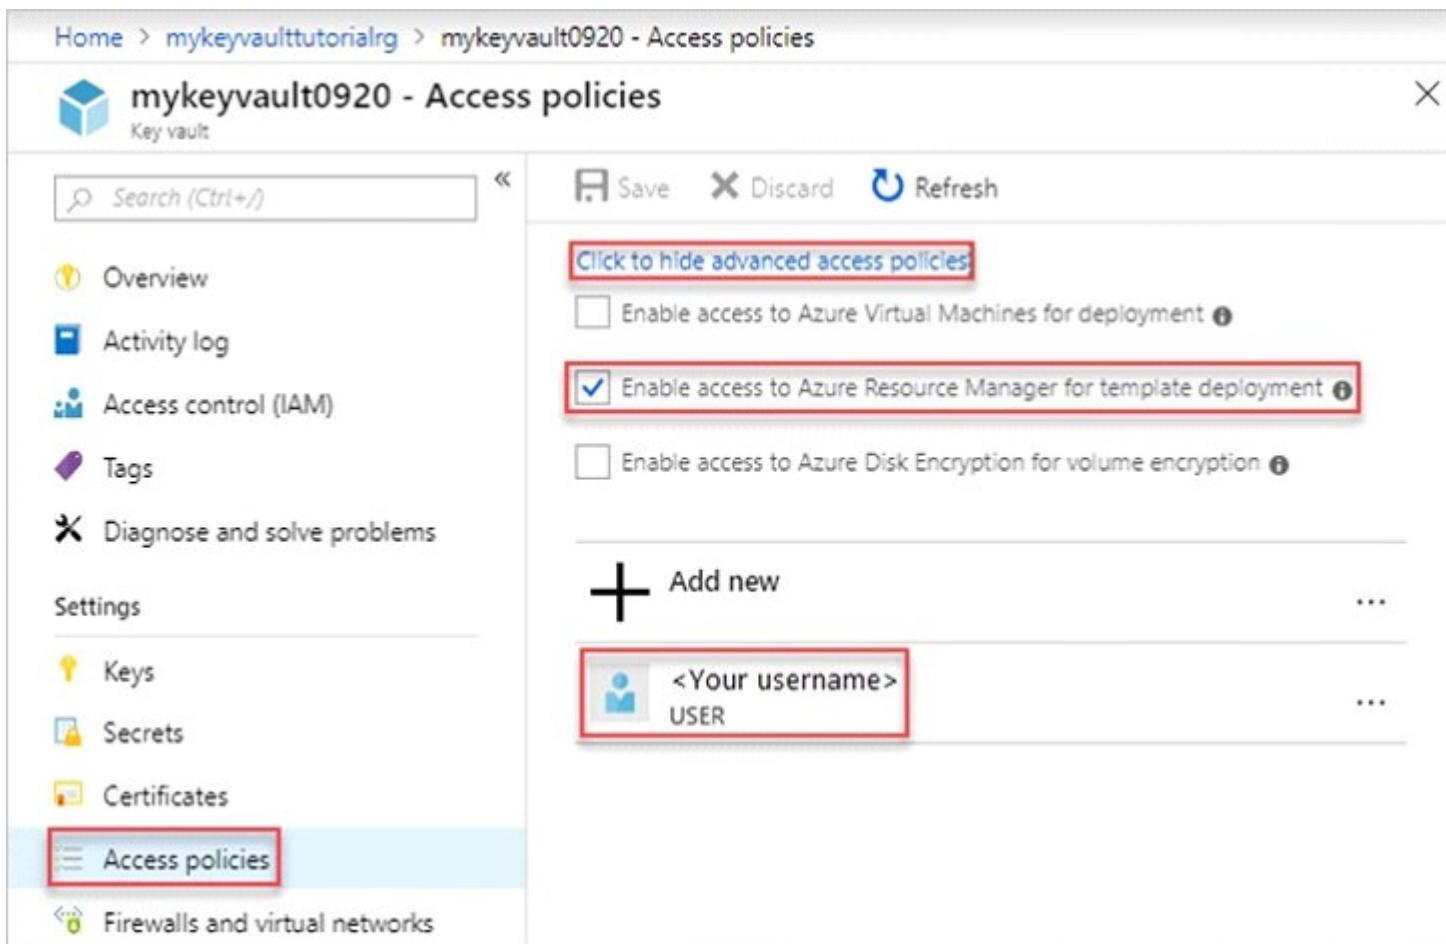
A. Answer:

Configurations	Answer Area
A Key Vault access policy	Enable key vaults for template deployment by using: A Key Vault advanced access policy
	Restrict access to the secrets in Key Vault by using: RBAC

Explanation

Explanation/Reference:
Explanation:

Box 1: A key Vault advanced access policy



Box 2: RBAC

Management plane access control uses RBAC.

The management plane consists of operations that affect the key vault itself, such as:

Creating or deleting a key vault.

Getting a list of vaults in a subscription.

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation

Box 1: A key Vault advanced access policy

The screenshot shows the 'Access policies' section of the Azure Key Vault 'mykeyvault0920'. The left sidebar lists 'Overview', 'Activity log', 'Access control (IAM)', 'Tags', 'Diagnose and solve problems', 'Settings' (with 'Keys', 'Secrets', 'Certificates'), 'Access policies' (which is selected and highlighted in blue), and 'Firewalls and virtual networks'. The main pane displays an 'Advanced access policies' section with three checkboxes:

- 'Enable access to Azure Virtual Machines for deployment' (unchecked)
- 'Enable access to Azure Resource Manager for template deployment' (checked)
- 'Enable access to Azure Disk Encryption for volume encryption' (unchecked)

Below this is an 'Add new' button followed by a row for a user named '<Your username>' (USER), which is also highlighted with a red box.

Box 2: RBAC

Management plane access control uses RBAC.

The management plane consists of operations that affect the key vault itself, such as:

Creating or deleting a key vault.

Getting a list of vaults in a subscription.

Retrieving Key Vault properties (such as SKU and tags).

Setting Key Vault access policies that control user and application access to keys and secrets.

References:
<https://docs.microsoft.com/en-us/azure/azure-resource-manager/resource-manager-tutorial-use-key-vault>

QUESTION 58

DRAG DROP

You need to configure access to Azure DevOps agent pools to meet the following requirements:

Use a project agent pool when authoring build or release pipelines.

View the agent pool and agents of the organization.

Use the principle of least privilege.

Which role memberships are required for the Azure DevOps organization and the project? To answer, drag the appropriate role memberships to the correct targets. Each role membership may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Roles	Answer Area
<input checked="" type="checkbox"/> Administrator	Organization: <input type="text"/>
<input type="checkbox"/> Reader	Project: <input type="text"/>
<input type="checkbox"/> Service Account	
<input type="checkbox"/> User	

A. Answer:

Roles	Answer Area
Administrator	Organization: Reader Project: Service Account
User	

Explanation

Explanation/Reference:

Explanation:

Box 1: Reader

Members of the Reader role can view the organization agent pool as well as agents. You typically use this to add operators that are responsible for monitoring the agents and their health.

Box 2: Service account

Members of the Service account role can use the organization agent pool to create a project agent pool in a project. If you follow the guidelines above for creating new project agent pools, you typically do not have to add any members here.

Incorrect Answers:

In addition to all the permissions given the Reader and the Service Account role, members of the administrator role can register or unregister agents from the organization agent pool. They can also refer to the organization agent pool when creating a project agent pool in a project. Finally, they can also manage membership for all roles of the organization agent pool. The user that created the organization agent pool is automatically added to the Administrator role for that pool.

References: <https://docs.microsoft.com/en-us/azure/devops/pipelines/agents/pools-queues>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box 1: Reader

Members of the Reader role can view the organization agent pool as well as agents. You typically use this to add operators that are responsible for monitoring the agents and their health.

Box 2: Service account

Members of the Service account role can use the organization agent pool to create a project agent pool in a project. If you follow the guidelines above for creating new project agent pools, you typically do not have to add any members here.

Incorrect Answers:

In addition to all the permissions given the Reader and the Service Account role, members of the administrator role can register or unregister agents from the organization agent pool. They can also refer to the organization agent pool when creating a project agent pool in a project. Finally, they can also manage membership for all roles of the organization agent pool. The user that created the organization agent pool is automatically added to the Administrator role for that pool.

References: <https://docs.microsoft.com/en-us/azure/devops/pipelines/agents/pools-queues>

QUESTION 59

You have a branch policy in a project in Azure DevOps. The policy requires that code always builds successfully.

You need to ensure that a specific user can always merge changes to the master branch, even if the code fails to compile. The solution must use the principle of least privilege.

What should you do?

- A. Add the user to the Build Administrators group.
- B. Add the user to the Project Administrators group.
- C. From the Security settings of the repository, modify the access control for the user.
- D. From the Security settings of the branch, modify the access control for the user.

Answer: D

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

In some cases, you need to bypass policy requirements so you can push changes to the branch directly or complete a pull request even if branch policies are not satisfied. For these situations, grant the desired permission from the previous list to a user or group. You can scope this permission to an entire project, a repo, or a single branch. Manage this permission along with other Git permissions.

References: <https://docs.microsoft.com/en-us/azure/devops/repos/git/branch-policies>

QUESTION 60

Your company uses a Git repository in Azure Repos to manage the source code of a web application. The master branch is protected from direct updates. Developers work on new features in the topic branches.

Because of the high volume of requested features, it is difficult to follow the history of the changes to the master branch.

You need to enforce a pull request merge strategy. The strategy must meet the following requirements:

Consolidate commit histories.

Merge the changes into a single commit.

Which merge strategy should you use in the branch policy?

- A. squash merge
- B. fast-forward merge
- C. Git fetch
- D. no-fast-forward merge

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Squash merging is a merge option that allows you to condense the Git history of topic branches when you complete a pull request. Instead of each commit on the topic branch being added to the history of the default branch, a squash merge takes all the file changes and adds them to a single new commit on the default branch.

A simple way to think about this is that squash merge gives you just the file changes, and a regular merge gives you the file changes and the commit history.

Note: Squash merging keeps your default branch histories clean and easy to follow without demanding any workflow changes on your team. Contributors to the topic branch work how they want in the topic branch, and the default branches keep a linear history through the use of squash merges. The commit history of a master branch updated with squash merges will have one commit for each merged branch. You can step through this history commit by commit to find out exactly when work was done.

References: <https://docs.microsoft.com/en-us/azure/devops/repos/git/merging-with-squash>

QUESTION 61

Your company uses cloud-hosted Jenkins for builds.

You need to ensure that Jenkins can retrieve source code from Azure Repos.

Which three actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Create a webhook in Jenkins.
- B. Add the Team Foundation Server (TFS) plug-in to Jenkins.
- C. Add the dev.azure.com hostname to your Jenkins account.
- D. Create a personal access token in your Azure DevOps account.
- E. Create a service hook in Azure DevOps.

Answer: BDE

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

B: Jenkins' built-in Git Plugin or Team Foundation Server Plugin can poll a Team Services repository every few minutes and queue a job when changes are detected.

D: Use Azure DevOps/ Visual Studio Team Services to create an access token, and use th

E: For those who need tighter integration, Team Services provides two additional ways to achieve it: 1) the Jenkins Service Hook, and 2) Jenkins build and release tasks.)

Reference:

<https://blogs.msdn.microsoft.com/devops/2017/04/25/vsts-visual-studio-team-services-integration-with-jenkins/>

<http://www.aisoftwarellc.com/blog/post/how-to-setup-automated-builds-using-jenkins-and-visual-studio-team-foundation-server/2044>

QUESTION 62

DRAG DROP

Your company has four projects. The version control requirements for each project are shown in the following table.

Project	Requirement
Project 1	Project leads must be able to restrict access to individual files and folders in the repository.
Project 2	The version control system must enforce the following rules on the server before merging any changes to the main branch: <ul style="list-style-type: none"> Changes must be reviewed by at least two project members. Changes must be associated by at least one work item
Project 3	The project members must be able to work in Azure Repos directly from Xcode.
Project 4	The release branch must only be viewable or editable by the project leads.

You plan to use Azure Repos for all the projects.

Which version control system should you use for each project? To answer, drag the appropriate version control systems to the correct projects. Each version control system may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Version Control Systems	Answer Area
<input type="checkbox"/> Git	Project 1: <input type="text" value="0"/>
<input type="checkbox"/> Perforce	Project 2: <input type="text" value="0"/>
<input type="checkbox"/> Subversion	Project 3: <input type="text" value="0"/>
<input type="checkbox"/> Team Foundation Version Control	Project 4: <input type="text" value="0"/>

A. Answer:

Version Control Systems	Answer Area
Git	Project 1: Team Foundation Version Control
Perforce	Project 2: Git
Subversion	Project 3: Subversion
Team Foundation Version Control	Project 4: Git

Explanation

Explanation/Reference:

Explanation:

Box 1: Team Foundation Version Control

TFVC lets you apply granular permissions and restrict access down to a file level.

Box 2: Git

Git is the default version control provider for new projects. You should use Git for version control in your projects unless you have a specific need for centralized version control features in TFVC.

Box 3: Subversion

Note: Xcode is an integrated development environment (IDE) for macOS containing a suite of software development tools developed by Apple

Box 4: Git

Note: Perforce: Due to its multitenant nature, many groups can work on versioned files. The server tracks changes in a central database of MD5 hashes of file content, along with descriptive meta data and separately retains a master repository of file versions that can be verified through the hashes.

References:

<https://searchitoperations.techtarget.com/definition/Perforce-Software>

<https://docs.microsoft.com/en-us/azure/devops/repos/git/share-your-code-in-git-xcode>

<https://docs.microsoft.com/en-us/azure/devops/repos/tfvc/overview>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box 1: Team Foundation Version Control

TFVC lets you apply granular permissions and restrict access down to a file level.

Box 2: Git

Git is the default version control provider for new projects. You should use Git for version control in your projects unless you have a specific need for centralized version control features in TFVC.

Box 3: Subversion

Note: Xcode is an integrated development environment (IDE) for macOS containing a suite of software development tools developed by Apple

Box 4: Git

Note: Perforce: Due to its multitenant nature, many groups can work on versioned files. The server tracks changes in a central database of MD5 hashes of file content, along with descriptive meta data and separately retains a master repository of file versions that can be verified through the hashes.

References:

<https://searchitoperations.techtarget.com/definition/Perforce-Software>

<https://docs.microsoft.com/en-us/azure/devops/repos/git/share-your-code-in-git-xcode>

<https://docs.microsoft.com/en-us/azure/devops/repos/tfvc/overview>

QUESTION 63

You have an Azure Resource Manager template that deploys a multi-tier application.

You need to prevent the user who performs the deployment from viewing the account credentials and connection strings used by the application.

What should you use?

- A. Azure Key Vault
- B. a Web.config file
- C. an Appsettings.json file
- D. an Azure Storage table
- E. an Azure Resource Manager parameter file

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

When you need to pass a secure value (like a password) as a parameter during deployment, you can retrieve the value from an Azure Key Vault. You retrieve the value by referencing the key vault and secret in your parameter file. The value is never exposed

because you only reference its key vault ID. The key vault can exist in a different subscription than the resource group you are deploying to.

References: <https://docs.microsoft.com/en-us/azure/azure-resource-manager/resource-manager-keyvault-parameter>

QUESTION 64

You are automating the build process for a Java-based application by using Azure DevOps.

You need to add code coverage testing and publish the outcomes to the pipeline.

What should you use?

- A. Bullseye Coverage
- B. JUnit
- C. JaCoCo

D. NUnit

Answer: C

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Use Publish Code Coverage Results task in a build pipeline to publish code coverage results to Azure Pipelines or TFS, which were produced by a build in Cobertura or JaCoCo format.

Incorrect Answers:

A: Bullseye Coverage is used for C++ code, and not for Java.

Reference:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/tasks/test/publish-code-coverage-results>

QUESTION 65

SIMULATION

Your company plans to implement a new compliance strategy that will require all Azure web apps to be backed up every five hours.

You need to back up an Azure web app named az400-11566895-main every five hours to an Azure Storage account in your resource group.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

With the storage account ready, you can configure backs up in the web app or App Service.

1. Open the App Service az400-11566895-main, which you want to protect, in the Azure Portal and browse to Settings > Backups. Click Configure and a Backup Configuration blade should appear.
2. Select the storage account.
3. Click + to create a private container. You could name this container after the web app or App Service.
4. Select the container.
5. If you want to schedule backups, then set Scheduled Backup to On and configure a schedule: every five hours
6. Select your retention. Note that 0 means never delete backups.
7. Decide if at least one backup should always be retained.
8. Choose if any connected databases should be included in the web app backup.
9. Click Save to finalize the backup configuration.

The screenshot shows the 'Backup Configuration' blade in the Azure portal. It is divided into three main sections: 'Backup Storage', 'Backup Schedule', and 'Backup Database'.

Backup Storage: This section shows a storage account named 'petri' under 'Storage Settings'. The 'Storage Account' dropdown is set to 'petriaspbackup.blob.core.windows.net'. A note says 'Select the target container to store your app backup.'

Backup Schedule: This section is titled 'Backup Schedule' and includes a clock icon. It allows configuring a scheduled backup. The 'Scheduled backup' switch is set to 'On'. The 'Backup Every' field shows '1 Days'. The 'Start backup schedule from' field is set to '2018-01-20 16:31:38 UTC - Coordinated Universal Time'. The 'Retention (Days)' field is set to '3655'. A note says 'Configure the schedule for your app backup.'

Backup Database: This section shows a database icon and includes a note 'Keep at least one backup' with options 'No' and 'Yes'.

The top right corner of the blade has a close ('X') button.

Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

With the storage account ready, you can configure backs up in the web app or App Service.

1. Open the App Service az400-11566895-main, which you want to protect, in the Azure Portal and browse to Settings > Backups. Click Configure and a Backup Configuration blade should appear.
2. Select the storage account.
3. Click + to create a private container. You could name this container after the web app or App Service.
4. Select the container.
5. If you want to schedule backups, then set Scheduled Backup to On and configure a schedule: every five hours
6. Select your retention. Note that 0 means never delete backups.
7. Decide if at least one backup should always be retained.
8. Choose if any connected databases should be included in the web app backup.
9. Click Save to finalize the backup configuration.

Backup Configuration

Backup Storage

Select the target container to store your app backup.

Storage Settings >
petri

Storage Account: petriaspbackup.blob.core.windows.net

Backup Schedule

Configure the schedule for your app backup.

Scheduled backup On Off

* Backup Every
1 Days Hours

* Start backup schedule from
2018-01-20 16:31:38
UTC - Coordinated Universal Time

* Retention (Days) 3655

Keep at least one backup No Yes

Backup Database

Select the databases you to include with your backup. The backup database list is based on the apps configured connection strings.

INCLUDE IN BACKUP	CONNECTION STRING NAME	DATABASE TYPE
<input checked="" type="checkbox"/> Included	defaultConnection	Sql Database

Reference:
<https://petri.com/backing-azure-app-service>

QUESTION 66

SIMULATION

You need to configure a virtual machine named VM1 to securely access stored secrets in an Azure Key Vault named az400-11566895-kv.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

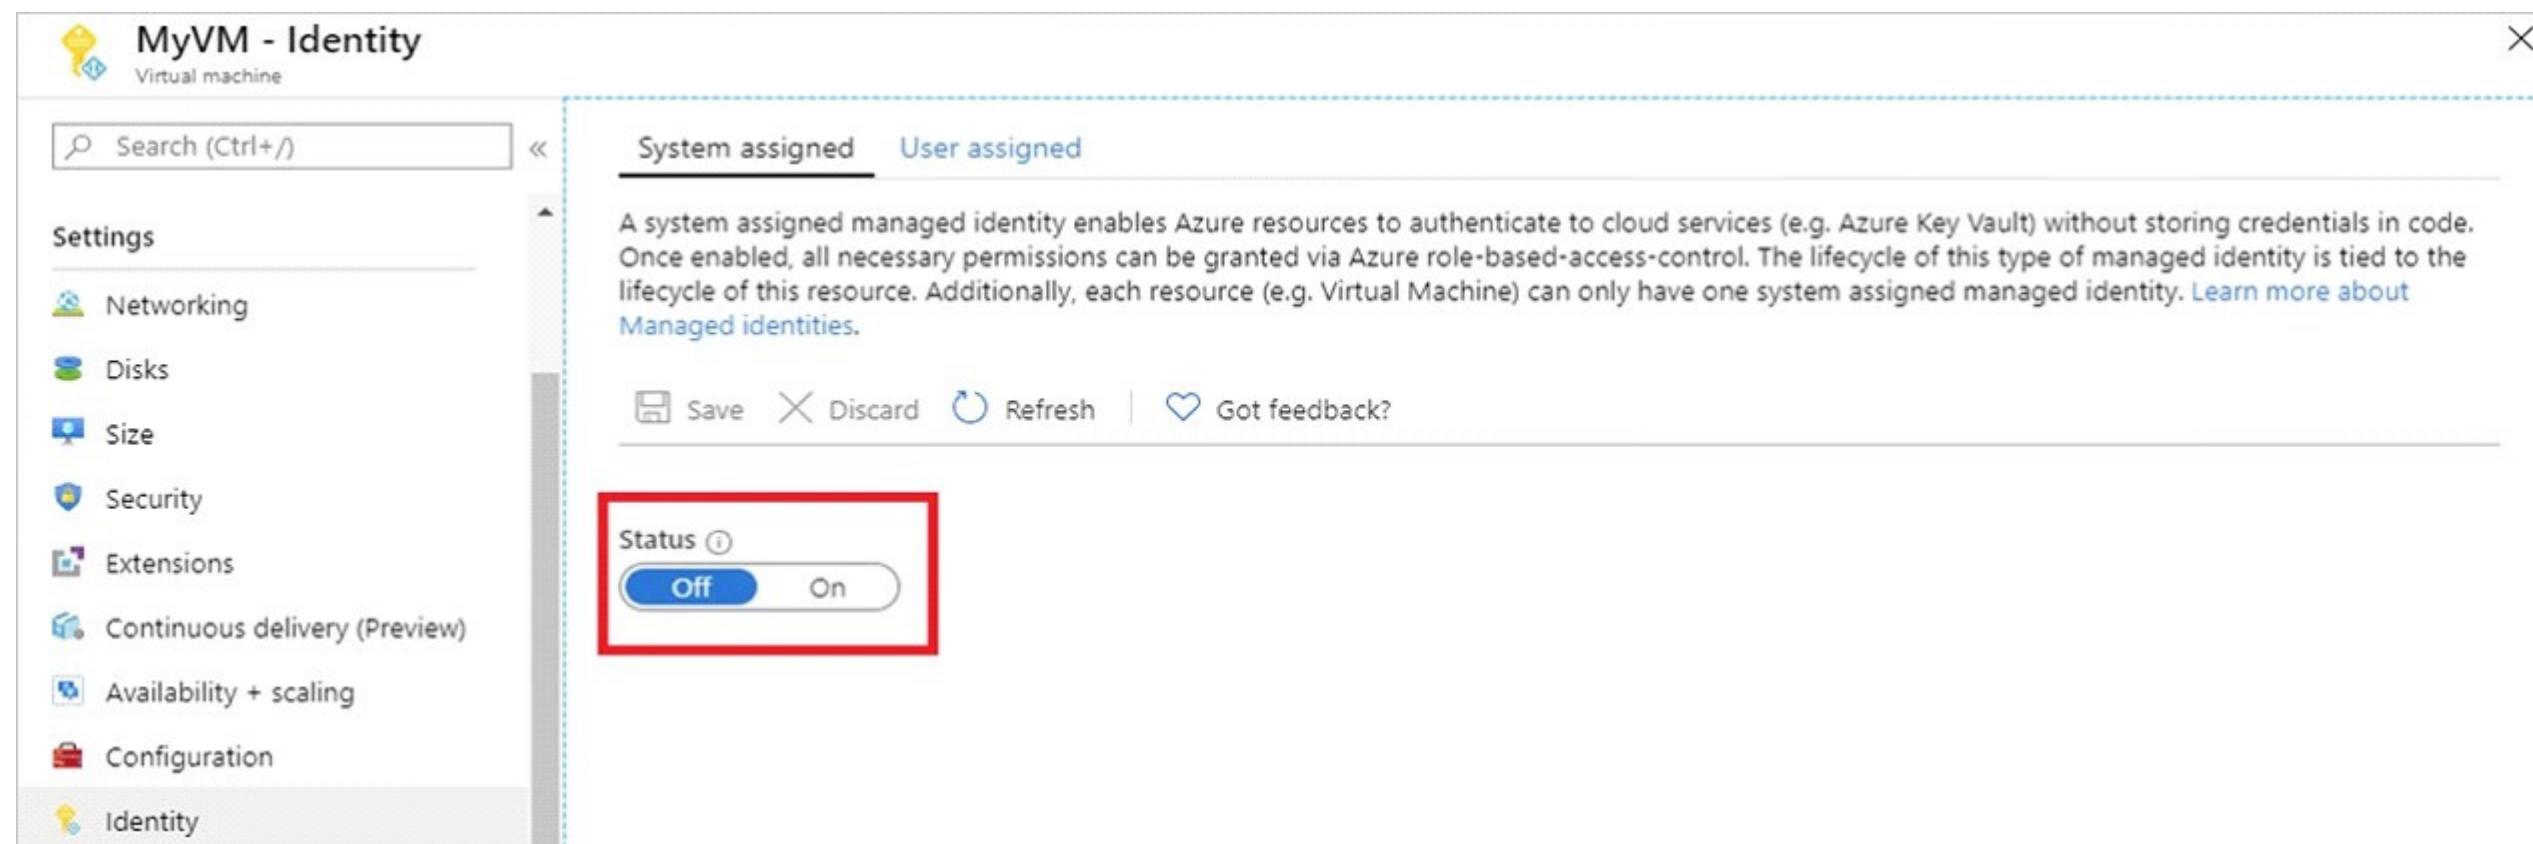
Explanation

Explanation/Reference:

Explanation:

You can use a system-assigned managed identity for a Windows virtual machine (VM) to access Azure Key Vault.

1. Sign in to Azure portal
2. Locate virtual machine VM1.
3. Select Identity
4. Enable the system-assigned identity for VM1 by setting the Status to On.



Note: Enabling a system-assigned managed identity is a one-click experience. You can either enable it during the creation of a VM or in the properties of an existing VM.

Reference:

<https://docs.microsoft.com/en-us/azure/active-directory/managed-identities-azure-resources/tutorial-windows-vm-access-nonaad>

Answer:

Explanation/Reference:

Answer: See solution below.

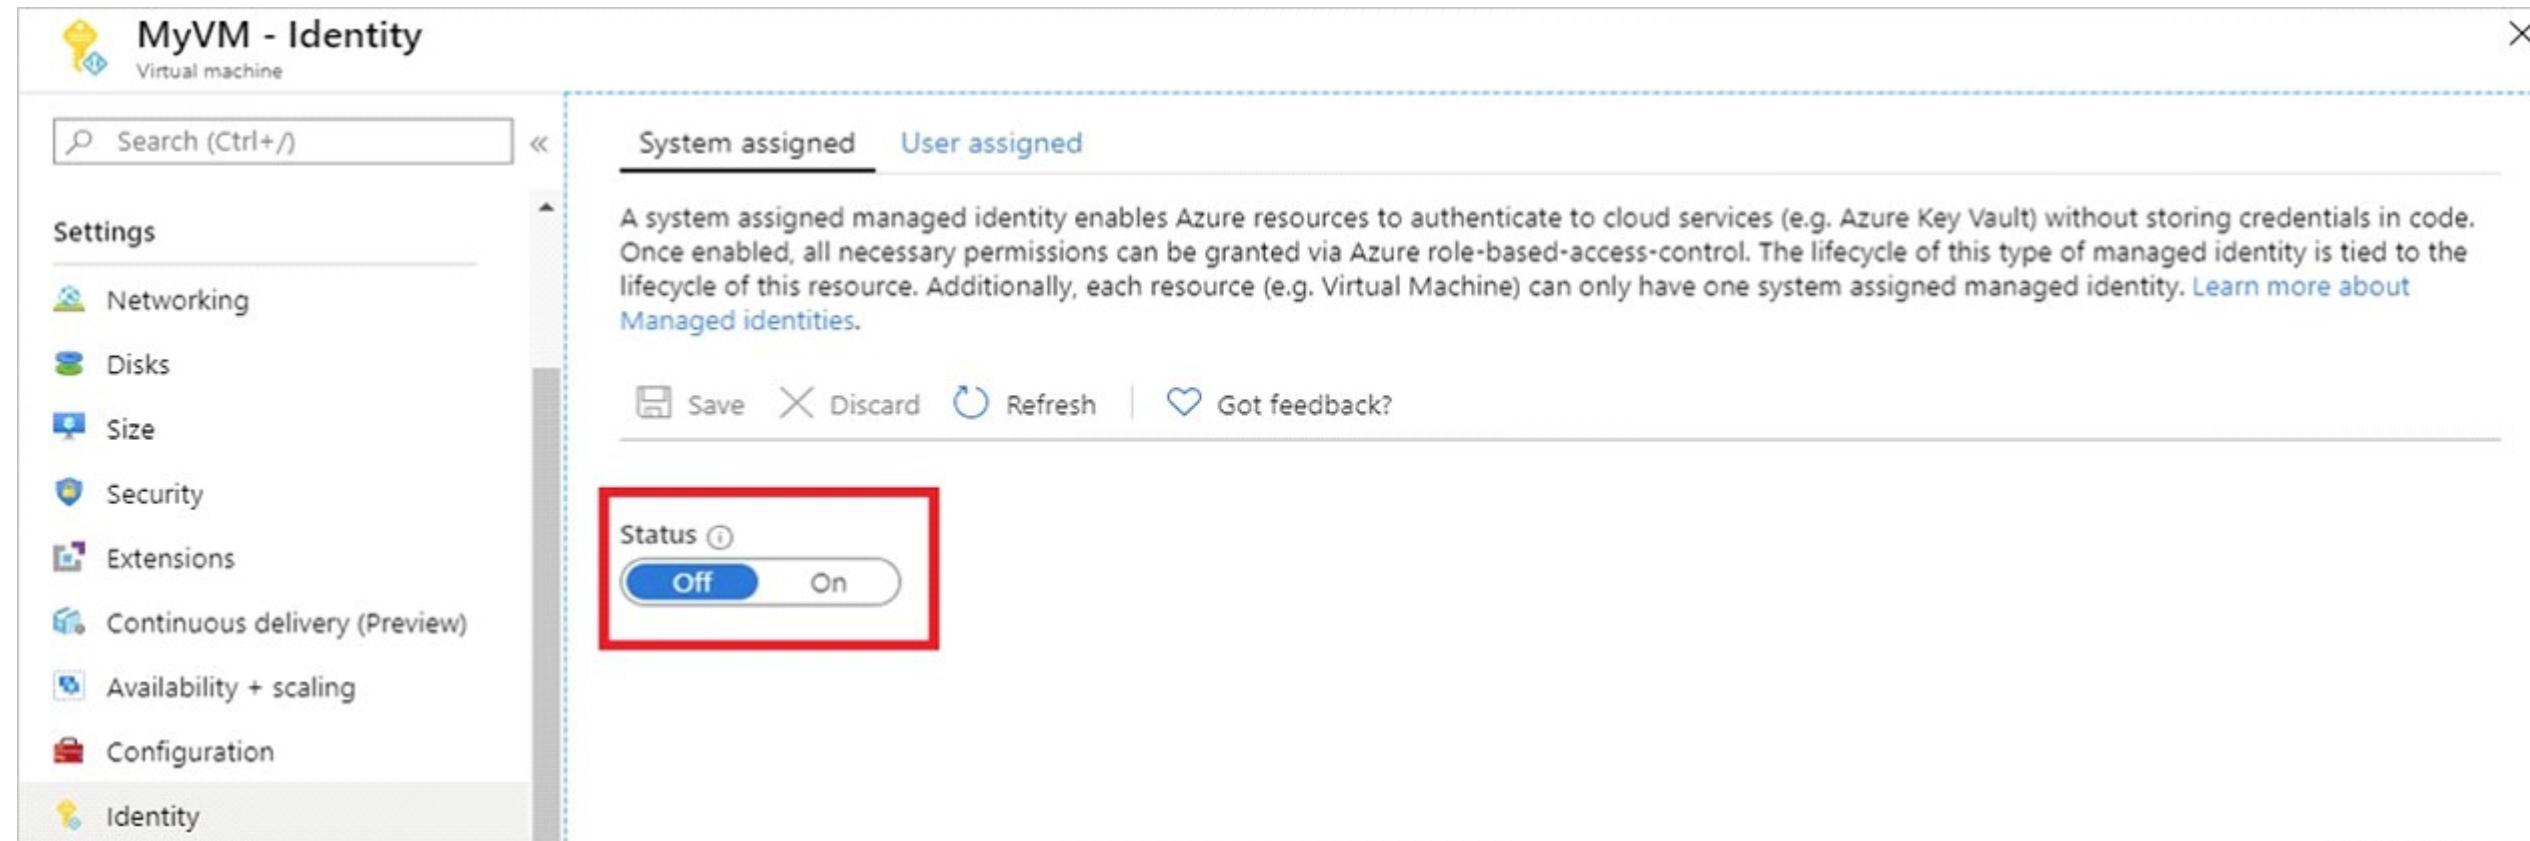
Explanation

Explanation/Reference:

Explanation:

You can use a system-assigned managed identity for a Windows virtual machine (VM) to access Azure Key Vault.

1. Sign in to Azure portal
2. Locate virtual machine VM1.
3. Select Identity
4. Enable the system-assigned identity for VM1 by setting the Status to On.



Note: Enabling a system-assigned managed identity is a one-click experience. You can either enable it during the creation of a VM or in the properties of an existing VM.

Reference:

<https://docs.microsoft.com/en-us/azure/active-directory/managed-identities-azure-resources/tutorial-windows-vm-access-nonaad>

QUESTION 67

SIMULATION

You need to ensure that Microsoft Visual Studio 2017 can remotely attach to an Azure Function named fa-11566895.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

Enable Remote Debugging

Before we start a debugging session to our Azure Function app we need to enable the functionality.

1. Navigate in the Azure portal to your function app fa-11566895
2. Go to the "Application settings"
3. Under "Debugging" set Remote Debugging to On and set Remote Visual Studio version to 2017.

Reference:

<https://www.locktar.nl/uncategorized/azure-remote-debugging-manually-in-visual-studio-2017/>

Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

Enable Remote Debugging

Before we start a debugging session to our Azure Function app we need to enable the functionality.

1. Navigate in the Azure portal to your function app fa-11566895
2. Go to the "Application settings"
3. Under "Debugging" set Remote Debugging to On and set Remote Visual Studio version to 2017.

Reference:

<https://www.locktar.nl/uncategorized/azure-remote-debugging-manually-in-visual-studio-2017/>

QUESTION 68

HOTSPOT

You company uses Azure DevOps to deploy infrastructures to Azure.

Pipelines are developed by using YAML.

You execute a pipeline and receive the results in the web portal for Azure Pipelines as shown in the following exhibit.

Azure DevOps

Fast Track

Overview

Boards

Repos

Pipelines

Pipelines

Environments

Releases

Library

Task groups

Deployment groups

WhiteSource Bolt

Test Plans

Artifacts

← Jobs in run #20191120.1

Fast Track

build vm

initial_build

1 Pool: Azure Pipelines

2 Image: Ubuntu-18.04

3 Agent: Hosted Agent

4 Started: Just now

5 Duration: 7s

6

7 ▶ Job preparation parameters

initialize build

Initialize job <1s

Checkout 4s

CmdLine 2s

Post-job: Ccheckout <1s

Finalize Job <1s

deploy_to_dev

> deploy_to_dev_server 2s

deploy_to_uat

> deploy_to_uat_server 2s

Finalize build

Report build status <1s

The screenshot shows the Azure DevOps interface with the 'Pipelines' section selected. A specific pipeline named 'Fast Track' is running, with the ID #20191120.1. The 'initial_build' job is expanded, showing its tasks: Initialize job, Checkout, CmdLine, Post-job: Ccheckout, and Finalize Job. The duration of this job is 7 seconds. Below the initial build, there are two deployment steps: 'deploy_to_dev' and 'deploy_to_uat', each consisting of a single task ('deploy_to_dev_server' and 'deploy_to_uat_server' respectively) taking 2 seconds. Finally, the 'Finalize build' step contains a single task ('Report build status') which completed very quickly (<1 second). The entire run was started just now.

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

The pipeline contains

- one stage
- two stages
- three stages
- four stages
- five stages

Build_vm contains

- one job
- two jobs
- three jobs
- four jobs
- five jobs

A. Answer:

Answer Area

The pipeline contains

one stage	▼
two stages	
three stages	
four stages	
five stages	

Build_vm contains

one job	▼
two jobs	
three jobs	
four jobs	
five jobs	

Explanation

Explanation/Reference:

Reference:

<https://dev.to/rajikaimal/azure-devops-ci-cd-yaml-pipeline-4glj>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Reference:

<https://dev.to/rajikaimal/azure-devops-ci-cd-yaml-pipeline-4glj>

QUESTION 69

DRAG DROP

Your company has an Azure subscription named Subscription1. Subscription1 is associated to an Azure Active Directory tenant named contoso.com.

You need to provision an Azure Kubernetes Services (AKS) cluster in Subscription1 and set the permissions for the cluster by using RBAC roles that reference the identities in contoso.com.

Which three objects should you create in sequence? To answer, move the appropriate objects from the list of objects to the answer area and arrange them in the correct order.

Select and Place:

Answer Area

Objects

a system-assigned managed identity

a cluster

an application registration in contoso.com

an RBAC binding

A. Answer:

Answer Area	
Objects	
a system-assigned managed identity	
a cluster	a cluster
an application registration in contoso.com	a system-assigned managed identity
an RBAC binding	an RBAC binding

Explanation

Explanation/Reference:

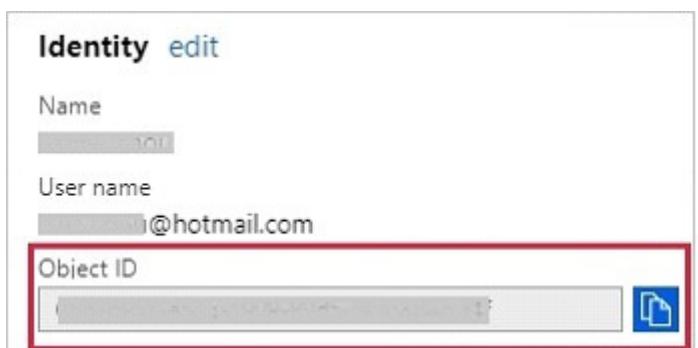
Explanation:

Step 1: Create an AKS cluster

Step 2: a system-assigned managed identity

To create an RBAC binding, you first need to get the Azure AD Object ID.

1. Sign in to the Azure portal.
2. In the search field at the top of the page, enter Azure Active Directory.
3. Click Enter.
4. In the Manage menu, select Users.
5. In the name field, search for your account.
6. In the Name column, select the link to your account.
7. In the Identity section, copy the Object ID.



Step 3: a RBAC binding

Reference:

<https://docs.microsoft.com/en-us/azure/developer/ansible/aks-configure-rbac>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

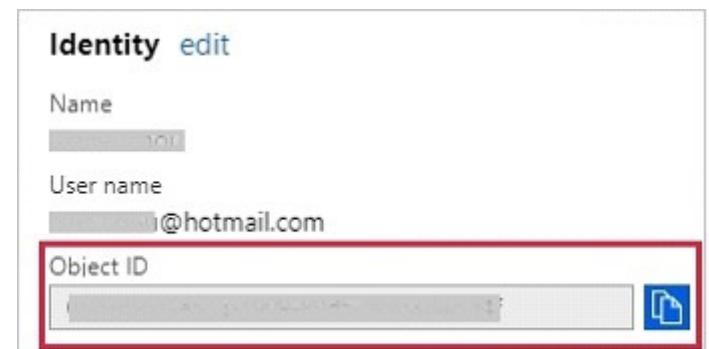
Explanation

Step 1: Create an AKS cluster

Step 2: a system-assigned managed identity

To create an RBAC binding, you first need to get the Azure AD Object ID.

1. Sign in to the Azure portal.
2. In the search field at the top of the page, enter Azure Active Directory.
3. Click Enter.
4. In the Manage menu, select Users.
5. In the name field, search for your account.
6. In the Name column, select the link to your account.
7. In the Identity section, copy the Object ID.



Step 3: a RBAC binding

Reference:

<https://docs.microsoft.com/en-us/azure/developer/ansible/aks-configure-rbac>

QUESTION 70

HOTSPOT

You manage build and release pipelines by using Azure DevOps. Your entire managed environment resides in Azure.

You need to configure a service endpoint for accessing Azure Key Vault secrets. The solution must meet the following requirements:

Ensure that the secrets are retrieved by Azure DevOps.

Avoid persisting credentials and tokens in Azure DevOps.

How should you configure the service endpoint? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Service connection type:

- Azure Resource Manager
- Generic service
- Team Foundation Server / Azure Pipelines service connection

Authentication/authorization method for the connection:

- Azure Active Directory OAuth 2.0
- Grant authorization
- Managed Service Identity Authentication

A. Answer:

Answer Area				
Service connection type:	<table border="1"><tr><td>Azure Resource Manager</td></tr><tr><td>Generic service</td></tr><tr><td>Team Foundation Server / Azure Pipelines service connection</td></tr></table>	Azure Resource Manager	Generic service	Team Foundation Server / Azure Pipelines service connection
Azure Resource Manager				
Generic service				
Team Foundation Server / Azure Pipelines service connection				
Authentication/authorization method for the connection:	<table border="1"><tr><td>Azure Active Directory OAuth 2.0</td></tr><tr><td>Grant authorization</td></tr><tr><td>Managed Service Identity Authentication</td></tr></table>	Azure Active Directory OAuth 2.0	Grant authorization	Managed Service Identity Authentication
Azure Active Directory OAuth 2.0				
Grant authorization				
Managed Service Identity Authentication				

Explanation

Explanation/Reference:

Explanation:

Box 1: Azure Pipelines service connection

Box 2: Managed Service Identity Authentication

The managed identities for Azure resources feature in Azure Active Directory (Azure AD) provides Azure services with an automatically managed identity in Azure AD. You can use the identity to authenticate to any service that supports Azure AD authentication, including Key Vault, without any credentials in your code.

Reference:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/tasks/deploy/azure-key-vault>

<https://docs.microsoft.com/en-us/azure/active-directory/managed-identities-azure-resources/overview>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box 1: Azure Pipelines service connection

Box 2: Managed Service Identity Authentication

The managed identities for Azure resources feature in Azure Active Directory (Azure AD) provides Azure services with an automatically managed identity in Azure AD. You can use the identity to authenticate to any service that supports Azure AD authentication, including Key Vault, without any credentials in your code.

Reference:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/tasks/deploy/azure-key-vault>

<https://docs.microsoft.com/en-us/azure/active-directory/managed-identities-azure-resources/overview>

QUESTION 71

You are deploying a server application that will run on a Server Core installation of Windows Server 2019.

You create an Azure key vault and a secret.

You need to use the key vault to secure API secrets for third-party integrations.

Which three actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Configure RBAC for the key vault.
- B. Modify the application to access the key vault.
- C. Configure a Key Vault access policy.
- D. Deploy an Azure Desired State Configuration (DSC) extension.
- E. Deploy a virtual machine that uses a system-assigned managed identity.

Answer: BCE

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

BE: An app deployed to Azure can take advantage of Managed identities for Azure resources, which allows the app to authenticate with Azure Key Vault using Azure AD authentication without credentials (Application ID and Password/Client Secret) stored in the app.

C:

1. Select Add Access Policy.
2. Open Secret permissions and provide the app with Get and List permissions.
3. Select Select principal and select the registered app by name. Select the Select button.
4. Select OK.
5. Select Save.
6. Deploy the app.

References:

<https://docs.microsoft.com/en-us/aspnet/core/security/key-vault-configuration>

QUESTION 72

Testlet 1

Case Study

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other question in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview

Contoso, Ltd. is a manufacturing company that has a main office in Chicago.

Existing Environment

Contoso plans to improve its IT development and operations processes by implementing Azure DevOps principles. Contoso has an Azure subscription and creates an Azure DevOps organization.

The Azure DevOps organization includes:

The Docker extension

A deployment pool named Pool7 that contains 10 Azure virtual machines that run Windows Server 2016

The Azure subscription contains an Azure Automation account.

Requirements

Planned changes

Contoso plans to create projects in Azure DevOps as shown in the following table.

Project name	Project details
Project 1	Project1 will provide support for incremental builds and third-party SDK components
Project 2	Project2 will use an automatic build policy. A small team of developers named Team2 will work independently on changes to the project. The Team2 members will not have permissions to Project2.
Project 3	Project3 will be integrated with SonarQube
Project 4	Project4 will provide support for a build pipeline that creates a Docker image and pushes the image to the Azure Container Registry. Project4 will use an existing Dockerfile.
Project 5	Project5 will contain a Git repository in Azure Repos and a continuous integration trigger that will initiate a build in response to any change except for changes within /folder1 of the repository.
Project 6	Project6 will provide support for build and deployment pipelines. Deployment will be allowed only if the number of current work items representing active software bugs is 0.
Project 7	Project7 will contain a target deployment group named Group7 that maps to Pool7. Project7 will use Azure Automation State Configuration to maintain the desired state of the computers in Group7.

Technical requirements

Contoso identifies the following technical requirements:

Implement build agents for Project1.

Whenever possible, use Azure resources.

Avoid using deprecated technologies.

Implement a code flow strategy for Project2 that will:

- Enable Team2 to submit pull requests for Project2.
 - Enable Team2 to work independently on changes to a copy of Project2.
 - Ensure that any intermediary changes performed by Team2 on a copy of Project2 will be subject to the same restrictions as the ones defined in the build policy of Project2.
- Whenever possible, implement automation and minimize administrative effort.

Implement Project3, Project5, Project6, and Project7 based on the planned changes.

Implement Project4 and configure the project to push Docker images to Azure Container Registry.

HOTSPOT

How should you configure the filters for the Project5 trigger? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Set a  /folder1.

- branch filter to exclude
- branch filter to include
- path filter to exclude
- path filter to include

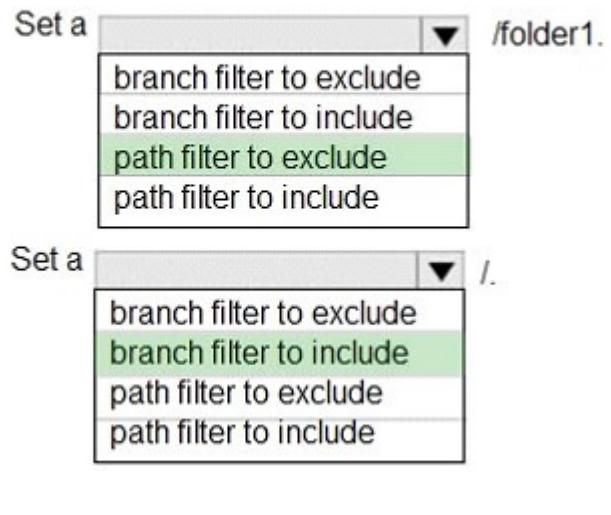
Set a  /.

- branch filter to exclude
- branch filter to include
- path filter to exclude
- path filter to include

@

A. Answer:

Answer Area



@

Explanation

Explanation/Reference:

Explanation:

Scenario:

Project5 will contain a Git repository in Azure Reports and a continuous integration trigger that will initiate a build in response to any change except for changes within /folder1 of the repository.

Continuous integration (CI) triggers cause a build to run whenever a push is made to the specified branches or a specified tag is pushed.

Box 2: branch filter to include

You can specify branches to include and exclude. For example:

```
# specific branch build  
trigger:  
branches:  
include:  
- master  
- releases/*  
exclude:  
- releases/old*
```

References:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/build/triggers>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Scenario:

Project5 will contain a Git repository in Azure Reports and a continuous integration trigger that will initiate a build in response to any change except for changes within /folder1 of the repository.

Continuous integration (CI) triggers cause a build to run whenever a push is made to the specified branches or a specified tag is pushed.

Box 2: branch filter to include

You can specify branches to include and exclude. For example:

```
# specific branch build  
trigger:  
branches:  
include:  
- master  
- releases/*  
exclude:  
- releases/old*
```

References:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/build/triggers>

QUESTION 73

In Azure DevOps, you create Project3.

You need to meet the requirements of the project.

What should you do first?

- A. From Azure DevOps, modify the build definition.
- B. From SonarQube, obtain an authentication token.
- C. From Azure DevOps, create a service endpoint.
- D. From SonarQube, create a project.

Answer: C

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

The first thing to do is to declare your SonarQube server as a service endpoint in your VSTS/DevOps project settings.

References: <https://docs.sonarqube.org/display/SCAN/Analyzing+with+SonarQube+Extension+for+vsts-TFS>

QUESTION 74

You need to implement Project4.

What should you do first?

- A. Add the FROM instruction in the Dockerfile file.
- B. Add a Copy and Publish Build Artifacts task to the build pipeline.
- C. Add a Docker task to the build pipeline.
- D. Add the MAINTAINER instruction in the Dockerfile file.

Answer: C

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Scenario: Implement Project4 and configure the project to push Docker images to Azure Container Registry.

Project 4	Project4 will provide support for a build pipeline that creates a Docker image and pushes the image to the Azure Container Registry. Project4 will use an existing Dockerfile.
-----------	--

You use Azure Container Registry Tasks commands to quickly build, push, and run a Docker container image natively within Azure, showing how to offload your "inner-loop" development cycle to the cloud. ACR Tasks is a suite of features within Azure Container Registry to help you manage and modify container images across the container lifecycle.

References:

<https://docs.microsoft.com/en-us/azure/container-registry/container-registry-quickstart-task-cli>

QUESTION 75

DRAG DROP

You need to recommend a procedure to implement the build agent for Project1.

Which three actions should you recommend be performed in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions

Sign in to Azure DevOps by using an account that is assigned the Administrator service connection security role.

Install the Azure Pipelines agent on on-premises virtual machine.

Create a personal access token in the Azure DevOps organization of Contoso.

Install and register the Azure Pipelines agent on an Azure virtual machine.

Sign in to Azure DevOps by using an account that is assigned the agent pool administrator role.

Answer Area

A. Answer:

Actions	Answer Area
	Sign in to Azure DevOps by using an account that is assigned the Administrator service connection security role.
Install the Azure Pipelines agent on on-premises virtual machine.	Create a personal access token in the Azure DevOps organization of Contoso.
	Install and register the Azure Pipelines agent on an Azure virtual machine.
Sign in to Azure DevOps by using an account that is assigned the agent pool administrator role.	

Explanation

Explanation/Reference:

Explanation:

Scenario:

Project 1	Project1 will provide support for incremental builds and third-party SDK components
-----------	---

Step 1: Sign in to Azure Devops by using an account that is assigned the Administrator service connection security role.

Note: Under Agent Phase, click Deploy Service Fabric Application. Click Docker Settings and then click Configure Docker settings. In Registry Credentials Source, select Azure Resource Manager Service Connection. Then select your Azure subscription.

Step 2: Create a personal access token..

A personal access token or PAT is required so that a machine can join the pool created with the Agent Pools (read, manage) scope.

Step 3: Install and register the Azure Pipelines agent on an Azure virtual machine.

By running a Azure Pipeline agent in the cluster, we make it possible to test any service, regardless of type.

References:

<https://docs.microsoft.com/en-us/azure/service-fabric/service-fabric-tutorial-deploy-container-app-with-cicd-vsts>

<https://mohitgoyal.co/2019/01/10/run-azure-devops-private-agents-in-kubernetes-clusters/>

Question Set 2

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation

Scenario:

Project 1	Project1 will provide support for incremental builds and third-party SDK components
-----------	---

Step 1: Sign in to Azure Devops by using an account that is assigned the Administrator service connection security role.

Note: Under Agent Phase, click Deploy Service Fabric Application. Click Docker Settings and then click Configure Docker settings. In Registry Credentials Source, select Azure Resource Manager Service Connection. Then select your Azure subscription.

Step 2: Create a personal access token..

A personal access token or PAT is required so that a machine can join the pool created with the Agent Pools (read, manage) scope.

Step 3: Install and register the Azure Pipelines agent on an Azure virtual machine.

By running a Azure Pipeline agent in the cluster, we make it possible to test any service, regardless of type.

References:

<https://docs.microsoft.com/en-us/azure/service-fabric/service-fabric-tutorial-deploy-container-app-with-cicd-vsts>

<https://mohitgoyal.co/2019/01/10/run-azure-devops-private-agents-in-kubernetes-clusters/>

Question Set 2

QUESTION 76

You have 50 Node.js-based projects that you scan by using WhiteSource. Each project includes Package.json, Package-lock.json, and Npm-shrinkwrap.json files.

You need to minimize the number of libraries reports by WhiteSource to only the libraries that you explicitly reference.

What should you do?

- A. Configure the File System Agent plug-in.
- B. Add a devDependencies section to Package-lock.json.
- C. Configure the Artifactory plug-in.
- D. Delete Package-lock.json.

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Separate Your Dependencies

Within your package.json file be sure you split out your npm dependencies between devDependencies and (production) dependencies. The key part is that you must then make use of the --production flag when installing the npm packages. The --production flag

will exclude all packages defined in the devDependencies section.

References: <https://blogs.msdn.microsoft.com/visualstudioalmrangers/2017/06/08/manage-your-open-source-usage-and-security-as-reported-by-your-cicd-pipeline/>

QUESTION 77

Your company deploys applications in Docker containers.

You want to detect known exploits in the Docker images used to provision the Docker containers.

You need to integrate image scanning into the application lifecycle. The solution must expose the exploits as early as possible during the application lifecycle.

What should you configure?

- A. a task executed in the continuous integration pipeline and a scheduled task that analyzes the image registry
- B. manual tasks performed during the planning phase and the deployment phase
- C. a task executed in the continuous deployment pipeline and a scheduled task against a running production container
- D. a task executed in the continuous integration pipeline and a scheduled task that analyzes the production container

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

You can use the Docker task to sign into ACR and then use a subsequent script to pull an image and scan the container image for vulnerabilities.
Use the docker task in a build or release pipeline. This task can be used with Docker or Azure Container registry.

Incorrect Answers:

C: We should not wait until deployment. We want to detect the exploits as early as possible.

D: We should wait until the image is in the product container. We want to detect the exploits as early as possible.

References: <https://docs.microsoft.com/en-us/azure/devops/articles/security-validation-cicd-pipeline?view=vsts>

QUESTION 78

Your company uses Azure DevOps for the build pipelines and deployment pipelines of Java-based projects.

You need to recommend a strategy for managing technical debt.

Which two actions should you include in the recommendation? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Configure post-deployment approvals in the deployment pipeline.
- B. Configure pre-deployment approvals in the deployment pipeline.
- C. Integrate Azure DevOps and SonarQube.
- D. Integrate Azure DevOps and Azure DevTest Labs.

Answer: BC

Explanation/Reference:

Explanation

Explanation/Reference:

QUESTION 79

Your company has a hybrid cloud between Azure and Azure Stack.

The company uses Azure DevOps for its full CI/CD pipelines. Some applications are built by using Erlang and Hack.

You need to ensure that Erlang and Hack are supported as part of the build strategy across the hybrid cloud. The solution must minimize management overhead.

What should you use to execute the build pipeline?

- A. a Microsoft-hosted agent
- B. Azure DevOps self-hosted agents on Azure DevTest Labs virtual machines.
- C. Azure DevOps self-hosted agents on Hyper-V virtual machines
- D. Azure DevOps self-hosted agents on virtual machines that run on Azure Stack

Answer: D

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Azure Stack offers virtual machines (VMs) as one type of an on-demand, scalable computing resource. You can choose a VM when you need more control over the computing environment.

References: <https://docs.microsoft.com/en-us/azure/azure-stack/user/azure-stack-compute-overview>

QUESTION 80

You need to recommend a Docker container build strategy that meets the following requirements:

Minimizes image sizes

Minimizes the security surface area of the final image

What should you include in the recommendation?

- A. multi-stage builds
- B. PowerShell Desired State Configuration (DSC)
- C. Docker Swarm
- D. single-stage builds

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Multi-stage builds are a new feature requiring Docker 17.05 or higher on the daemon and client. Multistage builds are useful to anyone who has struggled to optimize Dockerfiles while keeping them easy to read and maintain.

Incorrect Answers:

C: A swarm consists of multiple Docker hosts which run in swarm mode and act as managers (to manage membership and delegation) and workers (which run swarm services).

References: <https://docs.docker.com/develop/develop-images/multistage-build/>

QUESTION 81

You plan to create an image that will contain a .NET Core application.

You have a Dockerfile file that contains the following code. (Line numbers are included for reference only.)

```
01 FROM microsoft/dotnet:2.1-sdk
02 COPY ./
03 RUN dotnet publish -c Release -o out
04 FROM microsoft/dotnet:2.1-sdk
05 COPY --from=0 /out /
06 WORKDIR /
07 ENTRYPOINT ["dotnet", "app1.dll"]
```

You need to ensure that the image is as small as possible when the image is built.

Which line should you modify in the file?

- A. 1
- B. 3
- C. 4
- D. 7

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Multi-stage builds (in Docker 17.05 or higher) allow you to drastically reduce the size of your final image, without struggling to reduce the number of intermediate layers and files.

With multi-stage builds, you use multiple FROM statements in your Dockerfile. Each FROM instruction can use a different base, and each of them begins a new stage of the build. You can selectively copy artifacts from one stage to another, leaving behind everything you don't want in the final image.

References: <https://docs.docker.com/develop/develop-images/multistage-build/#use-multi-stage-builds>

QUESTION 82

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

Your company has a project in Azure DevOps for a new web application.

You need to ensure that when code is checked in, a build runs automatically.

Solution: From the Triggers tab of the build pipeline, you select Batch changes while a build is in progress.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Instead, In Visual Designer you enable continuous integration (CI) by:

1. Select the Triggers tab.
2. Enable Continuous integration.

Note: Batch changes

Select this check box if you have many team members uploading changes often and you want to reduce the number of builds you are running. If you select this option, when a build is running, the system waits until the build is completed and then queues another build of all changes that have not yet been built.

References:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/get-started-designer>

QUESTION 83

HOTSPOT

You need to deploy Azure Kubernetes Service (AKS) to host an application. The solution must meet the following requirements:

Containers must only be published internally.

AKS clusters must be able to create and manage containers in Azure.

What should you use for each requirement? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Containers must only be published internally:

	▼
Azure Container Instances	
Azure Container Registry	
Dockerfile	

AKS clusters must be able to create and manage containers in Azure:

	▼
An Azure Active Directory (Azure AD) group	
An Azure Automation account	
An Azure service principal	

A. Answer:

Answer Area

Containers must only be published internally:

Azure Container Instances
Azure Container Registry
Dockerfile

AKS clusters must be able to create and manage containers in Azure:

An Azure Active Directory (Azure AD) group
An Azure Automation account
An Azure service principal

Explanation

Explanation/Reference:

Explanation:

Box 1: Azure Container Registry

Azure services like Azure Container Registry (ACR) and Azure Container Instances (ACI) can be used and connected from independent container orchestrators like kubernetes (k8s). You can set up a custom ACR and connect it to an existing k8s cluster to ensure images will be pulled from the private container registry instead of the public docker hub.

Box 2: An Azure service principal

When you're using Azure Container Registry (ACR) with Azure Kubernetes Service (AKS), an authentication mechanism needs to be established. You can set up AKS and ACR integration during the initial creation of your AKS cluster. To allow an AKS cluster to interact with ACR, an Azure Active Directory service principal is used.

References:

<https://thorsten-hans.com/how-to-use-private-azure-container-registry-with-kubernetes>

<https://docs.microsoft.com/en-us/azure/aks/cluster-container-registry-integration>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box 1: Azure Container Registry

Azure services like Azure Container Registry (ACR) and Azure Container Instances (ACI) can be used and connected from independent container orchestrators like kubernetes (k8s). You can set up a custom ACR and connect it to an existing k8s cluster to ensure images will be pulled from the private container registry instead of the public docker hub.

Box 2: An Azure service principal

When you're using Azure Container Registry (ACR) with Azure Kubernetes Service (AKS), an authentication mechanism needs to be established. You can set up AKS and ACR integration during the initial creation of your AKS cluster. To allow an AKS cluster to interact with ACR, an Azure Active Directory service principal is used.

References:
<https://thorsten-hans.com/how-to-use-private-azure-container-registry-with-kubernetes>

<https://docs.microsoft.com/en-us/azure/aks/cluster-container-registry-integration>

QUESTION 84

You are designing an Azure DevOps strategy for your company's development team.

You suspect that the team's productivity is low due to accumulate technical debt.

You need to recommend a metric to assess the amount of the team's technical debt.

What should you recommend?

- A. the number of code modules in an application
- B. the number of unit test failures
- C. the percentage of unit test failures
- D. the percentage of overall time spent on rework

Answer: D

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Technical Debt is the estimated cost to fix code elements issues.

Technical Debt Ratio: Ratio between the cost to develop the software and the cost to fix it. The Technical Debt Ratio formula is:

Remediation cost / Development cost

Which can be restated as:

Remediation cost / (Cost to develop 1 line of code * Number of lines of code)

References:

<http://www.azure365.co.in/devops/3PDevOps-4>

QUESTION 85

You are developing an open source solution that uses a GitHub repository.

You create a new public project in Azure DevOps.

You plan to use Azure Pipelines for continuous build. The solution will use the GitHub Checks API.

Which authentication type should you use?

- A. OpenID
- B. GitHub App
- C. a personal access token (PAT)
- D. SAML

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

You can authenticate as a GitHub App.

Reference:

<https://developer.github.com/apps/building-github-apps/authenticating-with-github-apps/>

QUESTION 86

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

Your company has a project in Azure DevOps for a new web application.

You need to ensure that when code is checked in, a build runs automatically.

Solution: From the Continuous deployment trigger settings of the release pipeline, you enable the Pull request trigger setting.

Does this meet the goal?

A. Yes

B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

In Visual Designer you enable continuous integration (CI) by:

1. Select the Triggers tab.
2. Enable Continuous integration.

References:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/get-started-designer>

QUESTION 87

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

Your company has a project in Azure DevOps for a new web application.

You need to ensure that when code is checked in, a build runs automatically.

Solution: From the Pre-deployment conditions settings of the release pipeline, you select After stage.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Instead, In Visual Designer you enable continuous integration (CI) by:

1. Select the Triggers tab.
2. Enable Continuous integration.

References:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/get-started-designer>

QUESTION 88

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

Your company has a project in Azure DevOps for a new web application.

You need to ensure that when code is checked in, a build runs automatically.

Solution: From the Pre-deployment conditions settings of the release pipeline, you select Batch changes while a build is in progress.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Instead, In Visual Designer you enable continuous integration (CI) by:

1. Select the Triggers tab.
2. Enable Continuous integration.

References:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/get-started-designer>

QUESTION 89

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

Your company has a project in Azure DevOps for a new web application.

You need to ensure that when code is checked in, a build runs automatically.

Solution: From the Triggers tab of the build pipeline, you select Enable continuous integration.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

In Visual Designer you enable continuous integration (CI) by:

1. Select the Triggers tab.
2. Enable Continuous integration.

A continuous integration trigger on a build pipeline indicates that the system should automatically queue a new build whenever a code change is committed.

References:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/get-started-designer>

QUESTION 90

You have a project in Azure DevOps. You have an Azure Resource Group deployment project in Microsoft Visual Studio that is checked in to the Azure DevOps project.

You need to create a release pipeline that will deploy resources by using Azure Resource Manager templates. The solution must minimize administrative effort.

Which task type should you include in the solution?

- A. Azure Cloud Service Deployment
- B. Azure RM Web App Deployment
- C. Azure PowerShell
- D. Azure App Service Manage

Answer: C

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

There are two different ways to deploy templates to Azure DevOps Services. Both methods provide the same results, so choose the one that best fits your workflow.

1. Add a single step to your build pipeline that runs the PowerShell script that's included in the Azure Resource Group deployment project (Deploy-AzureResourceGroup.ps1). The script copies artifacts and then deploys the template.
2. Add multiple Azure DevOps Services build steps, each one performing a stage task.

The first option has the advantage of using the same script used by developers in Visual Studio and providing consistency throughout the lifecycle.

References:

<https://docs.microsoft.com/en-us/azure/vs-azure-tools-resource-groups-ci-in-vsts>

QUESTION 91

Your company is building a new solution in Java.

The company currently uses a SonarQube server to analyze the code of .NET solutions.

You need to analyze and monitor the code quality of the Java solution.

Which task types should you add to the build pipeline?

- A. Chef
- B. Gradle
- C. Octopus
- D. xCODE

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

SonarQube is a set of static analyzers that can be used to identify areas of improvement in your code. It allows you to analyze the technical debt in your project and keep track of it in the future. With Maven and Gradle build tasks, you can run SonarQube analysis with minimal setup in a new or existing Azure DevOps Services build task.

Reference:

<https://docs.microsoft.com/en-us/azure/devops/java/sonarqube?view=azure-devops>

QUESTION 92

DRAG DROP

You are developing a full Microsoft .NET Framework solution that includes unit tests.

You need to configure SonarQube to perform a code quality validation of the C# code as part of the build pipelines.

Which four tasks should you perform in sequence? To answer, move the appropriate tasks from the list of tasks to the answer area and arrange them in the correct order.

Select and Place:

Actions Commands Cmdlets Statements

Run Code Analysis

Visual Studio Test

Publish Build Artifacts

Visual Studio Build

Prepare Analysis Configuration

Answer Area

A. Answer:

Actions	Commands	Cmdlets	Statements	Answer Area
				Prepare Analysis Configuration
				Visual Studio Build
Publish Build Artifacts				Visual Studio Test
				Run Code Analysis

Explanation

Explanation/Reference:

Explanation:

Step 1: Prepare Analysis Configuration

Prepare Analysis Configuration task, to configure all the required settings before executing the build.

This task is mandatory.

In case of .NET solutions or Java projects, it helps to integrate seamlessly with MSBuild, Maven and Gradle tasks.

Step 2: Visual Studio Build

Reorder the tasks to respect the following order:

Prepare Analysis Configuration task before any MSBuild or Visual Studio Build task.

Step 3: Visual Studio Test

Reorder the tasks to respect the following order:

Run Code Analysis task after the Visual Studio Test task.

Step 4: Run Code Analysis

Run Code Analysis task, to actually execute the analysis of the source code.

This task is not required for Maven or Gradle projects, because scanner will be run as part of the Maven/Gradle build.

Note:



NuGet restore

NuGet



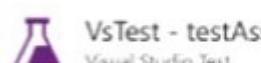
Prepare analysis on SonarQube

Prepare Analysis Configuration



Build solution ***.sln

Visual Studio Build



VsTest - testAssemblies

Visual Studio Test



Run Code Analysis

Run Code Analysis

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation

Step 1: Prepare Analysis Configuration

Prepare Analysis Configuration task, to configure all the required settings before executing the build.

This task is mandatory.

In case of .NET solutions or Java projects, it helps to integrate seamlessly with MSBuild, Maven and Gradle tasks.

Step 2: Visual Studio Build

Reorder the tasks to respect the following order:

Prepare Analysis Configuration task before any MSBuild or Visual Studio Build task.

Step 3: Visual Studio Test

Reorder the tasks to respect the following order:

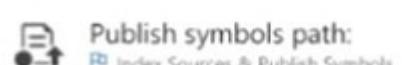
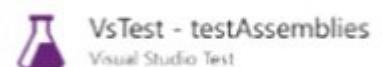
Run Code Analysis task after the Visual Studio Test task.

Step 4: Run Code Analysis

Run Code Analysis task, to actually execute the analysis of the source code.

This task is not required for Maven or Gradle projects, because scanner will be run as part of the Maven/Gradle build.

Note:



References: <https://docs.sonarqube.org/display/SCAN/Analyzing+with+SonarQube+Extension+for+VSTS-TFS>

QUESTION 93

You have an Azure DevOps organization named Contoso and an Azure DevOps project named Project1.

You plan to use Microsoft-hosted agents to build container images that will host full Microsoft .NET Framework apps in a YAML pipeline in Project1.

What are two possible virtual machine images that you can use for the Microsoft-hosted agent pool? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. vs2017-win2016
- B. ubuntu-16.04
- C. win1803
- D. macOS-10.13
- E. vs.2015-win2012r2

Answer: BC

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

The Microsoft-hosted agent pool provides 7 virtual machine images to choose from:

Ubuntu 16.04 (ubuntu-16.04)

Windows Server 1803 (win1803) - for running Windows containers

Visual Studio 2019 Preview on Windows Server 2019 (windows-2019)

Visual Studio 2017 on Windows Server 2016 (vs2017-win2016)

Visual Studio 2015 on Windows Server 2012R2 (vs2015-win2012r2)

macOS X Mojave 10.14 (macOS-10.14)

macOS X High Sierra 10.13 (macOS-10.13)

References:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/agents/hosted?view=azure-devops>

QUESTION 94

HOTSPOT

You currently use JIRA, Jenkins, and Octopus as part of your DevOps processes.

You plan to use Azure DevOps to replace these tools.

Which Azure DevOps service should you use to replace each tool? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

JIRA:

A dropdown menu for JIRA with the following options: Boards, Build pipelines, Release pipelines, and Repos. The 'Boards' option is highlighted with a black border.

Jenkins:

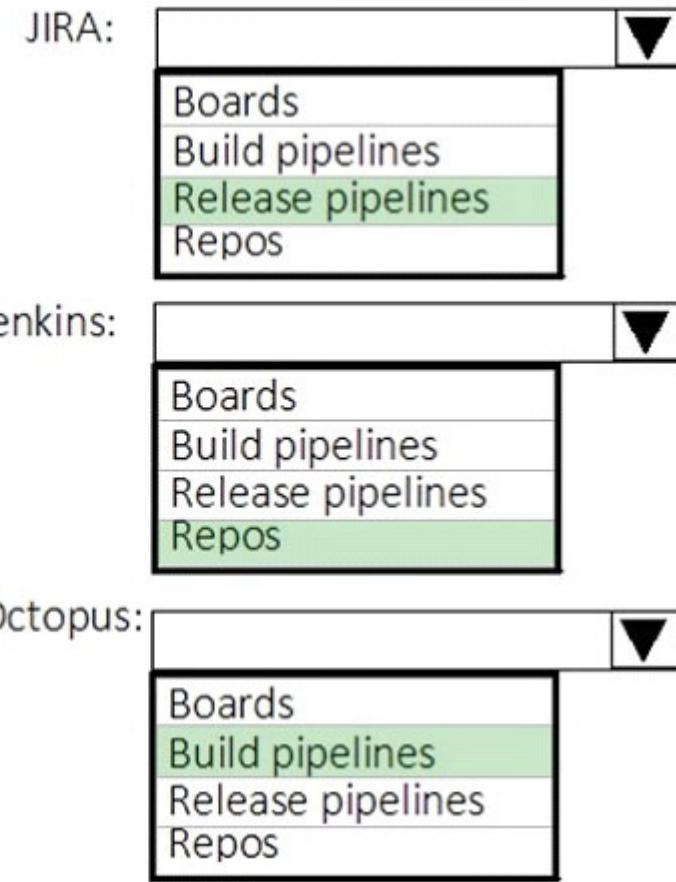
A dropdown menu for Jenkins with the following options: Boards, Build pipelines, Release pipelines, and Repos. The 'Boards' option is highlighted with a black border.

Octopus:

A dropdown menu for Octopus with the following options: Boards, Build pipelines, Release pipelines, and Repos. The 'Boards' option is highlighted with a black border.

A. Answer:

Answer Area



Explanation

Explanation/Reference:
Explanation:

JIRA: Release pipelines
Atlassian's Jira Software is a popular application that helps teams to plan, track, and manage software releases, whereas Octopus Deploy helps teams automate their development and operations processes in a fast, repeatable, and reliable manner. Together, they enable teams to get better end-to-end visibility into their software pipelines from idea to production.

Jenkins: Repos
One way to integrate Jenkins with Azure Pipelines is to run CI jobs in Jenkins separately. This involves configuration of a CI pipeline in Jenkins and a web hook in Azure DevOps that invokes the CI process when source code is pushed to a repository or a branch.

Octopus: Build pipelines

References:
<https://octopus.com/blog/octopus-jira-integration>

<https://www.azuredevopslabs.com/labs/vstsextend/jenkins/>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

JIRA: Release pipelines

Atlassian's Jira Software is a popular application that helps teams to plan, track, and manage software releases, whereas Octopus Deploy helps teams automate their development and operations processes in a fast, repeatable, and reliable manner. Together, they enable teams to get better end-to-end visibility into their software pipelines from idea to production.

Jenkins: Repos

One way to integrate Jenkins with Azure Pipelines is to run CI jobs in Jenkins separately. This involves configuration of a CI pipeline in Jenkins and a web hook in Azure DevOps that invokes the CI process when source code is pushed to a repository or a branch.

Octopus: Build pipelines**References:**

<https://octopus.com/blog/octopus-jira-integration>

<https://www.azuredevopslabs.com/labs/vstsextend/jenkins/>

QUESTION 95

Your company has a project in Azure DevOps.

You need to ensure that when there are multiple builds pending deployment, only the most recent build is deployed.

What should you use?

- A. deployment conditions
- B. deployment queue settings
- C. release gates
- D. pull request triggers

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

The options you can choose for a queuing policy are:

Number of parallel deployments

If you specify a maximum number of deployments, two more options appear:

- Deploy all in sequence
- Deploy latest and cancel the others: Use this option if you are producing releases faster than builds, and you only want to deploy the latest build.

Incorrect Answers:

C: Release gates allow automatic collection of health signals from external services, and then promote the release when all the signals are successful at the same time or stop the deployment on timeout. Typically, gates are used in connection with

incident management, problem management, change management, monitoring, and external approval systems.

References:
<https://docs.microsoft.com/en-us/azure/devops/pipelines/process/stages?tabs=classic&view=azure-devops#queuing-policies>

QUESTION 96

Your company develops a client banking application that processes a large volume of data.

Code quality is an ongoing issue for the company. Recently, the code quality has deteriorated because of an increase in time pressure on the development team.

You need to implement static code analysis.

During which phase should you use static code analysis?

- A. integration testing
- B. staging
- C. production release
- D. build

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

The Secure Development Lifecycle (SDL) Guidelines recommend that teams perform static analysis during the implementation phase of their development cycle.

Note: The company should focus in particular on the implementation of DevOps tests to assess the quality of the software from the planning stage to the implementation phase of the project.

References:

<https://secdevtools.azurewebsites.net/>

QUESTION 97

Your company is building a new solution in Java.

The company currently uses a SonarQube server to analyze the code of .NET solutions.

You need to analyze and monitor the code quality of the Java solution.

Which task types should you add to the build pipeline?

- A. Grunt
- B. Chef
- C. Maven
- D. Gulp

Answer: C

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

SonarQube is a set of static analyzers that can be used to identify areas of improvement in your code. It allows you to analyze the technical debt in your project and keep track of it in the future. With Maven and Gradle build tasks, you can run SonarQube analysis with minimal setup in a new or existing Azure DevOps Services build task.

References:

<https://docs.microsoft.com/en-us/azure/devops/java/sonarqube?view=azure-devops>

QUESTION 98

DRAG DROP

Your company has a project in Azure DevOps.

You plan to create a release pipeline that will deploy resources by using Azure Resource Manager templates. The templates will reference secrets stored in Azure Key Vault.

You need to recommend a solution for accessing the secrets stored in the key vault during deployments. The solution must use the principle of least privilege.

What should you include in the recommendation? To answer, drag the appropriate configurations to the correct targets. Each configuration may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Configurations

an Azure Key Vault access policy

a personal access token (PAT)

RBAC

Answer Area

Restrict access to delete the key vault:

Restrict access to the secrets in Key Vault by using:

A. Answer:

Configurations

an Azure Key Vault access policy

a personal access token (PAT)

RBAC

Answer Area

Restrict access to delete the key vault: an Azure Key Vault access policy

Restrict access to the secrets in Key Vault by using: RBAC

Explanation

Explanation/Reference:

Explanation:

Box 1: An Azure Key Vault access policy

The screenshot shows the 'mykeyvault0920 - Access policies' page in the Azure portal. The left sidebar lists 'Overview', 'Activity log', 'Access control (IAM)', 'Tags', 'Diagnose and solve problems', 'Settings' (with 'Keys', 'Secrets', 'Certificates'), 'Access policies' (which is selected and highlighted in blue), and 'Firewalls and virtual networks'. The main content area has a header with 'Save', 'Discard', and 'Refresh' buttons. It includes a note 'Click to hide advanced access policies'. Underneath are several checkboxes: 'Enable access to Azure Virtual Machines for deployment' (unchecked), 'Enable access to Azure Resource Manager for template deployment' (checked), and 'Enable access to Azure Disk Encryption for volume encryption' (unchecked). At the bottom, there's a 'Add new' button and a list entry for a user: '<Your username>' under 'USER'.

Box 2: RBAC

Management class access control uses RBAC

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation

Box 1: An Azure Key Vault access policy

The screenshot shows the 'Access policies' section of the Azure Key Vault 'mykeyvault0920'. The left sidebar lists 'Overview', 'Activity log', 'Access control (IAM)', 'Tags', 'Diagnose and solve problems', 'Settings' (with 'Keys', 'Secrets', 'Certificates'), 'Access policies' (which is selected and highlighted in blue), and 'Firewalls and virtual networks'. The main area has a search bar, save, discard, and refresh buttons. It displays a list of access policies. One policy is shown with a checked checkbox and a red border around the entire row. The policy details are: 'Enable access to Azure Resource Manager for template deployment' (checkbox checked). Below this is another policy with an unchecked checkbox: 'Enable access to Azure Disk Encryption for volume encryption'. At the bottom, there is an 'Add new' button and a list item for '<Your username> USER'.

Box 2: RBAC

Management plane access control uses RBAC.

The management plane consists of operations that affect the key vault itself, such as:

Creating or deleting a key vault.

Getting a list of vaults in a subscription.

Retrieving Key Vault properties (such as SKU and tags).

Setting Key Vault access policies that control user and application access to keys and secrets.

References:

<https://docs.microsoft.com/en-us/azure/azure-resource-manager/resource-manager-tutorial-use-key-vault>

QUESTION 99

SIMULATION

You have an Azure function hosted in an App Service plan named az400-9940427-func1.

You need to configure az400-9940427-func1 to upgrade the functions automatically whenever new code is committed to the master branch of <https://github.com/Azure-Samples/functions-quickstart>.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

1. Open Microsoft Azure Portal
2. Log into your Azure account, select App Services in the Azure portal left navigation, and then select configure az400-9940427-func1.
3. On the app page, select Deployment Center in the left menu.
4. On the Build provider page, select Azure Pipelines (Preview), and then select Continue.
5. On the Configure page, in the Code section:
For GitHub, drop down and select the Organization, Repository, and Branch you want to deploy continuously.
6. Select Continue.
7. On the Test page, choose whether to enable load tests, and then select Continue.
8. Depending on your App Service plan pricing tier, you may see a Deploy to staging page. Choose whether to enable deployment slots, and then select Continue.
9. After you configure the build provider, review the settings on the Summary page, and then select Finish.

References:

<https://docs.microsoft.com/en-us/azure/app-service/deploy-continuous-deployment>

Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

1. Open Microsoft Azure Portal

2. Log into your Azure account, select App Services in the Azure portal left navigation, and then select configure az400-9940427-func1.
3. On the app page, select Deployment Center in the left menu.
4. On the Build provider page, select Azure Pipelines (Preview), and then select Continue.
5. On the Configure page, in the Code section:
For GitHub, drop down and select the Organization, Repository, and Branch you want to deploy continuously.
6. Select Continue.
7. On the Test page, choose whether to enable load tests, and then select Continue.
8. Depending on your App Service plan pricing tier, you may see a Deploy to staging page. Choose whether to enable deployment slots, and then select Continue.
9. After you configure the build provider, review the settings on the Summary page, and then select Finish.

References:

<https://docs.microsoft.com/en-us/azure/app-service/deploy-continuous-deployment>

QUESTION 100

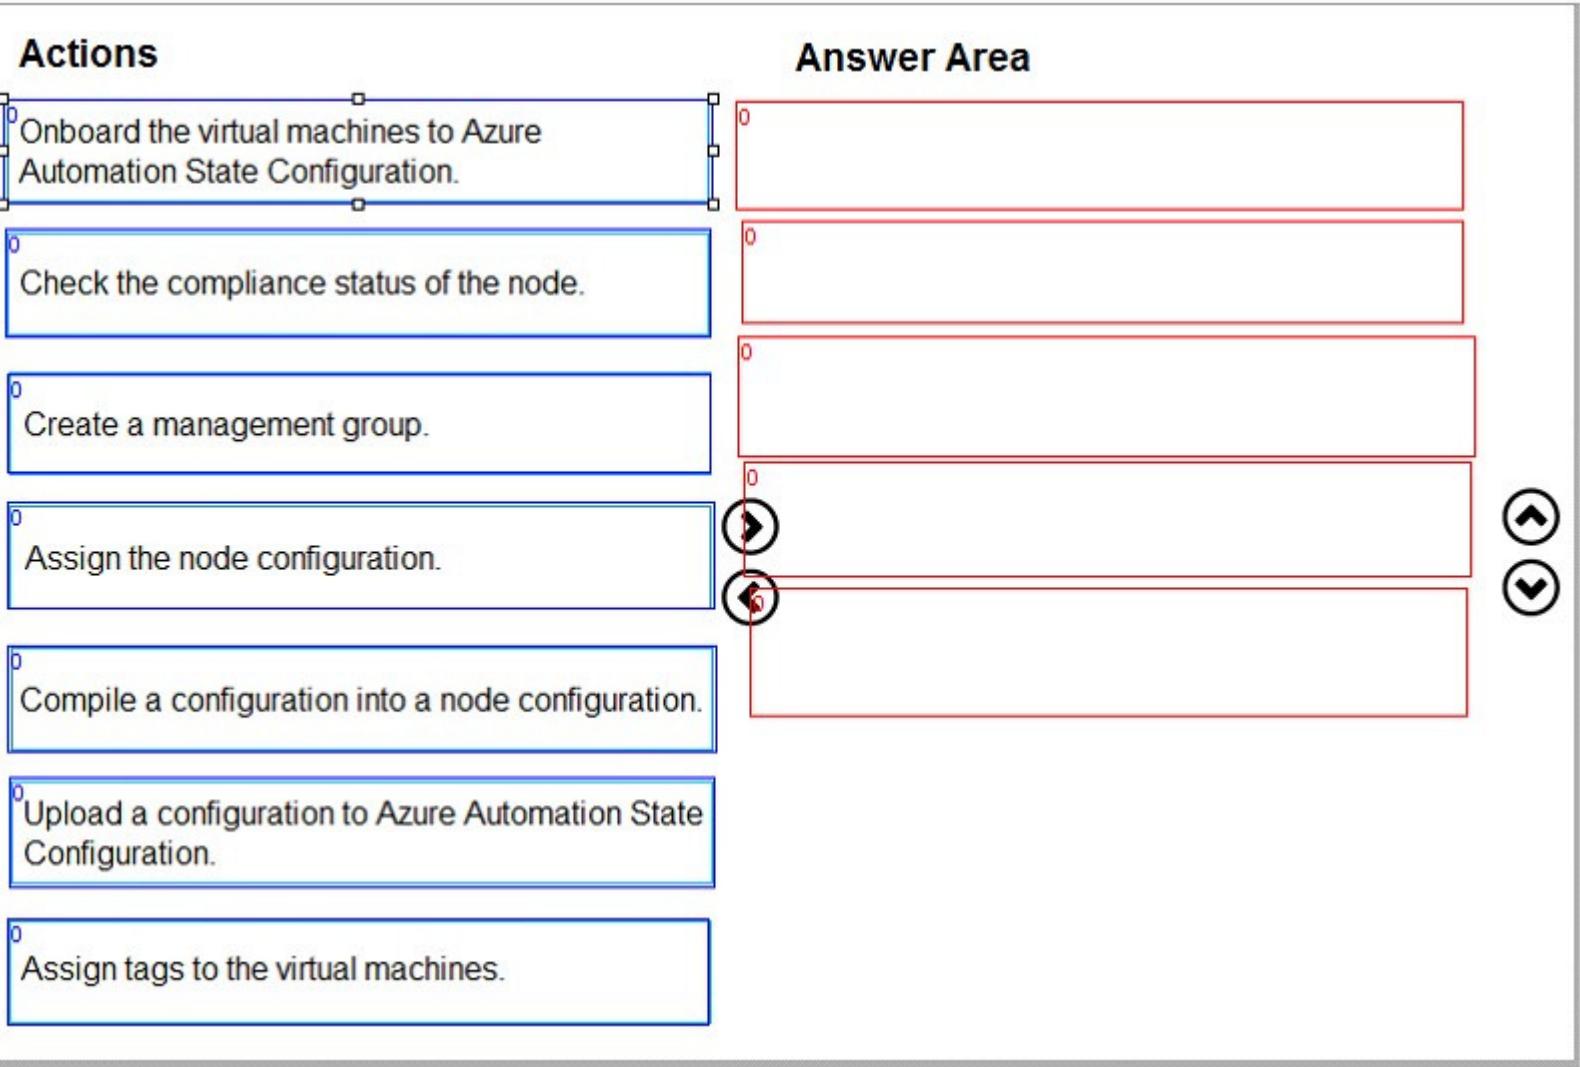
DRAG DROP

You need to use Azure Automation State Configuration to manage the ongoing consistency of virtual machine configurations.

Which five actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

NOTE: More than one order of answer choices is correct. You will receive credit for any of the correct orders you select.

Select and Place:



A. Answer:

Actions	Answer Area
	Assign the node configuration.
	Upload a configuration to Azure Automation State Configuration.
Create a management group.	Compile a configuration into a node configuration.
	Onboard the virtual machines to Azure Automation State Configuration. 
	Check the compliance status of the node. 
Assign tags to the virtual machines.	

Explanation

Explanation/Reference:

Explanation:

Step 1: Assign the node configuration.

You create a simple DSC configuration that ensures either the presence or absence of the Web-Server Windows Feature (IIS), depending on how you assign nodes.

Step 2: Upload a configuration to Azure Automation State Configuration.

You import the configuration into the Automation account.

Step 3: Compiling a configuration into a node configuration

Compiling a configuration in Azure Automation

Before you can apply a desired state to a node, a DSC configuration defining that state must be compiled into one or more node configurations (MOF document), and placed on the Automation DSC Pull Server.

Step 4: Onboard the virtual machines to Azure State Configuration

Onboarding an Azure VM for management with Azure Automation State Configuration

Step 5: Check the compliance status of the node.

Viewing reports for managed nodes. Each time Azure Automation State Configuration performs a consistency check on a managed node, the node sends a status report back to the pull server. You can view these reports on the page for that node.

On the blade for an individual report, you can see the following status information for the corresponding consistency check:

The report status -- whether the node is "Compliant", the configuration "Failed", or the node is "Not Compliant" (when the node is in ApplyandMonitor mode and the machine is not in the desired state).

References: <https://docs.microsoft.com/en-us/azure/automation/automation-dsc-getting-started>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation

Step 1: Assign the node configuration.

You create a simple DSC configuration that ensures either the presence or absence of the Web-Server Windows Feature (IIS), depending on how you assign nodes.

Step 2: Upload a configuration to Azure Automation State Configuration.

You import the configuration into the Automation account.

Step 3: Compiling a configuration into a node configuration

Compiling a configuration in Azure Automation

Before you can apply a desired state to a node, a DSC configuration defining that state must be compiled into one or more node configurations (MOF document), and placed on the Automation DSC Pull Server.

Step 4: Onboard the virtual machines to Azure State Configuration

Onboarding an Azure VM for management with Azure Automation State Configuration

Step 5: Check the compliance status of the node.

Viewing reports for managed nodes. Each time Azure Automation State Configuration performs a consistency check on a managed node, the node sends a status report back to the pull server. You can view these reports on the page for that node.

On the blade for an individual report, you can see the following status information for the corresponding consistency check:

The report status -- whether the node is "Compliant", the configuration "Failed", or the node is "Not Compliant" (when the node is in ApplyandMonitor mode and the machine is not in the desired state).

References: <https://docs.microsoft.com/en-us/azure/automation/automation-dsc-getting-started>

QUESTION 101

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

Your company has a project in Azure DevOps for a new web application.

You need to ensure that when code is checked in, a build runs automatically.

Solution: From the Continuous deployment trigger settings of the release pipeline, you enable the Pull request trigger setting.

Does this meet the goal?

A. Yes

B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation

Instead, In Visual Designer you enable continuous integration (CI) by:

1. Select the Triggers tab.
2. Enable Continuous integration.

Reference:
<https://docs.microsoft.com/en-us/azure/devops/pipelines/get-started-designer>

QUESTION 102

You use WhiteSource Bolt to scan a Node.js application.

The WhiteSource Bolt scan identifies numerous libraries that have invalid licenses. The libraries are used only during development and are not part of a production deployment.

You need to ensure that WhiteSource Bolt only scans production dependencies.

Which two actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Run npm install and specify the --production flag.
- B. Modify the WhiteSource Bolt policy and set the action for the licenses used by the development tools to Reassign.
- C. Modify the devDependencies section of the project's Package.json file.
- D. Configure WhiteSource Bolt to scan the node_modules directory only.

Answer: AC

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

A: To resolve NPM dependencies, you should first run "npm install" command on the relevant folders before executing the plugin.

C: All npm packages contain a file, usually in the project root, called package.json - this file holds various metadata relevant to the project. This file is used to give information to npm that allows it to identify the project as well as handle the project's dependencies. It can also contain other metadata such as a project description, the version of the project in a particular distribution, license information, even configuration data - all of which can be vital to both npm and to the end users of the package.

Reference:

<https://whitesource.atlassian.net/wiki/spaces/WD/pages/34209870/NPM+Plugin>

<https://nodejs.org/en/knowledge/getting-started/npm/what-is-the-file-package-json>

QUESTION 103

SIMULATION

You plan to deploy a runbook that will create Azure AD user accounts.

You need to ensure that runbooks can run the Azure PowerShell cmdlets for Azure Active Directory.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

Azure Automation now ships with the Azure PowerShell module of version 0.8.6, which introduced the ability to non-interactively authenticate to Azure using OrgId (Azure Active Directory user) credential-based authentication. Using the steps below, you can set up Azure Automation to talk to Azure using this authentication type.

Step 1: Find the Azure Active Directory associated with the Azure subscription to manage:

1. Log in to the Azure portal as the service administrator for the Azure subscription you want to manage using Azure Automation. You can find this user by logging in to the Azure portal as any user with access to this Azure subscription, then clicking Settings, then Administrators.



2. Note the name of the directory associated with the Azure subscription you want to manage. You can find this directory by clicking Settings, then Subscriptions.

A screenshot of the 'Subscriptions' page in the Azure portal. The top navigation bar includes 'SUBSCRIPTIONS', 'MANAGEMENT CERTIFICATES', 'ADMINISTRATORS', 'AFFINITY GROUPS', and 'USAGE'. Below this is a table with columns: 'SUBSCRIPTION', 'SUBSCRIPTION ID', 'ACCOUNT ADMINISTRATOR', and 'DIRECTORY'. The first row shows 'Windows Azure MSDN - Visual Studio Ultimate' in the 'SUBSCRIPTION' column, a redacted 'SUBSCRIPTION ID', and 'Joe Levy' in the 'ACCOUNT ADMINISTRATOR' column, which is highlighted with a red box. The 'DIRECTORY' column is also redacted.

Step 2: Create an Azure Active Directory user in the directory associated with the Azure subscription to manage:

You can skip this step if you already have an Azure Active Directory user in this directory. and plan to use this OrgId to manage Azure.

1. In the Azure portal click on Active Directory service.



Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

Azure Automation now ships with the Azure PowerShell module of version 0.8.6, which introduced the ability to non-interactively authenticate to Azure using OrgId (Azure Active Directory user) credential-based authentication. Using the steps below, you can set up Azure Automation to talk to Azure using this authentication type.

Step 1: Find the Azure Active Directory associated with the Azure subscription to manage:

1. Log in to the Azure portal as the service administrator for the Azure subscription you want to manage using Azure Automation. You can find this user by logging in to the Azure portal as any user with access to this Azure subscription, then clicking Settings, then Administrators.



2. Note the name of the directory associated with the Azure subscription you want to manage. You can find this directory by clicking Settings, then Subscriptions.

A screenshot of the Azure portal's Subscriptions page. At the top, there are tabs for Subscriptions, Management Certificates, Administrators, Affinity Groups, and Usage. Below these are four columns: Subscription, Subscription ID, Account Administrator, and Directory. The 'Subscription' column shows 'Windows Azure MSDN - Visual Studio Ultimate'. The 'Account Administrator' column shows 'Joe Levy', which is highlighted with a red box. The 'Directory' column shows a redacted URL.

Step 2: Create an Azure Active Directory user in the directory associated with the Azure subscription to manage:

You can skip this step if you already have an Azure Active Directory user in this directory. and plan to use this OrgId to manage Azure.

1. In the Azure portal click on Active Directory service.



2. Click the directory name that is associated with this Azure subscription.
3. Click on the Users tab and then click the Add User button.
4. For type of user, select "New user in your organization." Enter a username for the user to create.
5. Fill out the user's profile. For role, pick "User." Don't enable multi-factor authentication. Multi-factor accounts cannot be used with Azure Automation.
6. Click Create.
7. Jot down the full username (including part after @ symbol) and temporary password.

Step 3: Allow this Azure Active Directory user to manage this Azure subscription.

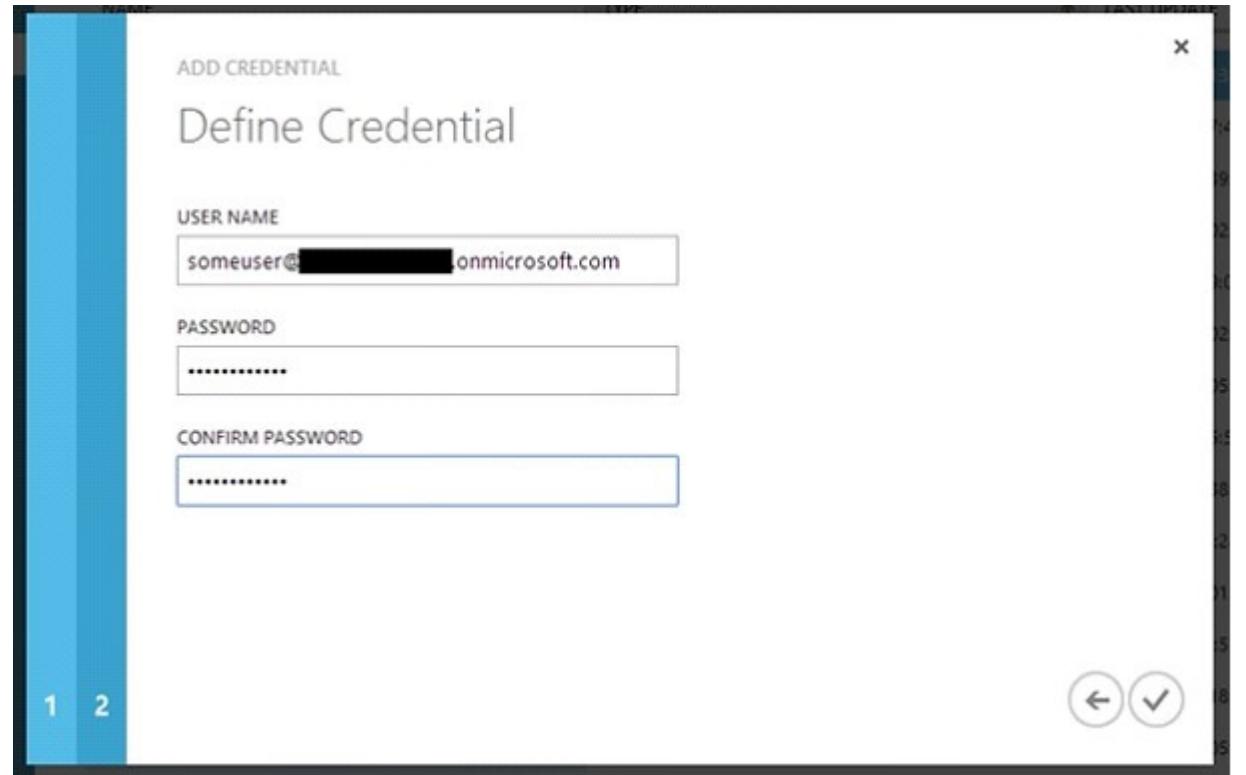
1. Click on Settings (bottom Azure tab under StorSimple)



2. Click Administrators
3. Click the Add button. Type the full user name (including part after @ symbol) of the Azure Active Directory user you want to set up to manage Azure. For subscriptions, choose the Azure subscriptions you want this user to be able to manage. Click the check mark.

Step 4: Configure Azure Automation to use this Azure Active Directory user to manage this Azure subscription

Create an Azure Automation credential asset containing the username and password of the Azure Active Directory user that you have just created. You can create a credential asset in Azure Automation by clicking into an Automation Account and then clicking the Assets tab, then the Add Setting button.



Note: Once you have set up the Azure Active Directory credential in Azure and Azure Automation, you can now manage Azure from Azure Automation runbooks using this credential.

References:

<https://azure.microsoft.com/sv-se/blog/azure-automation-authenticating-to-azure-using-azure-active-directory/>

QUESTION 104

DRAG DROP

You are creating a container for an ASP.NET Core app.

You need to create a Dockerfile file to build the image. The solution must ensure that the size of the image is minimized.

How should you configure the file? To answer, drag the appropriate values to the correct targets. Each value must be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Answer Area

Values

```
dotnet publish -c Release -o out
```

```
dotnet restore
```

```
microsoft/dotnet:2.2-aspnetcore-runtime
```

```
Microsoft/dotnet:2.2-sdk
```

FROM [REDACTED] As build-env

COPY . /app/

WORKDIR /app

RUN [REDACTED]

FROM [REDACTED]

COPY --from=build-env /app/out /app

WORKDIR /app

ENTRYPOINT ["dotnet", "MvcMovie.dll"]

A. Answer:

Answer Area

Values	FROM Microsoft/dotnet:2.2-sdk As build-env
dotnet publish -c Release -o out	COPY . /app/
dotnet restore	WORKDIR /app
microsoft/dotnet:2.2-aspnetcore-runtime	RUN dotnet restore
Microsoft/dotnet:2.2-sdk	FROM microsoft/dotnet:2.2-aspnetcore-runtime
	COPY --from=build-env /app/out /app
	WORKDIR /app
	ENTRYPOINT ["dotnet", "MvcMovie.dll"]

Explanation

Explanation/Reference:

Explanation:

Box 1: microsoft.com/dotnet/sdk:2.3

The first group of lines declares from which base image we will use to build our container on top of. If the local system does not have this image already, then docker will automatically try and fetch it. The mcr.microsoft.com/dotnet/core/sdk:2.1 comes packaged

with the .NET core 2.1 SDK installed, so it's up to the task of building ASP .NET core projects targeting version 2.1

Box 2: dotnet restore

The next instruction changes the working directory in our container to be /app, so all commands following this one execute under this context.

```
COPY *.csproj ./  
RUN dotnet restore
```

Box 3: microsoft.com/dotnet/2.2-aspnetcore-runtime

When building container images, it's good practice to include only the production payload and its dependencies in the container image. We don't want the .NET core SDK included in our final image because we only need the .NET core runtime, so the dockerfile

is written to use a temporary container that is packaged with the SDK called build-env to build the app.

Reference:

<https://docs.microsoft.com/de-DE/virtualization/windowscontainers/quick-start/building-sample-app>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box 1: microsoft.com/dotnet/sdk:2.3

The first group of lines declares from which base image we will use to build our container on top of. If the local system does not have this image already, then docker will automatically try and fetch it. The mcr.microsoft.com/dotnet/core/sdk:2.1 comes packaged with the .NET core 2.1 SDK installed, so it's up to the task of building ASP .NET core projects targeting version 2.1

Box 2: dotnet restore

The next instruction changes the working directory in our container to be /app, so all commands following this one execute under this context.

```
COPY *.csproj ./  
RUN dotnet restore
```

Box 3: microsoft.com/dotnet/2.2-aspnetcore-runtime

When building container images, it's good practice to include only the production payload and its dependencies in the container image. We don't want the .NET core SDK included in our final image because we only need the .NET core runtime, so the dockerfile

is written to use a temporary container that is packaged with the SDK called build-env to build the app.

Reference:

<https://docs.microsoft.com/de-DE/virtualization/windowscontainers/quick-start/building-sample-app>

QUESTION 105

HOTSPOT

You plan to create alerts that will be triggered based on the page load performance of a home page.

You have the Application Insights log query shown in the following exhibit.

Application Insights

Demo



Analytics

Home Pa...

New Query 1*

New Query...*X +

Demo

Run

Time range: Set in query

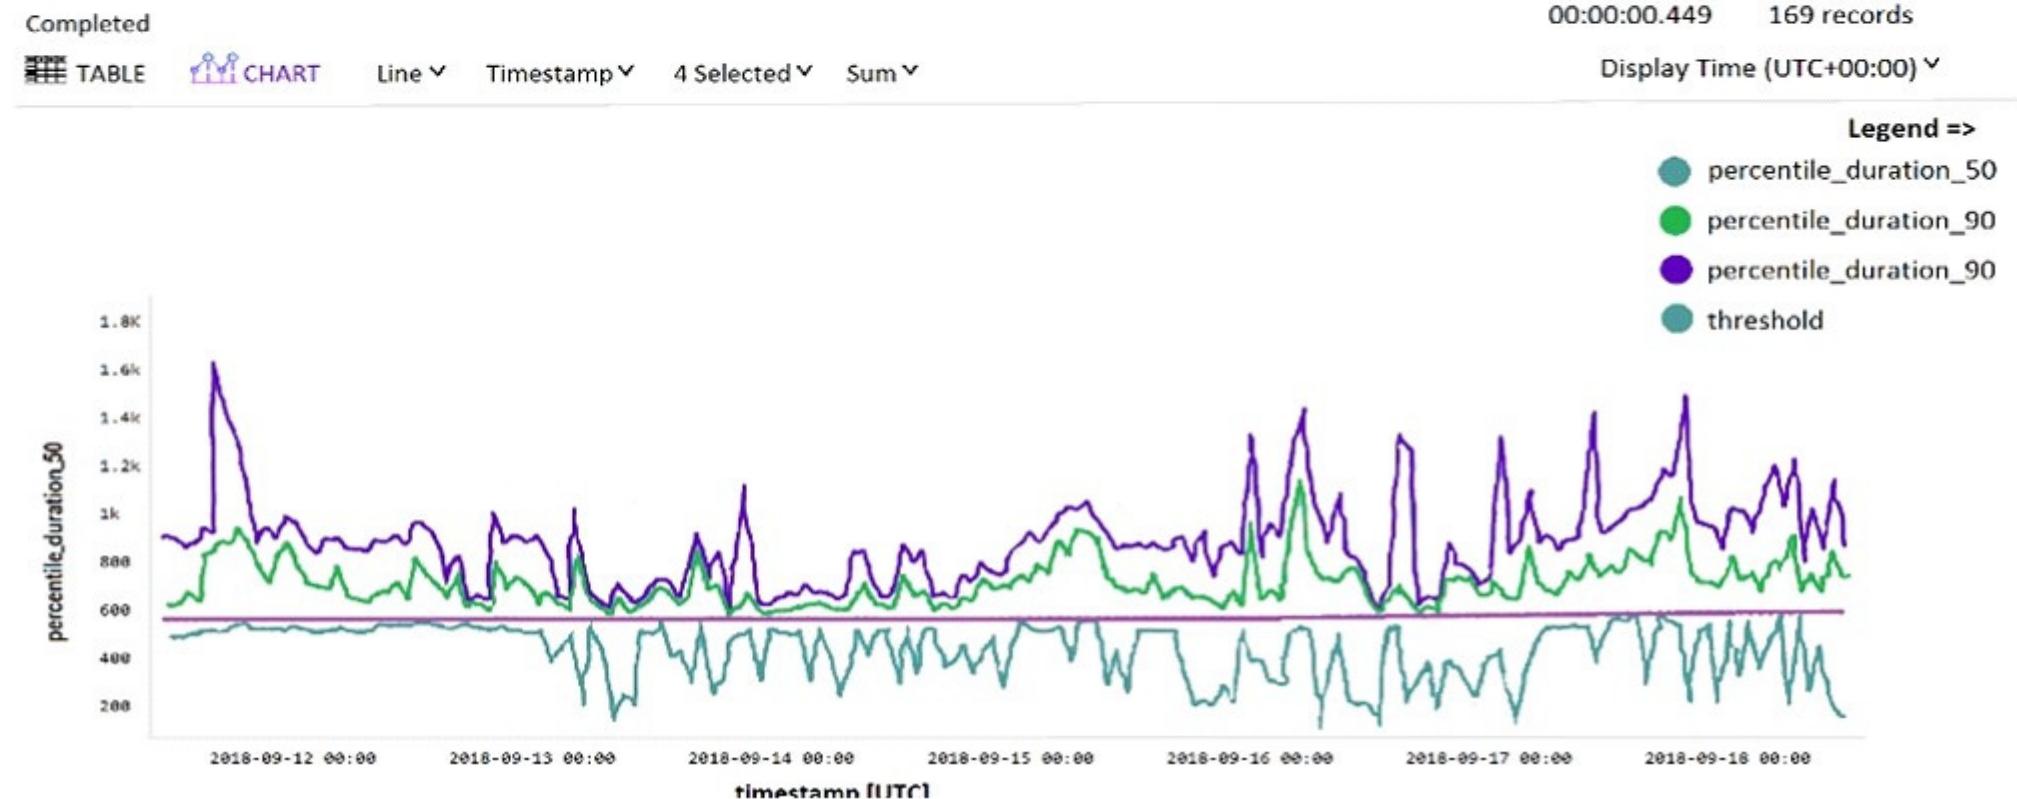
Schema Filter (preview) <<

Filter by name 🔍

① Showing top 10 values for each facet

- ✓ name (1)
 - █ GET Home/Index (38852)
- ✓ operation_Name (1)
 - █ GET Home/Index (38852)
- ✓ performanceBucket (8)
 - █ 500ms-1sec (18690)
 - █ <250ms (17098)
 - █ 1sec-3sec (1620)
 - █ 250ms-500ms (1354)
 - █ 15sec-30sec (57)
 - █ 3sec-7sec (19)
 - █ 7sec-15sec (10)
 - █ 30sec-1min (4)
- ✓ resultCode (1)
 - █ 200 (38852)
- ✓ success (1)
 - █ True (38852)

```
requests
| where timestamp >= ago(7d)
| where operation_Name endswith( 'Home/Index' )
| where operation_Name startswith( 'GET' )
| summarize percentiles(duration, 50, 90, 95) by bin(timestamp, 1h)
| extend threshold=675
| render timechart
```



Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

To create an alert based on the page load experience of most users, the alerting level must be based on [answer choice].

	▼
percentile_duration_50	
percentile_duration_90	
percentile_duration_95	
threshold	

To only create an alert when authentication error occurs on the server, the query must be filtered on [answer choice].

	▼
item Type	
resultCode	
source	
success	

A. Answer:

Answer Area

To create an alert based on the page load experience of most users, the alerting level must be based on [answer choice].

percentile_duration_50
percentile_duration_90
percentile_duration_95
threshold

To only create an alert when authentication error occurs on the server, the query must be filtered on [answer choice].

item Type
resultCode
source
success

Explanation

Explanation/Reference:
Explanation:

Box 1: percentile_duration_95

Box 2: success

For example -

requests
| project name, url, success
| where success == "False"

This will return all the failed requests in my App Insights within the specified time range.

Reference:
<https://devblogs.microsoft.com/premier-developer/alerts-based-on-analytics-query-using-custom-log-search/>

Answer:

Explanation/Reference:
Explanation

Explanation/Reference:
Explanation:

Box 1: percentile_duration_95

Box 2: success

For example -

requests
| project name, url, success
| where success == "False"

This will return all the failed requests in my App Insights within the specified time range.

Reference:
<https://devblogs.microsoft.com/premier-developer/alerts-based-on-analytics-query-using-custom-log-search/>

QUESTION 106

DRAG DROP

You are configuring the settings of a new Git repository in Azure Repos.

You need to ensure that pull requests in a branch meet the following criteria before they are merged:
Committed code must compile successfully.

Pull requests must have a Quality Gate status of Passed in SonarCloud.

Which policy type should you configure for each requirement? To answer, drag the appropriate policy types to the correct requirements. Each policy type may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Answer Area

Policy Types

A build policy

A check-in policy

A status policy

Committed code must compile successfully:

Pull requests must have a Quality Gate status
of Passed in SonarCloud:

A. Answer:

Answer Area

Policy Types

A build policy

A check-in policy

A status policy

Committed code must compile successfully:

A check-in policy

Pull requests must have a Quality Gate status
of Passed in SonarCloud:

A build policy

Explanation

Explanation/Reference:

Explanation:

Box 1: A check-in policy

Administrators of Team Foundation version control can add check-in policy requirements. These check-in policies require the user to take actions when they conduct a check-in to source control.

By default, the following check-in policy types are available:

Builds Requires that the last build was successful before a check-in.

Code Analysis Requires that code analysis is run before check-in.

Work Items Requires that one or more work items be associated with the check-in.

Box 2: Build policy

Reference:

<https://docs.microsoft.com/en-us/azure/devops/repos/tfvc/add-check-policies>

<https://azuredevopslabs.com/labs/vstsextend/sonarcloud/>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box 1: A check-in policy

Administrators of Team Foundation version control can add check-in policy requirements. These check-in policies require the user to take actions when they conduct a check-in to source control.

By default, the following check-in policy types are available:

Builds Requires that the last build was successful before a check-in.

Code Analysis Requires that code analysis is run before check-in.

Work Items Requires that one or more work items be associated with the check-in.

Box 2: Build policy

Reference:

<https://docs.microsoft.com/en-us/azure/devops/repos/tfvc/add-check-policies>

<https://azuredevopslabs.com/labs/vstsextend/sonarcloud/>

QUESTION 107

You use a Git repository in Azure Repos to manage the source code of a web application. Developers commit changes directly to the master branch.

You need to implement a change management procedure that meets the following requirements:

The master branch must be protected, and new changes must be built in the feature branches first.

Changes must be reviewed and approved by at least one release manager before each merge.

Changes must be brought into the master branch by using pull requests.

What should you configure in Azure Repos?

- A. branch policies of the master branch
- B. Services in Project Settings
- C. Deployment pools in Project Settings
- D. branch security of the master branch

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Branch policies help teams protect their important branches of development. Policies enforce your team's code quality and change management standards.

Reference:

<https://docs.microsoft.com/en-us/azure/devops/repos/git/branch-policies>

QUESTION 108

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You plan to update the Azure DevOps strategy of your company.

You need to identify the following issues as they occur during the company's development process:

Licensing violations

Prohibited libraries

Solution: You implement continuous integration.

Does this meet the goal?

- A. Yes
- B. No

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

WhiteSource is the leader in continuous open source software security and compliance management. WhiteSource integrates into your build process, irrespective of your programming languages, build tools, or development environments. It works automatically, continuously, and silently in the background, checking the security, licensing, and quality of your open source components against WhiteSource constantly-updated definitive database of open source repositories.

Reference:

<https://azuredovplabs.com/labs/vstsexpand/whitesource/>

QUESTION 109

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You plan to update the Azure DevOps strategy of your company.

You need to identify the following issues as they occur during the company's development process:

Licensing violations

Prohibited libraries

Solution: You implement pre-deployment gates.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Instead use implement continuous integration.

Note: WhiteSource is the leader in continuous open source software security and compliance management. WhiteSource integrates into your build process, irrespective of your programming languages, build tools, or development environments. It works automatically, continuously, and silently in the background, checking the security, licensing, and quality of your open source components against WhiteSource constantly-updated definitive database of open source repositories.

Reference:

<https://azuredevopslabs.com/labs/vstsextend/whitesource/>

QUESTION 110

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You plan to update the Azure DevOps strategy of your company.

You need to identify the following issues as they occur during the company's development process:

Licensing violations

Prohibited libraries

Solution: You implement automated security testing.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Instead use implement continuous integration.

Note: WhiteSource is the leader in continuous open source software security and compliance management. WhiteSource integrates into your build process, irrespective of your programming languages, build tools, or development environments. It works automatically, continuously, and silently in the background, checking the security, licensing, and quality of your open source components against WhiteSource constantly-updated definitive database of open source repositories.

Reference:
<https://azuredevopslabs.com/labs/vstsexpand/whitesource/>

QUESTION 111

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

Your company uses Azure DevOps to manage the build and release processes for applications.

You use a Git repository for applications source control.

You need to implement a pull request strategy that reduces the history volume in the master branch.

Solution: You implement a pull request strategy that uses fast-forward merges.

Does this meet the goal?

- A. Yes
- B. No

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

No fast-forward merge - This option merges the commit history of the source branch when the pull request closes and creates a merge commit in the target branch.

Reference:

<https://docs.microsoft.com/en-us/azure/devops/repos/git/branch-policies>

QUESTION 112

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

Your company uses Azure DevOps to manage the build and release processes for applications.

You use a Git repository for applications source control.

You need to implement a pull request strategy that reduces the history volume in the master branch.

Solution: You implement a pull request strategy that uses squash merges.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Instead use fast-forward merge.

Note:

Squash merge - Complete all pull requests with a squash merge, creating a single commit in the target branch with the changes from the source branch.

No fast-forward merge - This option merges the commit history of the source branch when the pull request closes and creates a merge commit in the target branch.

Reference:

<https://docs.microsoft.com/en-us/azure/devops/repos/git/branch-policies>

QUESTION 113

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

Your company uses Azure DevOps to manage the build and release processes for applications.

You use a Git repository for applications source control.

You need to implement a pull request strategy that reduces the history volume in the master branch.

Solution: You implement a pull request strategy that uses an explicit merge.

Does this meet the goal?

A. Yes

B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Instead use fast-forward merge.

Note:

No fast-forward merge - This option merges the commit history of the source branch when the pull request closes and creates a merge commit in the target branch.

Reference:

<https://docs.microsoft.com/en-us/azure/devops/repos/git/branch-policies>

QUESTION 114

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

Your company uses Azure DevOps to manage the build and release processes for applications.

You use a Git repository for applications source control.

You need to implement a pull request strategy that reduces the history volume in the master branch.

Solution: You implement a pull request strategy that uses a three-way merge.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Instead use fast-forward merge.

Note:

No fast-forward merge - This option merges the commit history of the source branch when the pull request closes and creates a merge commit in the target branch.

Reference:

<https://docs.microsoft.com/en-us/azure/devops/repos/git/branch-policies>

QUESTION 115

You are developing an application. The application source has multiple branches.

You make several changes to a branch used for experimentation.

You need to update the main branch to capture the changes made to the experimentation branch and override the history of the Git repository.

Which Git option should you use?

- A. Rebase
- B. Fetch
- C. Merge
- D. Push

Answer: C

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Create pull requests to review and merge code in a Git project. Pull requests let your team review code and give feedback on changes before merging it into the master branch.

Incorrect Answers:

A: Use rebase to address the problem of updating your branch with the latest changes from the main branch. Rebase takes the changes made in the commits in your current branch and replays them on the history of another branch. The commit history of your current branch will be rewritten so that it starts from the most recent commit in the target branch of the rebase. Rebasing your changes in your feature branch off the latest changes in the main branch lets you test your changes on the most recent version in the

main branch while keeping a clean Git history.

D: Share changes made in commits and branches using the push command. Push your branches to the remote repository. Git adds your commits to an existing branch on the remote or creates a new branch with the same commits as your local branch.

Reference:
<https://docs.microsoft.com/en-us/azure/devops/repos/git/pull-requests>

QUESTION 116

Testlet 1

Case Study

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other question in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview

Litware, Inc. is an independent software vendor (ISV). Litware has a main office and five branch offices.

Existing Environment

Application Architecture

The company's primary application is a single monolithic retirement fund management system based on ASP.NET web forms that use logic written in VB.NET. Some new sections of the application are written in C#.

Variations of the application are created for individual customers. Currently, there are more than 80 live code branches in the application's code base.

The application was developed by using Microsoft Visual Studio. Source code is stored in Team Foundation Server (TFS) in the main office. The branch offices access the source code by using TFS proxy servers.

Architectural Issues

Litware focuses on writing new code for customers. No resources are provided to refactor or remove existing code. Changes to the code base take a long time, as dependencies are not obvious to individual developers.

Merge operations of the code often take months and involve many developers. Code merging frequently introduces bugs that are difficult to locate and resolve.

Customers report that ownership costs of the retirement fund management system increase continually. The need to merge unrelated code makes even minor code changes expensive.

Customers report that bug reporting is overly complex.

Requirements

Planned changes

Litware plans to develop a new suite of applications for investment planning. The investment planning applications will require only minor integration with the existing retirement fund management system.

The investment planning applications suite will include one multi-tier web application and two iOS mobile applications. One mobile application will be used by employees; the other will be used by customers.

Litware plans to move to a more agile development methodology. Shared code will be extracted into a series of packages.

Litware has started an internal cloud transformation process and plans to use cloud-based services whenever suitable.

Litware wants to become proactive in detecting failures, rather than always waiting for customer bug reports.

Technical requirements

The company's investment planning applications suite must meet the following requirements:

New incoming connections through the firewall must be minimized.

Members of a group named Developers must be able to install packages.

The principle of least privilege must be used for all permission assignments.

A branching strategy that supports developing new functionality in isolation must be used.

Members of a group named Team Leaders must be able to create new packages and edit the permissions of package feeds.

Visual Studio App Center must be used to centralize the reporting of mobile application crashes and device types in use.

By default, all releases must remain available for 30 days, except for production releases, which must be kept for 60 days.

Code quality and release quality are critical. During release, deployments must not proceed between stages if any active bugs are logged against the release.

The mobile applications must be able to call the share pricing service of the existing retirement fund management system. Until the system is upgraded, the service will only support basic authentication over HTTPS.

The required operating system configuration for the test servers changes weekly. Azure Automation State Configuration must be used to ensure that the operating system on each test server is configured the same way when the servers are created and

checked periodically.

Current Technical Issue

The test servers are configured correctly when first deployed, but they experience configuration drift over time. Azure Automation State Configuration fails to correct the configurations.

Azure Automation State Configuration nodes are registered by using the following command.

```
Register-AzureRmAutomationDscNode
-ResourceGroupName 'TestResourceGroup'
-AutomationAccountName 'LitwareAutomationAccount'
-AzureVMName $vname
-ConfigurationMode 'ApplyOnly'
```

What should you use to implement the code quality restriction on the release pipeline for the investment planning applications suite?

- A. a pre-deployment approval
- B. a deployment gate
- C. a post-deployment approval
- D. a trigger

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

When a release is created from a release pipeline that defines approvals, the deployment stops at each point where approval is required until the specified approver grants approval or rejects the release (or re-assigns the approval to another user).

Scenario: Code quality and release quality are critical. During release, deployments must not proceed between stages if any active bugs are logged against the release.

References: <https://docs.microsoft.com/en-us/azure/devops/pipelines/release/approvals/approvals>

QUESTION 117

HOTSPOT

How should you configure the release retention policy for the investment planning applications suite? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Global release:

- Set the default retention policy to 30 days.
- Set the maximum retention policy to 30 days.
- Set the stage retention policy to 30 days.**
- Set the stage retention policy to 60 days.

Production stage:

- Set the default retention policy to 30 days.
- Set the maximum retention policy to 60 days.
- Set the stage retention policy to 30 days.**
- Set the stage retention policy to 60 days.

A. Answer:

Answer Area

Global release:

- Set the default retention policy to 30 days.
- Set the maximum retention policy to 30 days.
- Set the stage retention policy to 30 days.
- Set the stage retention policy to 60 days.

Production stage:

- Set the default retention policy to 30 days.
- Set the maximum retention policy to 60 days.
- Set the stage retention policy to 30 days.
- Set the stage retention policy to 60 days.

Explanation

Explanation/Reference:

Explanation:

Scenario: By default, all releases must remain available for 30 days, except for production releases, which must be kept for 60 days.

Box 1: Set the default retention policy to 30 days

The Global default retention policy sets the default retention values for all the build pipelines. Authors of build pipelines can override these values.

Box 2: Set the stage retention policy to 60 days

You may want to retain more releases that have been deployed to specific stages.

References: <https://docs.microsoft.com/en-us/azure/devops/pipelines/policies/retention>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Scenario: By default, all releases must remain available for 30 days, except for production releases, which must be kept for 60 days.

Box 1: Set the default retention policy to 30 days

The Global default retention policy sets the default retention values for all the build pipelines. Authors of build pipelines can override these values.

Box 2: Set the stage retention policy to 60 days

You may want to retain more releases that have been deployed to specific stages.

References: <https://docs.microsoft.com/en-us/azure/devops/pipelines/policies/retention>

QUESTION 118

HOTSPOT

Where should the build and release agents for the investment planning application suite run? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Build agent:

- A hosted service
- A source control system
- The developers' computers

Release agent:

- A hosted service
- A source control system
- The developers' computers

A. Answer:

Answer Area

Build agent:

- A hosted service
- A source control system
- The developers' computers

Release agent:

- A hosted service
- A source control system
- The developers' computers

Explanation

Explanation/Reference:

Explanation:

Box 1: A source control system

A source control system, also called a version control system, allows developers to collaborate on code and track changes. Source control is an essential tool for multi-developer projects.

Box 2: A hosted service

To build and deploy Xcode apps or Xamarin.iOS projects, you'll need at least one macOS agent. If your pipelines are in Azure Pipelines and a Microsoft-hosted agent meets your needs, you can skip setting up a self-hosted macOS agent.

Scenario: The investment planning applications suite will include one multi-tier web application and two iOS mobile applications. One mobile application will be used by employees; the other will be used by customers.

References:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/agents/v2-osx?view=azure-devops>

Question Set 2

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box 1: A source control system

A source control system, also called a version control system, allows developers to collaborate on code and track changes. Source control is an essential tool for multi-developer projects.

Box 2: A hosted service

To build and deploy Xcode apps or Xamarin.iOS projects, you'll need at least one macOS agent. If your pipelines are in Azure Pipelines and a Microsoft-hosted agent meets your needs, you can skip setting up a self-hosted macOS agent.

Scenario: The investment planning applications suite will include one multi-tier web application and two iOS mobile applications. One mobile application will be used by employees; the other will be used by customers.

References:

QUESTION 119

DRAG DROP

You have an Azure Kubernetes Service (AKS) cluster.

You need to deploy an application to the cluster by using Azure DevOps.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions	Answer Area
Create a service account in the cluster.	
Create a service principal in Azure Active Directory (Azure AD).	
Add an Azure Function App for Container task to the deployment pipeline.	
Add a Helm package and deploy a task to the deployment pipeline.	
Add a Docker Compose task to the deployment pipeline.	
Configure RBAC roles in the cluster.	

A. Answer:

Answer Area	
Actions	
Create a service account in the cluster.	Create a service principal in Azure Active Directory (Azure AD).
	Add a Helm package and deploy a task to the deployment pipeline.
Add an Azure Function App for Container task to the deployment pipeline.	Add a Docker Compose task to the deployment pipeline.
Configure RBAC roles in the cluster.	

Explanation

Explanation/Reference:

Explanation:

You can set up a CI/CD pipeline to deploy your apps on a Kubernetes cluster with Azure DevOps by leveraging a Linux agent, Docker, and Helm.

Step 1: Create a service principle in Azure Active Directory (Azure AD)

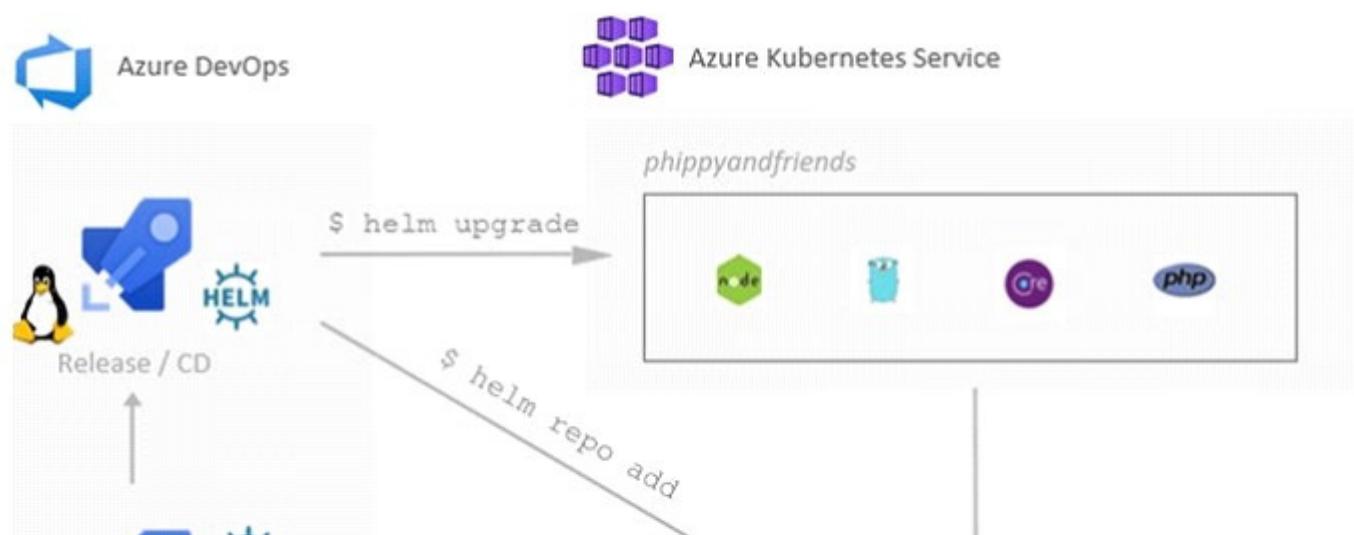
We need to assign 3 specific service principals with specific Azure Roles that need to interact with our ACR and our AKS.

Create a specific Service Principal for our Azure DevOps pipelines to be able to push and pull images and charts of our ACR.

Create a specific Service Principal for our Azure DevOps pipelines to be able to deploy our application in our AKS.

Step 2: Add a Helm package and deploy a task to the deployment pipeline

This is the DevOps workflow with containers:



Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation

You can set up a CI/CD pipeline to deploy your apps on a Kubernetes cluster with Azure DevOps by leveraging a Linux agent, Docker, and Helm.

Step 1: Create a service principle in Azure Active Directory (Azure AD)

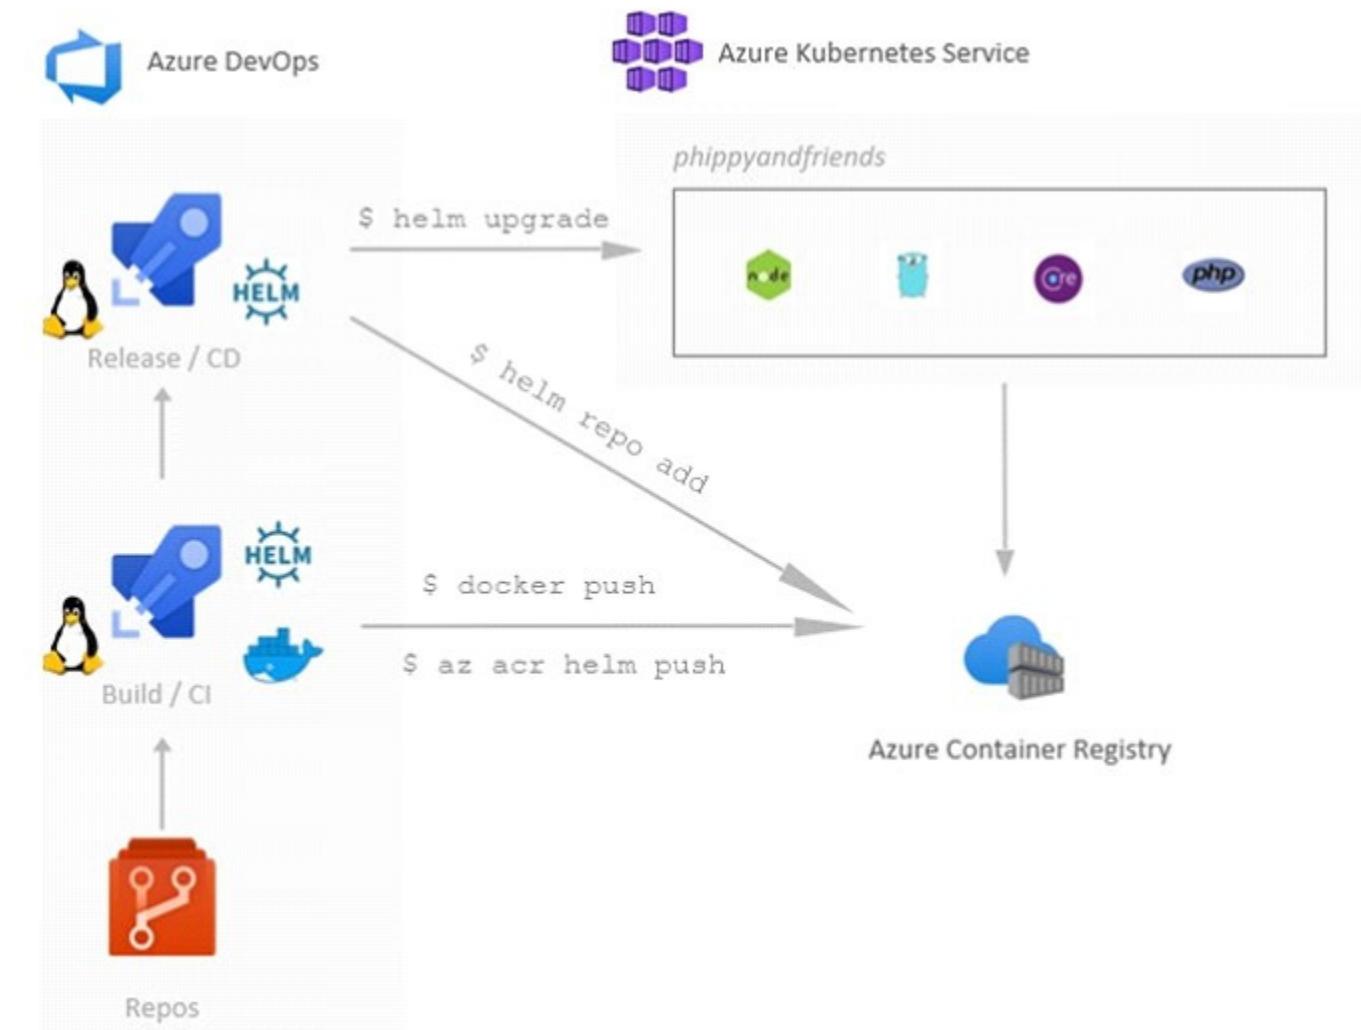
We need to assign 3 specific service principals with specific Azure Roles that need to interact with our ACR and our AKS.

Create a specific Service Principal for our Azure DevOps pipelines to be able to push and pull images and charts of our ACR.

Create a specific Service Principal for our Azure DevOps pipelines to be able to deploy our application in our AKS.

Step 2: Add a Helm package and deploy a task to the deployment pipeline

This is the DevOps workflow with containers:



Step 3: Add a Docker Compose task to the deployment pipeline.

Dockerfile file is a script leveraged by Docker, composed of various commands (instructions) and arguments listed successively to automatically perform actions on a base image in order to create a new Docker image by packaging the app.

Reference:

<https://cloudblogs.microsoft.com/opensource/2018/11/27/tutorial-azure-devops-setup-cicd-pipeline-kubernetes-docker-helm/>

QUESTION 120

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an approval process that contains a condition. The condition requires that releases be approved by a team leader before they are deployed.

You have a policy stating that approvals must occur within eight hours.

You discover that deployment fail if the approvals take longer than two hours.

You need to ensure that the deployments only fail if the approvals take longer than eight hours.

Solution: From Post-deployment conditions, you modify the Time between re-evaluation of gates option.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Use a gate From Pre-deployment conditions instead.

References: <https://docs.microsoft.com/en-us/azure/devops/pipelines/release/approvals/gates>

QUESTION 121

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an approval process that contains a condition. The condition requires that releases be approved by a team leader before they are deployed.

You have a policy stating that approvals must occur within eight hours.

You discover that deployment fail if the approvals take longer than two hours.

You need to ensure that the deployments only fail if the approvals take longer than eight hours.

Solution: From Pre-deployment conditions, you modify the Time between re-evaluation of gates option.

Does this meet the goal?

- A. Yes
- B. No

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Gates allow automatic collection of health signals from external services, and then promote the release when all the signals are successful at the same time or stop the deployment on timeout. Typically, gates are used in connection with incident management, problem management, change management, monitoring, and external approval systems.

Approvals and gates give you additional control over the start and completion of the deployment pipeline. Each stage in a release pipeline can be configured with pre-deployment and post-deployment conditions that can include waiting for users to manually approve or reject deployments, and checking with other automated systems until specific conditions are verified.

References: <https://docs.microsoft.com/en-us/azure/devops/pipelines/release/approvals/gates>

QUESTION 122

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an approval process that contains a condition. The condition requires that releases be approved by a team leader before they are deployed.

You have a policy stating that approvals must occur within eight hours.

You discover that deployment fail if the approvals take longer than two hours.

You need to ensure that the deployments only fail if the approvals take longer than eight hours.

Solution: From Pre-deployment conditions, you modify the Timeout setting for pre-deployment approvals.

Does this meet the goal?

A. Yes

B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Use a gate instead of an approval instead.

References: <https://docs.microsoft.com/en-us/azure/devops/pipelines/release/approvals/gates>

QUESTION 123

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You plan to create a release pipeline that will deploy Azure resources by using Azure Resource Manager templates. The release pipeline will create the following resources:

Two resource groups

Four Azure virtual machines in one resource group

Two Azure SQL databases in other resource group

You need to recommend a solution to deploy the resources.

Solution: Create two standalone templates, each of which will deploy the resources in its respective group.

Does this meet the goal?

A. Yes

B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Use a main template and two linked templates.

References: <https://docs.microsoft.com/en-us/azure/azure-resource-manager/resource-group-linked-templates>

QUESTION 124

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You plan to create a release pipeline that will deploy Azure resources by using Azure Resource Manager templates. The release pipeline will create the following resources:

Two resource groups

Four Azure virtual machines in one resource group

Two Azure SQL databases in other resource group

You need to recommend a solution to deploy the resources.

Solution: Create a single standalone template that will deploy all the resources.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Use two templates, one for each resource group, and link the templates.

References:

<https://docs.microsoft.com/en-us/azure/azure-resource-manager/resource-group-linked-templates>

QUESTION 125

HOTSPOT

Your company has an Azure subscription.

The company requires that all resource groups in the subscription have a tag named organization set to a value of Contoso.

You need to implement a policy to meet the tagging requirement.

How should you complete the policy? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
{  
  "policyRule": {  
    "if": {  
      "allOf": [  
        {  
          "field": "type",  
          "equals":  
            },  
        {  
          "field": "tags[\"organization\"]",  
          "equals": "Contoso"  
        }  
      ]  
    },  
    "then": {  
      "effect":  
        "details": [  
          "Append",  
          "Deny",  
          "DeployIfNotExists",  
          {  
            "field": "tags[\"organization\"]",  
            "value": "Contoso"  
          }  
        ]  
    }  
  }  
}
```

A. Answer:

Answer Area

```
{  
  "policyRule": {  
    "if": {  
      "allOf": [  
        {  
          "field": "type",  
          "equals":  
            },  
            {  
              "MicrosoftResources/deployments"  
              "MicrosoftResources/subscriptions"  
              "MicrosoftResources/subscriptions/resourceGroups"  
            }  
        ]  
      },  
      "not": {  
        "field": "tags['organization']",  
        "equals": "Contoso"  
      }  
    }  
  },  
  "then": {  
    "effect":  
      "details": [  
        "Append",  
        "Deny",  
        "DeployIfNotExists",  
        {  
          "field": "tags['organization']",  
          "value": "Contoso"  
        }  
      ]  
  }  
}
```

Explanation

Explanation/Reference:
Explanation:

Box 1: " Microsoft.Resources/subscriptions/resourceGroups"

Box 2: "Deny",
Sample - Enforce tag and its value on resource groups

```
},  
  "policyRule": {  
    "if": {  
      "allOf": [  
        {  
          "field": "type",  
          "equals":  
            },  
            {  
              "MicrosoftResources/deployments"  
              "MicrosoftResources/subscriptions"  
              "MicrosoftResources/subscriptions/resourceGroups"  
            }  
        ]  
      },  
      "not": {  
        "field": "tags['organization']",  
        "equals": "Contoso"  
      }  
    }  
  },  
  "then": {  
    "effect":  
      "details": [  
        "Append",  
        "Deny",  
        "DeployIfNotExists",  
        {  
          "field": "tags['organization']",  
          "value": "Contoso"  
        }  
      ]  
  }  
}
```

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation

Box 1: " Microsoft.Resources/subscriptions/resourceGroups"

Box 2: "Deny",
Sample - Enforce tag and its value on resource groups

```
},
"policyRule": {
  "if": {
    "allOf": [
      {
        "field": "type",
        "equals": "Microsoft.Resources/subscriptions/resourceGroups"
      },
      {
        "not": {
          "field": "[concat('tags[',parameters('tagName'), ']')]",
          "equals": "[parameters('tagValue')]"
        }
      }
    ],
    "then": {
      "effect": "deny"
    }
  }
}
```

References:

<https://docs.microsoft.com/en-us/azure/governance/policy/samples/enforce-tag-on-resource-groups>

QUESTION 126

DRAG DROP

You are defining release strategies for two applications as shown in the following table.

Application name	Goal
App1	Failure of App1 has a major impact on your company. You need a small group of users, who opted in to a testing App1, to test new releases of the application.
App2	You need to minimize the time it takes to deploy new releases of App2, and you must be able to roll back as quickly as possible.

Which release strategy should you use for each application? To answer, drag the appropriate release strategies to the correct applications. Each release strategy may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Release Strategies

Blue/Green deployment

Canary deployment

Rolling deployment

Answer Area:

App1:

App2:

A. Answer:

Release Strategies

Blue/Green deployment

[Redacted]

[Redacted]

Answer Area:

App1: Canary deployment

App2: Rolling deployment

Explanation

Explanation/Reference:

Explanation:

App1: Canary deployment

With canary deployment, you deploy a new application code in a small part of the production infrastructure. Once the application is signed off for release, only a few users are routed to it. This minimizes any impact.

With no errors reported, the new version can gradually roll out to the rest of the infrastructure.

App2: Rolling deployment:

In a rolling deployment, an application's new version gradually replaces the old one. The actual deployment happens over a period of time. During that time, new and old versions will coexist without affecting functionality or user experience. This process makes it easier to roll back any new component incompatible with the old components.

Incorrect Answers:

Blue/Green deployment

A blue/green deployment is a change management strategy for releasing software code. Blue/green deployments, which may also be referred to as A/B deployments require two identical hardware environments that are configured exactly the same way. While one environment is active and serving end users, the other environment remains idle.

Blue/green deployments are often used for consumer-facing applications and applications with critical uptime requirements. New code is released to the inactive environment, where it is thoroughly tested. Once the code has been vetted, the team makes the idle environment active, typically by adjusting a router configuration to redirect application program traffic. The process reverses when the next software iteration is ready for release.

References:

<https://dev.to/mostlyjason/intro-to-deployment-strategies-blue-green-canary-and-more-3a3>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

App1: Canary deployment

With canary deployment, you deploy a new application code in a small part of the production infrastructure. Once the application is signed off for release, only a few users are routed to it. This minimizes any impact.

With no errors reported, the new version can gradually roll out to the rest of the infrastructure.

App2: Rolling deployment:

In a rolling deployment, an application's new version gradually replaces the old one. The actual deployment happens over a period of time. During that time, new and old versions will coexist without affecting functionality or user experience. This process makes it easier to roll back any new component incompatible with the old components.

Incorrect Answers:

Blue/Green deployment

A blue/green deployment is a change management strategy for releasing software code. Blue/green deployments, which may also be referred to as A/B deployments require two identical hardware environments that are configured exactly the same way. While one environment is active and serving end users, the other environment remains idle.

Blue/green deployments are often used for consumer-facing applications and applications with critical uptime requirements. New code is released to the inactive environment, where it is thoroughly tested. Once the code has been vetted, the team makes the idle environment active, typically by adjusting a router configuration to redirect application program traffic. The process reverses when the next software iteration is ready for release.

References:

<https://dev.to/mostlyjason/intro-to-deployment-strategies-blue-green-canary-and-more-3a3>

QUESTION 127

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an approval process that contains a condition. The condition requires that releases be approved by a team leader before they are deployed.

You have a policy stating that approvals must occur within eight hours.

You discover that deployment fail if the approvals take longer than two hours.

You need to ensure that the deployments only fail if the approvals take longer than eight hours.

Solution: From Post-deployment conditions, you modify the Timeout setting for post-deployment approvals.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Use Pre-deployments conditions instead.

Use a gate instead of an approval instead.

References:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/release/approvals/gates>

QUESTION 128

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.
You have an Azure DevOps project.

Your build process creates several artifacts.

You need to deploy the artifacts to on-premises servers.

Solution: You deploy a Kubernetes cluster on-premises. You deploy a Helm agent to the cluster. You add a Download Build Artifacts task to the deployment pipeline.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Instead you should deploy an Azure self-hosted agent to an on-premises server.

Note: To build your code or deploy your software using Azure Pipelines, you need at least one agent.

If your on-premises environments do not have connectivity to a Microsoft-hosted agent pool (which is typically the case due to intermediate firewalls), you'll need to manually configure a self-hosted agent on on-premises computer(s).

Note 2: As we [Microsoft] are launching this new experience in preview, we are currently optimizing it for Azure Kubernetes Service (AKS) and Azure Container Registry (ACR). Other Kubernetes clusters, for example running on-premises or in other clouds, as well as other container registries, can be used, but require setting up a Service Account and connection manually.

References:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/agents/agents?view=azure-devops>

QUESTION 129

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an Azure DevOps project.

Your build process creates several artifacts.

You need to deploy the artifacts to on-premises servers.

Solution: You deploy a Docker build to an on-premises server. You add a Download Build Artifacts task to the deployment pipeline.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Instead you should deploy an Azure self-hosted agent to an on-premises server.

Note: To build your code or deploy your software using Azure Pipelines, you need at least one agent.

If your on-premises environments do not have connectivity to a Microsoft-hosted agent pool (which is typically the case due to intermediate firewalls), you'll need to manually configure a self-hosted agent on on-premises computer(s).

References:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/agents/agents?view=azure-devops>

QUESTION 130

This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You have an Azure DevOps project.

Your build process creates several artifacts.

You need to deploy the artifacts to on-premises servers.

Solution: You deploy an Azure self-hosted agent to an on-premises server. You add a Copy and Publish Build Artifacts task to the deployment pipeline.

Does this meet the goal?

- A. Yes
- B. No

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

To build your code or deploy your software using Azure Pipelines, you need at least one agent.

If your on-premises environments do not have connectivity to a Microsoft-hosted agent pool (which is typically the case due to intermediate firewalls), you'll need to manually configure a self-hosted agent on on-premises computer(s). The agents must have connectivity to the target on-premises environments, and access to the Internet to connect to Azure Pipelines or Team Foundation Server.

References:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/agents/agents?view=azure-devops>

QUESTION 131

Your company hosts a web application in Azure. The company uses Azure Pipelines for the build and release management of the application.

Stakeholders report that the past few releases have negatively affected system performance.

You configure alerts in Azure Monitor.

You need to ensure that new releases are only deployed to production if the releases meet defined performance baseline criteria in the staging environment first.

What should you use to prevent the deployment of releases that fail to meet the performance baseline?

- A. an Azure Scheduler job
- B. a trigger
- C. a gate
- D. an Azure function

Answer: C

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Scenarios and use cases for gates include:

Quality validation. Query metrics from tests on the build artifacts such as pass rate or code coverage and deploy only if they are within required thresholds.

Use Quality Gates to integrate monitoring into your pre-deployment or post-deployment. This ensures that you are meeting the key health/performance metrics (KPIs) as your applications move from dev to production and any differences in the infrastructure environment or scale is not negatively impacting your KPIs.

Note: Gates allow automatic collection of health signals from external services, and then promote the release when all the signals are successful at the same time or stop the deployment on timeout. Typically, gates are used in connection with incident management, problem management, change management, monitoring, and external approval systems.

References:

<https://docs.microsoft.com/en-us/azure/azure-monitor/continuous-monitoring>

<https://docs.microsoft.com/en-us/azure/devops/pipelines/release/approvals/gates?view=azure-devops>

QUESTION 132

SIMULATION

You need to ensure that an Azure web app named az400-9940427-main supports rolling upgrades. The solution must ensure that only 10 percent of users who connect to az400-9940427-main use update versions of the app.

The solution must minimize administrative effort.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

Set up staging environments in Azure App Service

1. Open Microsoft Azure Portal

2. Log into your Azure account, select your app's resource page, in the left pane, select Deployment slots > Add Slot.

my-demo-app - Deployment slots

Search (Ctrl+ /)

Save Discard Swap Refresh

+ Add Slot

You haven't added any deployment slots. Click here to get started.

Deployment Slots

Deployment slots are live apps with their own hostnames. App content and configurations elements can be swapped between two deployment slots, including the production slot.

NAME	STATUS	APP SERVICE PLAN	TRAFFIC %
my-demo-app PRODUCTION	Running	myAppServicePlan	100

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Security

Deployment

Quickstart

Deployment slots **selected**

Deployment Center

Settings

Configuration

3. In the Add a slot dialog box, give the slot a name, and select whether to clone an app configuration from another deployment slot. Select Add to continue.

Add a slot

Name:

Clone settings from:

Add Close

4. After the slot is added, select Close to close the dialog box. The new slot is now shown on the Deployment slots page. By default, Traffic % is set to 0 for the new slot, with all customer traffic routed to the production slot.

5. Select the new deployment slot to open that slot's resource page.

Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

Set up staging environments in Azure App Service

1. Open Microsoft Azure Portal

2. Log into your Azure account, select your app's resource page, in the left pane, select Deployment slots > Add Slot.

The screenshot shows the 'Deployment slots' blade for the 'my-demo-app' app service in the Azure portal. The left sidebar has 'Deployment slots' selected. The top bar includes 'Save', 'Discard', 'Swap', and 'Refresh' buttons, with 'Add Slot' highlighted by a red box. A message at the top says 'You haven't added any deployment slots. Click here to get started.' Below is a 'Deployment Slots' section with a description: 'Deployment slots are live apps with their own hostnames. App content and configurations elements can be swapped between two deployment slots, including the production slot.' A table lists the single deployment slot:

NAME	STATUS	APP SERVICE PLAN	TRAFFIC %
my-demo-app	Running	myAppServicePlan	100

3. In the Add a slot dialog box, give the slot a name, and select whether to clone an app configuration from another deployment slot. Select Add to continue.

Add a slot

Name:

Clone settings from:

4. After the slot is added, select Close to close the dialog box. The new slot is now shown on the Deployment slots page. By default, Traffic % is set to 0 for the new slot, with all customer traffic routed to the production slot.

5. Select the new deployment slot to open that slot's resource page.

my-demo-app - Deployment slots

Deployment Slots

Deployment slots are live apps with their own hostnames. App content and configurations elements can be swapped between two deployment slots, including the production slot.

NAME	STATUS	APP SERVICE PLAN	TRAFFIC %
my-demo-app	Running	myAppServicePlan	100
my-demo-app-staging	Running	myAppServicePlan	0

6. Change TRAFFIC % to 10

References:

<https://docs.microsoft.com/en-us/azure/app-service/deploy-staging-slots>

QUESTION 133

You have an Azure DevOps project named Project1 and an Azure subscription named Sub1. Sub1 contains an Azure SQL database named DB1.

You need to create a release pipeline that uses the Azure SQL Database Deployment task to update DB1.

Which artifact should you deploy?

- A. a BACPAC
- B. a DACPAC
- C. an LDF file
- D. an MDF file

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Use Azure SQL Database Deployment task in a build or release pipeline to deploy to Azure SQL DB using a DACPAC or run scripts using SQLCMD.

Reference:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/tasks/deploy/sql-azure-dacpac-deployment>

QUESTION 134

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You plan to create a release pipeline that will deploy Azure resources by using Azure Resource Manager templates. The release pipeline will create the following resources:

Two resource groups

Four Azure virtual machines in one resource group

Two Azure SQL databases in other resource group

You need to recommend a solution to deploy the resources.

Solution: Create a main template that will deploy the resources in one resource group and a nested template that will deploy the resources in the other resource group.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Use two linked templates, instead of the nested template.

References: <https://docs.microsoft.com/en-us/azure/azure-resource-manager/resource-group-linked-templates>

QUESTION 135

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You plan to create a release pipeline that will deploy Azure resources by using Azure Resource Manager templates. The release pipeline will create the following resources:

Two resource groups

Four Azure virtual machines in one resource group

Two Azure SQL databases in other resource group

You need to recommend a solution to deploy the resources.

Solution: Create a main template that has two linked templates, each of which will deploy the resources in its respective group.

Does this meet the goal?

- A. Yes
- B. No

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

To deploy your solution, you can use either a single template or a main template with many related templates. The related template can be either a separate file that is linked to from the main template, or a template that is nested within the main template.

References: <https://docs.microsoft.com/en-us/azure/azure-resource-manager/resource-group-linked-templates>

QUESTION 136

HOTSPOT

You have a project in Azure DevOps.

You plan to create a build pipeline that will deploy resources by using Azure Resource Manager templates. The templates will reference secrets stored in Azure Key Vault.

You need to ensure that you can dynamically generate the resource ID of the key vault during template deployment.

What should you include in the template? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

```
"resources": [
    {
        "apiVersion": "2018-05-01",
        "name" : "secrets",
        "type": "Microsoft.KeyVault/vaults",
        "Microsoft.Resources/deployments",
        "Microsoft.Subscription/subscriptions"
    },
    "properties": {
        "mode" : "Incremental",
        "deployment" :{
            "template"
            "templateLink"
        }
    },
    "contentVersion": "1.0.0.0",
        "uri" : "[uri(parameters('_artifactsLocation'),
        concat('./nested/sqlserver.json',
        parameters('_artifactsLocationSasToken')))]"
    },
    "parameters":{
        "secret": {
            "reference": {
                "keyVault": {
                    "id": "[resourceId(parameters('vaultSubscription'),
                    parameters('vaultResourceGroupName'),
                    'Microsoft.KeyVault/vaults',
                    parameters('vaultName'))]"
                },
                "secretName": "[parameters('secretName')]"
            }
        }
    }
],
},
```

A. Answer:

Answer Area

```
"resources": [
    {
        "apiVersion": "2018-05-01",
        "name" : "secrets",
        "type": "Microsoft.Resources/deployments",
        "dependsOn": [
            "Microsoft.KeyVault/vaults",
            "Microsoft.Resources/deployments",
            "Microsoft.Subscription/subscriptions"
        ],
        "properties": {
            "mode" : "Incremental",
            "templateLink": {
                "uri" : "[uri(parameters('_artifactsLocation'),
                concat('./nested/sqlserver.json',
                parameters('_artifactsLocationSasToken')))]"
            },
            "parameters": {
                "secret": {
                    "reference": {
                        "keyVault": {
                            "id": "[resourceId(parameters('vaultSubscription'),
                            parameters('vaultResourceGroupName'),
                            'Microsoft.KeyVault/vaults',
                            parameters('vaultName'))]"
                        },
                        "secretName": "[parameters('secretName')]"
                    }
                }
            }
        }
    }
],
```

Explanation

Explanation/Reference:

Explanation:

Box 1: "Microsoft.Resources/deployments"

Reference a secret with dynamic ID. You need to reference a key vault secret that varies based on the current deployment.

Example:

```
"resources": [
{
    "apiVersion": "2018-05-01",
    "name": "dynamicSecret",
    "type": "Microsoft.Resources/deployments",
    "properties": {
```

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation

Box 1: "Microsoft.Resources/deployments"

Reference a secret with dynamic ID. You need to reference a key vault secret that varies based on the current deployment.

Example:

```
"resources": [  
  {  
    "apiVersion": "2018-05-01",  
    "name": "dynamicSecret",  
    "type": "Microsoft.Resources/deployments",  
    "properties": {  
      "mode": "Incremental",  
      "templateLink": {
```

Box 2: "templateLink"

In your parent template, you add the linked template and pass in a parameter that contains the dynamically generated resource ID.

References: <https://docs.microsoft.com/en-us/azure/azure-resource-manager/resource-manager-keyvault-parameter>

QUESTION 137

Your company has a project in Azure DevOps for a new web application.

The company uses ServiceNow for change management.

You need to ensure that a change request is processed before any components can be deployed to the production environment.

What are two ways to integrate ServiceNow into the Azure DevOps release pipeline? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. Define a deployment control that invokes the ServiceNow REST API.
- B. Define a pre-deployment gate before the deployment to the Prod stage.
- C. Define a deployment control that invokes the ServiceNow SOAP API.
- D. Define a post-deployment gate after the deployment to the QA stage.

Answer: BD

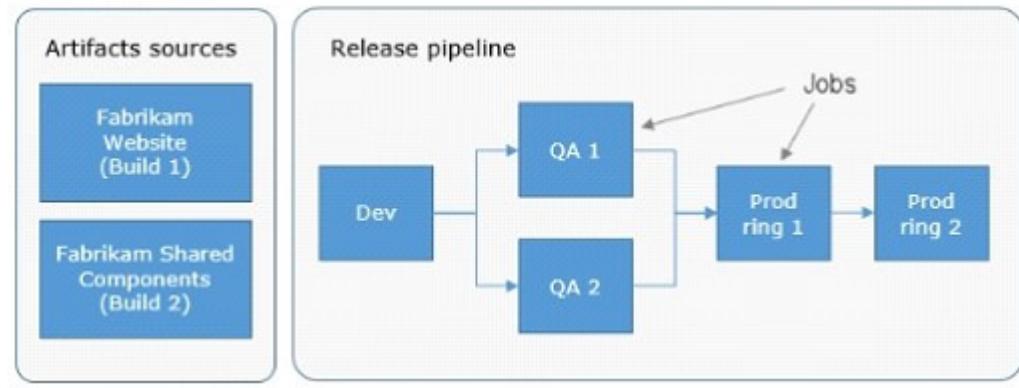
Explanation/Reference:

Explanation

Explanation/Reference:

Explanation

An example of a release pipeline that can be modeled through a release pipeline is shown below:



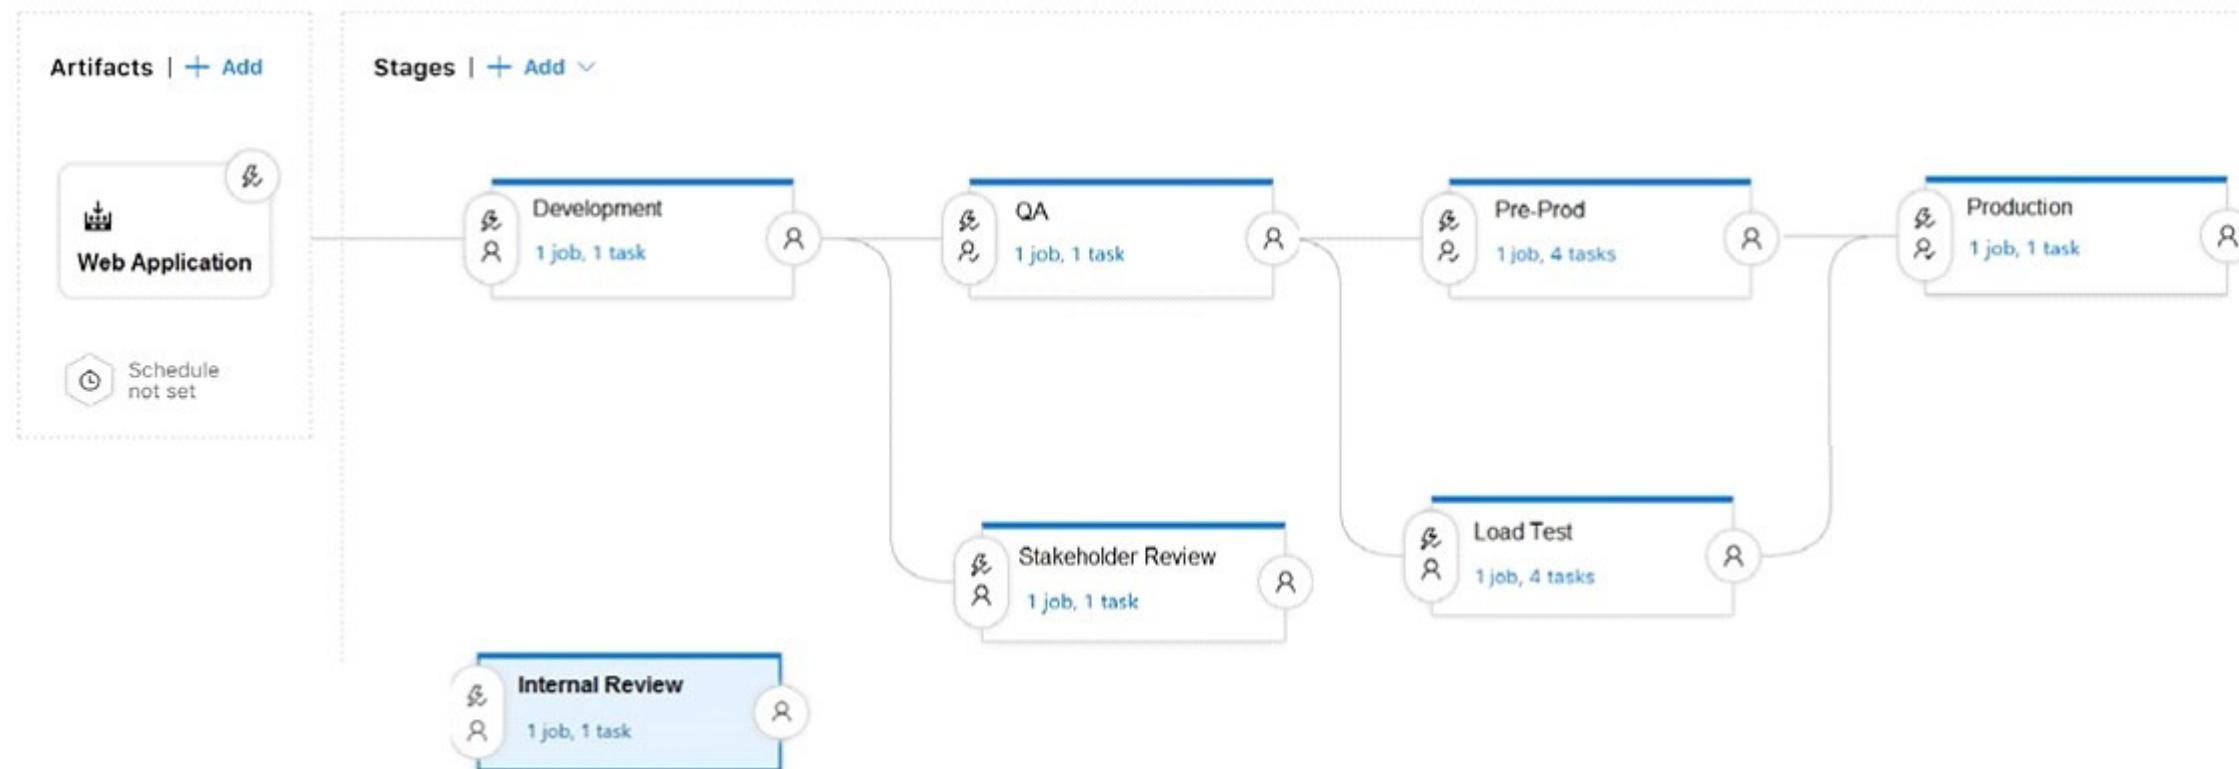
In this example, a release of a website is created by collecting specific versions of two builds (artifacts), each from a different build pipeline. The release is first deployed to a Dev stage and then forked to two QA stages in parallel. If the deployment succeeds in both the QA stages, the release is deployed to Prod ring 1 and then to Prod ring 2. Each production ring represents multiple instances of the same website deployed at various locations around the globe.

References: <https://docs.microsoft.com/en-us/azure/devops/pipelines/release>

QUESTION 138

HOTSPOT

You are configuring a release pipeline in Azure DevOps as shown in the exhibit.



Use the drop-down menus to select the answer choice that answers each question based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

How many stages have triggers set?

0
1
2
3
4
5
6
7

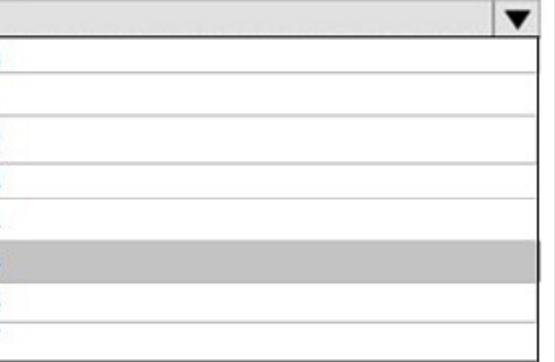
Which component should you modify to enable continuous delivery?

The Development stage
The Internal Review stage
The Production stage
The Web Application artifact

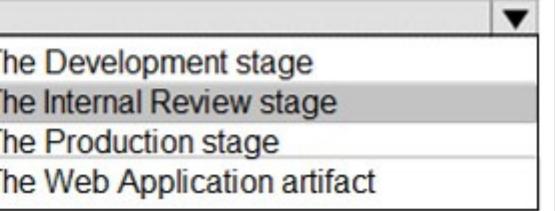
A. Answer:

Answer Area

How many stages have triggers set?



Which component should you modify to enable continuous delivery?



Explanation

Explanation/Reference:
Explanation:

Box 1: 5
There are five stages: Development, QA, Pre-production, Load Test and Production. They all have triggers.

Box 2: The Internal Review stage

References: <https://docs.microsoft.com/en-us/azure/devops/pipelines/release/triggers>

Answer:

Explanation/Reference:
Explanation

Explanation/Reference:
Explanation:

Box 1: 5
There are five stages: Development, QA, Pre-production, Load Test and Production. They all have triggers.

Box 2: The Internal Review stage

References: <https://docs.microsoft.com/en-us/azure/devops/pipelines/release/triggers>

QUESTION 139

DRAG DROP

Your company plans to deploy an application to the following endpoints:

Ten virtual machines hosted in Azure

Ten virtual machines hosted in an on-premises data center environment

All the virtual machines have the Azure Pipelines agent.

You need to implement a release strategy for deploying the application to the endpoints.

What should you recommend using to deploy the application to the endpoints? To answer, drag the appropriate components to the correct endpoints. Each component may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Components	Answer Area
<input checked="" type="checkbox"/> A deployment group	Ten virtual machines hosted in Azure: <input type="text"/>
<input type="checkbox"/> A management group	Ten virtual machines hosted in an on-premises data center environment: <input type="text"/>
<input type="checkbox"/> A resource group	
<input type="checkbox"/> Application roles	

A. Answer:

Components	Answer Area
A deployment group	
A management group	Ten virtual machines hosted in Azure: A deployment group
A resource group	Ten virtual machines hosted in an on-premises data center environment: A deployment group
Application roles	

Explanation

Explanation/Reference:

Explanation:

Box 1: A deployment group

When authoring an Azure Pipelines or TFS Release pipeline, you can specify the deployment targets for a job using a deployment group.

If the target machines are Azure VMs, you can quickly and easily prepare them by installing the Azure Pipelines Agent Azure VM extension on each of the VMs, or by using the Azure Resource Group Deployment task in your release pipeline to create a deployment group dynamically.

Box 2: A deployment group

References: <https://docs.microsoft.com/en-us/azure/devops/pipelines/release/deployment-groups>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box 1: A deployment group

When authoring an Azure Pipelines or TFS Release pipeline, you can specify the deployment targets for a job using a deployment group.

If the target machines are Azure VMs, you can quickly and easily prepare them by installing the Azure Pipelines Agent Azure VM extension on each of the VMs, or by using the Azure Resource Group Deployment task in your release pipeline to create a deployment group dynamically.

Box 2: A deployment group

References: <https://docs.microsoft.com/en-us/azure/devops/pipelines/release/deployment-groups>

QUESTION 140

You plan to use Terraform to deploy an Azure resource group.

You need to install the required frameworks to support the planned deployment.

Which two frameworks should you install? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Vault
- B. Terratest
- C. Node.js
- D. Yeoman
- E. Tiller

Answer: BD

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

You can use the combination of Terraform and Yeoman. Terraform is a tool for creating infrastructure on Azure. Yeoman makes it easy to create Terraform modules.

Terratest provides a collection of helper functions and patterns for common infrastructure testing tasks, like making HTTP requests and using SSH to access a specific virtual machine. The following list describes some of the major advantages of using Terratest:

Convenient helpers to check infrastructure - This feature is useful when you want to verify your real infrastructure in the real environment.

Organized folder structure - Your test cases are organized clearly and follow the standard Terraform module folder structure.

Test cases are written in Go - Many developers who use Terraform are Go developers. If you're a Go developer, you don't have to learn another programming language to use Terratest.

Extensible infrastructure - You can extend additional functions on top of Terratest, including Azure-specific features.

Reference:

<https://docs.microsoft.com/en-us/azure/developer/terraform/create-base-template-using-yeoman>

<https://docs.microsoft.com/en-us/azure/developer/terraform/test-modules-using-terratest>

QUESTION 141

SIMULATION

You manage a website that uses an Azure SQL Database named db1 in a resource group named RG1lod11566895.

You need to modify the SQL database to protect against SQL injection.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

Set up Advanced Threat Protection in the Azure portal

1. Sign into the Azure portal.
2. Navigate to the configuration page of the server you want to protect. In the security settings, select Advanced Data Security.
3. On the Advanced Data Security configuration page:

The screenshot shows the 'Advanced Data Security' configuration page for a SQL server named 'vanazuresqldbserver'. The left sidebar contains navigation links for Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, Settings (Quick start, Failover groups, Manage Backups, Active Directory admin, SQL databases, SQL elastic pools, Deleted databases, Import/Export history, DTU quota, Properties, Locks, Export template), and Security (Advanced Data Security). The 'Advanced Data Security' link is highlighted with a red box. The main content area has tabs for 'ADVANCED DATA SECURITY' (ON) and 'VULNERABILITY ASSESSMENT SETTINGS'. Under 'VULNERABILITY ASSESSMENT SETTINGS', there are sections for Subscription (SQL DB Content) and Storage account. Under 'ADVANCED THREAT PROTECTION SETTINGS', there are sections for Periodic recurring scans (OFF), Send scan reports to (empty input field), and an unchecked checkbox for 'Also send email notification to admins and subscription owners'. Below this is a red box highlighting the 'Send alerts to' section, which contains an input field with 'Email addresses' and a checked checkbox for 'Also send email notification to admins and subscription owners'. At the bottom, there is a section for Advanced Threat Protection types (All).

4. Enable Advanced Data Security on the server.

Note: Advanced Threat Protection for Azure SQL Database detects anomalous activities indicating unusual and potentially harmful attempts to access or exploit databases. Advanced Threat Protection can identify Potential SQL injection.

Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

Set up Advanced Threat Protection in the Azure portal

1. Sign into the Azure portal.
2. Navigate to the configuration page of the server you want to protect. In the security settings, select Advanced Data Security.
3. On the Advanced Data Security configuration page:

The screenshot shows the 'Advanced Data Security' blade for a SQL server named 'vanazuresqldbserver'. The left sidebar lists various management options like Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, and several database-related settings. The 'Advanced Data Security' option is selected and highlighted with a red box at the bottom of the sidebar.

The main content area is titled 'ADVANCED DATA SECURITY' with a status switch between 'ON' (purple) and 'OFF' (blue). Below it is the 'VULNERABILITY ASSESSMENT SETTINGS' section, which includes options for 'Subscription' (SQL DB Content) and 'Storage account'. Under 'Periodic recurring scans', the status is set to 'OFF' (blue).

The 'ADVANCED THREAT PROTECTION SETTINGS' section is highlighted with a red box. It contains fields for 'Send alerts to' (set to 'Email addresses') and a checked checkbox for 'Also send email notification to admins and subscription owners'. Below this section is an 'Advanced Threat Protection types' dropdown set to 'All'.

4. Enable Advanced Data Security on the server.

Note: Advanced Threat Protection for Azure SQL Database detects anomalous activities indicating unusual and potentially harmful attempts to access or exploit databases. Advanced Threat Protection can identify Potential SQL injection, Access from unusual location or data center, Access from unfamiliar principal or potentially harmful application, and Brute force SQL credentials

Reference:

<https://docs.microsoft.com/en-us/azure/storage/common/storage-account-create>

<https://docs.microsoft.com/en-us/azure/azure-sql/database/threat-detection-configure>

QUESTION 142**SIMULATION**

You plan to implement a CI/CD strategy for an Azure Web App named az400-11566895-main.

You need to configure a staging environment for az400-11566895-main.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

Add a slot

1. In the Azure portal, search for and select App Services and select your app az400-11566895-main.

The screenshot shows the Microsoft Azure portal homepage. The address bar at the top displays the URL <https://portal.azure.com/>. A search bar contains the text "app services". The main content area is titled "Services" and shows "All 60 results". On the left, there's a sidebar with sections for "Azure services" (Create a resource, Resource groups, Service Health), "Recent resources" (myFirstAzureWebA, WebApplicationAS, cs4316e81020662x), and "Navigate" (Subscriptions). The "App Services" item in the search results is highlighted with a red box.

2. In the left pane, select Deployment slots > Add Slot.

The screenshot shows the "Deployment slots" blade for the "my-demo-app" app service. At the top, there are buttons for "Save", "Discard", and a prominent blue "Add Slot" button, which is highlighted with a red box. Below these buttons, a message states: "You haven't added any deployment slots. Click here to get started." The left sidebar lists navigation options: Overview, Activity log, Access control (IAM), and Tags. The main content area is titled "Deployment Slots".

Answer:

Explanation/Reference:

Answer: See solution below.

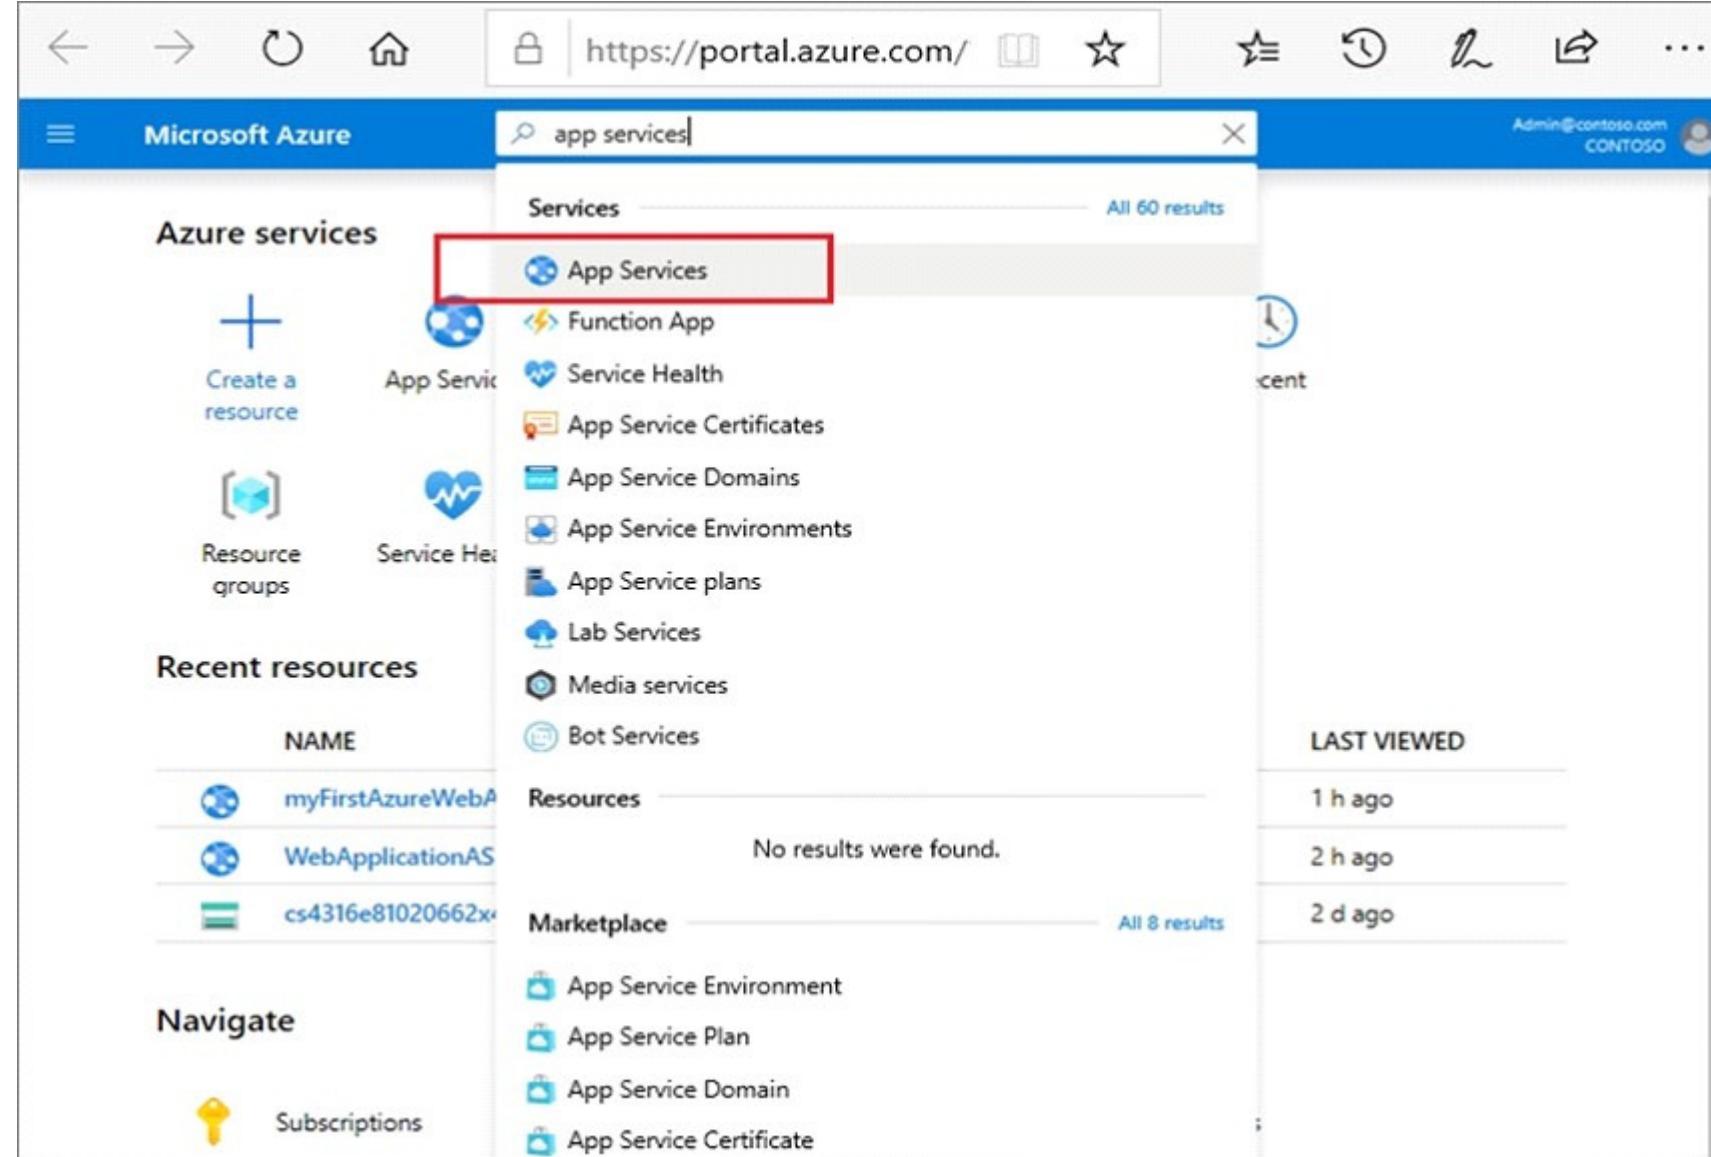
Explanation

Explanation/Reference:

Explanation:

Add a slot

1. In the Azure portal, search for and select App Services and select your app az400-11566895-main.



The screenshot shows the Microsoft Azure portal interface. The top navigation bar includes standard browser controls (back, forward, refresh, home) and a search bar containing the URL <https://portal.azure.com/>. To the right of the search bar are icons for pinned items, notifications, and user information (Admin@contoso.com, CONTOSO). The main content area has a blue header bar with the text "Microsoft Azure" and a search input field containing "app services". Below this is a sidebar with sections for "Azure services" (Create a resource, Resource groups), "Recent resources" (myFirstAzureWebA, WebApplicationAS, cs4316e81020662x), and "Navigate" (Subscriptions). The main pane displays search results under "Services" with a total of "All 60 results". The first result, "App Services", is highlighted with a red box. Other listed items include Function App, Service Health, App Service Certificates, App Service Domains, App Service Environments, App Service plans, Lab Services, Media services, Bot Services, and Marketplace items like App Service Environment, App Service Plan, App Service Domain, and App Service Certificate. On the right side of the main pane, there is a "LAST VIEWED" section showing items last viewed 1 h ago, 2 h ago, and 2 d ago.

2. In the left pane, select Deployment slots > Add Slot.

The screenshot shows the 'Deployment slots' page for an Azure App Service named 'my-demo-app'. On the left, there's a navigation sidebar with links like Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, Security, Deployment (Quickstart, Deployment slots, Deployment Center), Settings, and Configuration. The 'Deployment slots' link is highlighted with a red box. The main area has a heading 'Deployment Slots' with a sub-instruction: 'Deployment slots are live apps with their own hostnames. App content and configurations elements can be swapped between two deployment slots, including the production slot.' Below this is a table:

NAME	STATUS	APP SERVICE PLAN	TRAFFIC %
my-demo-app	Running	myAppServicePlan	100

3. In the Add a slot dialog box, give the slot a name, and select whether to clone an app configuration from another deployment slot. Select Add to continue.

The screenshot shows the 'Add a slot' dialog box. It has fields for 'Name' (containing 'staging') and 'Clone settings from:' (set to 'Do not clone settings'). At the bottom are 'Add' and 'Close' buttons, with 'Add' highlighted with a red box.

4. After the slot is added, select Close to close the dialog box. The new slot is now shown on the Deployment slots page.

my-demo-app - Deployment slots

Save Discard Add Slot Swap Refresh

Deployment Slots

Deployment slots are live apps with their own hostnames. App content and configurations elements can be swapped between two deployment slots, including the production slot.

NAME	STATUS	APP SERVICE PLAN	TRAFFIC %
my-demo-app PRODUCTION	Running	myAppServicePlan	100
my-demo-app-staging	Running	myAppServicePlan	0

Reference:
<https://docs.microsoft.com/en-us/azure/app-service/deploy-staging-slots>

QUESTION 143

SIMULATION

You have several apps that use an Azure SQL Database named db1.

You need to ensure that queries to db1 are tuned by Azure over time. The solution must only apply to db1.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

1. To enable automatic tuning on a single database, navigate to the database in the Azure portal and select Automatic tuning.

The screenshot shows the 'Automatic tuning' configuration page in the Azure portal. At the top, there's a 'Revert to defaults' button. Below it, an information icon says: 'Azure SQL Database built-in intelligence automatically tunes your databases to optimize performance. Click here to learn more about automatic tuning.' Under 'Inherit from:', the 'Server' tab is selected. A note states: 'The database is inheriting automatic tuning configuration from the server. You can set the configuration to be inherited by going to: Server tuning settings'. The main section, 'Configure the automatic tuning options', lists three options with their desired and current states:

OPTION	DESIRED STATE	CURRENT STATE
FORCE PLAN	ON OFF INHERIT	ON Inherited from server
CREATE INDEX	ON OFF INHERIT	ON Inherited from server
DROP INDEX	ON OFF INHERIT	ON Forced by user

At the bottom left is a blue 'Apply' button.

Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

1. To enable automatic tuning on a single database, navigate to the database in the Azure portal and select Automatic tuning.

Automatic tuning

 Revert to defaults



Azure SQL Database built-in intelligence automatically tunes your databases to optimize performance. Click here to learn more about automatic tuning.



Inherit from: 

Server Azure defaults Don't inherit

 The database is inheriting automatic tuning configuration from the server. You can set the configuration to be inherited by going to: [Server tuning settings](#)

Configure the automatic tuning options 

OPTION	DESIRED STATE	CURRENT STATE
 FORCE PLAN	<input type="radio"/> ON <input type="radio"/> OFF <input checked="" type="radio"/> INHERIT	ON Inherited from server
 CREATE INDEX	<input type="radio"/> ON <input type="radio"/> OFF <input checked="" type="radio"/> INHERIT	ON Inherited from server
 DROP INDEX	<input checked="" type="radio"/> ON <input type="radio"/> OFF <input type="radio"/> INHERIT	ON Forced by user

Apply

2. Select the automatic tuning options you want to enable and select Apply.

Note: Individual automatic tuning settings can be separately configured for each database. You can manually configure an individual automatic tuning option, or specify that an option inherits its settings from the server.

Reference:
<https://docs.microsoft.com/en-us/azure/azure-sql/database/automatic-tuning-enable>

QUESTION 144

Testlet 1

Case Study

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other question in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview

Litware, Inc. is an independent software vendor (ISV). Litware has a main office and five branch offices.

Existing Environment

Application Architecture

The company's primary application is a single monolithic retirement fund management system based on ASP.NET web forms that use logic written in VB.NET. Some new sections of the application are written in C#.

Variations of the application are created for individual customers. Currently, there are more than 80 live code branches in the application's code base.

The application was developed by using Microsoft Visual Studio. Source code is stored in Team Foundation Server (TFS) in the main office. The branch offices access the source code by using TFS proxy servers.

Architectural Issues

Litware focuses on writing new code for customers. No resources are provided to refactor or remove existing code. Changes to the code base take a long time, as dependencies are not obvious to individual developers.

Merge operations of the code often take months and involve many developers. Code merging frequently introduces bugs that are difficult to locate and resolve.

Customers report that ownership costs of the retirement fund management system increase continually. The need to merge unrelated code makes even minor code changes expensive.

Customers report that bug reporting is overly complex.

Requirements

Planned changes

Litware plans to develop a new suite of applications for investment planning. The investment planning applications will require only minor integration with the existing retirement fund management system.

The investment planning applications suite will include one multi-tier web application and two iOS mobile applications. One mobile application will be used by employees; the other will be used by customers.

Litware plans to move to a more agile development methodology. Shared code will be extracted into a series of packages.

Litware has started an internal cloud transformation process and plans to use cloud-based services whenever suitable.

Litware wants to become proactive in detecting failures, rather than always waiting for customer bug reports.

Technical requirements

The company's investment planning applications suite must meet the following requirements:

New incoming connections through the firewall must be minimized.

Members of a group named Developers must be able to install packages.

The principle of least privilege must be used for all permission assignments.

A branching strategy that supports developing new functionality in isolation must be used.

Members of a group named Team Leaders must be able to create new packages and edit the permissions of package feeds.

Visual Studio App Center must be used to centralize the reporting of mobile application crashes and device types in use.

By default, all releases must remain available for 30 days, except for production releases, which must be kept for 60 days.

Code quality and release quality are critical. During release, deployments must not proceed between stages if any active bugs are logged against the release.

The mobile applications must be able to call the share pricing service of the existing retirement fund management system. Until the system is upgraded, the service will only support basic authentication over HTTPS.

The required operating system configuration for the test servers changes weekly. Azure Automation State Configuration must be used to ensure that the operating system on each test server is configured the same way when the servers are created and

checked periodically.

Current Technical Issue

The test servers are configured correctly when first deployed, but they experience configuration drift over time. Azure Automation State Configuration fails to correct the configurations.

Azure Automation State Configuration nodes are registered by using the following command.

```
Register-AzureRmAutomationDscNode  
-ResourceGroupName 'TestResourceGroup'  
-AutomationAccountName 'LitwareAutomationAccount'  
-AzureVMName $vname  
-ConfigurationMode 'ApplyOnly'
```

DRAG DROP

Which package feed access levels should be assigned to the Developers and Team Leaders groups for the investment planning applications suite? To answer, drag the appropriate access levels to the correct groups. Each access level may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Access Levels	Answer Area
Collaborator	Developers: <input type="text"/>
Contributor	Team Leaders: <input type="text"/>
Owner	
Reader	

A. Answer:

Access Levels	Answer Area
Collaborator	Developers: Reader
Contributor	Team Leaders: Owner

Explanation

Explanation/Reference:

Explanation:

Box 1: Reader

Members of a group named Developers must be able to install packages.

Feeds have four levels of access: Owners, Contributors, Collaborators, and Readers. Owners can add any type of identity-individuals, teams, and groups-to any access level.

Box 2: Owner

Members of a group named Team Leaders must be able to create new packages and edit the permissions of package feeds.

Permission	Reader	Collaborator	Contributor	Owner
List and restore/install packages	✓	✓	✓	✓
Save packages from upstream sources	✓	✓	✓	
Push packages		✓	✓	
Unlist/deprecate packages		✓	✓	
Delete/unpublish package			✓	
Edit feed permissions				✓
Rename and delete feed				✓

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation

Box 1: Reader

Members of a group named Developers must be able to install packages.

Feeds have four levels of access: Owners, Contributors, Collaborators, and Readers. Owners can add any type of identity-individuals, teams, and groups-to any access level.

Box 2: Owner

Members of a group named Team Leaders must be able to create new packages and edit the permissions of package feeds.

Permission	Reader	Collaborator	Contributor	Owner
List and restore/install packages	✓	✓	✓	✓
Save packages from upstream sources	✓	✓	✓	
Push packages		✓	✓	
Unlist/deprecate packages		✓	✓	
Delete/unpublish package			✓	
Edit feed permissions			✓	
Rename and delete feed			✓	

QUESTION 145

Testlet 2

Case Study

This is a case study. Case studies are not timed separately. You can use as much exam time as you would like to complete each case. However, there may be additional case studies and sections on this exam. You must manage your time to ensure that you are able to complete all questions included on this exam in the time provided.

To answer the questions included in a case study, you will need to reference information that is provided in the case study. Case studies might contain exhibits and other resources that provide more information about the scenario that is described in the case study. Each question is independent of the other question in this case study.

At the end of this case study, a review screen will appear. This screen allows you to review your answers and to make changes before you move to the next section of the exam. After you begin a new section, you cannot return to this section.

To start the case study

To display the first question in this case study, click the Next button. Use the buttons in the left pane to explore the content of the case study before you answer the questions. Clicking these buttons displays information such as business requirements, existing environment, and problem statements. If the case study has an All Information tab, note that the information displayed is identical to the information displayed on the subsequent tabs. When you are ready to answer a question, click the Question button to return to the question.

Overview

Contoso, Ltd. is a manufacturing company that has a main office in Chicago.

Existing Environment

Contoso plans to improve its IT development and operations processes by implementing Azure DevOps principles. Contoso has an Azure subscription and creates an Azure DevOps organization.

The Azure DevOps organization includes:

The Docker extension

A deployment pool named Pool7 that contains 10 Azure virtual machines that run Windows Server 2016

The Azure subscription contains an Azure Automation account.

Requirements

Planned changes

Contoso plans to create projects in Azure DevOps as shown in the following table.

Project name	Project details
Project 1	Project1 will provide support for incremental builds and third-party SDK components
Project 2	Project2 will use an automatic build policy. A small team of developers named Team2 will work independently on changes to the project. The Team2 members will not have permissions to Project2.
Project 3	Project3 will be integrated with SonarQube
Project 4	Project4 will provide support for a build pipeline that creates a Docker image and pushes the image to the Azure Container Registry. Project4 will use an existing Dockerfile.
Project 5	Project5 will contain a Git repository in Azure Repos and a continuous integration trigger that will initiate a build in response to any change except for changes within /folder1 of the repository.
Project 6	Project6 will provide support for build and deployment pipelines. Deployment will be allowed only if the number of current work items representing active software bugs is 0.
Project 7	Project7 will contain a target deployment group named Group7 that maps to Pool7. Project7 will use Azure Automation State Configuration to maintain the desired state of the computers in Group7.

Technical requirements

Contoso identifies the following technical requirements:

Implement build agents for Project1.

Whenever possible, use Azure resources.

Avoid using deprecated technologies.

Implement a code flow strategy for Project2 that will:

- Enable Team2 to submit pull requests for Project2.
 - Enable Team2 to work independently on changes to a copy of Project2.
 - Ensure that any intermediary changes performed by Team2 on a copy of Project2 will be subject to the same restrictions as the ones defined in the build policy of Project2.
- Whenever possible, implement automation and minimize administrative effort.

Implement Project3, Project5, Project6, and Project7 based on the planned changes.

Implement Project4 and configure the project to push Docker images to Azure Container Registry.

DRAG DROP

You need to implement Project6.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions

Open the release pipeline editor.

Disable the continuous integration trigger.

Enable Gates.

Add a manual intervention task.

Open the **Triggers** tab.

Add Query Work Items.

Answer Area



A. Answer:

Actions

Disable the continuous integration trigger.

Add a manual intervention task.

Answer Area

Open the release pipeline editor.

Enable Gates.

Add Query Work Items.



Explanation

Explanation/Reference:

Explanation:

Scenario: Implement Project3, Project5, Project6, and Project7 based on the planned changes

Project 6	Project6 will provide support for build and deployment pipelines. Deployment will be allowed only if the number of current work items representing active software bugs is 0.
-----------	---

Step 1: Open the release pipeline editor.

In the Releases tab of Azure Pipelines, select your release pipeline and choose Edit to open the pipeline editor.

Step 2: Enable Gates.

Choose the pre-deployment conditions icon for the Production stage to open the conditions panel. Enable gates by using the switch control in the Gates section.

Step 3: Add Query Work items.

Choose + Add and select the Query Work Items gate.

Configure the gate by selecting an existing work item query.

The screenshot shows the 'Deployment gates' configuration page. At the top right, there's a '+ Add' button. Below it, a 'Query Work Items' gate is selected, indicated by a blue background. To the right of the gate name is a toggle switch labeled 'Enabled' which is turned on (blue). Below the gate list, there's a 'Task version' dropdown set to '0.*'.

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Scenario: Implement Project3, Project5, Project6, and Project7 based on the planned changes

Project 6	Project6 will provide support for build and deployment pipelines. Deployment will be allowed only if the number of current work items representing active software bugs is 0.
-----------	---

Step 1: Open the release pipeline editor.

In the Releases tab of Azure Pipelines, select your release pipeline and choose Edit to open the pipeline editor.

Step 2: Enable Gates.

Choose the pre-deployment conditions icon for the Production stage to open the conditions panel. Enable gates by using the switch control in the Gates section.

Step 3: Add Query Work items.

Choose + Add and select the Query Work Items gate.

Configure the gate by selecting an existing work item query.

Deployment gates ⓘ

+ Add ▾

Query Work Items

Enabled

Query Work Items ⓘ

Task version 0.* ▾

Display name *

Query Work Items

Query * ⓘ

Active Bugs ▾

Upper threshold * ⓘ

0

Advanced ^

Lower threshold * ⓘ

0

Output Variables ^

Reference name ⓘ

Variables list

There are no output variables associated with this task [more information](#)

Evaluation options ▾

The screenshot shows the 'Deployment gates' configuration for a specific task. The task is named 'Query Work Items' and is currently enabled. It is set to run at 'Task version 0.*'. The 'Display name' is also 'Query Work Items'. The 'Query' dropdown is set to 'Active Bugs'. The 'Upper threshold' is set to 0, and the 'Lower threshold' is also set to 0. Under the 'Output Variables' section, there is a reference name field which is empty. The note at the bottom states: 'Note: A case for release gate is: Incident and issues management. Ensure the required status for work items, incidents, and issues. For example, ensure deployment occurs only if no priority zero bugs exist, and validation that there are no active incidents takes place after deployment.' The 'Evaluation options' section is collapsed.

Note: A case for release gate is:

Incident and issues management. Ensure the required status for work items, incidents, and issues. For example, ensure deployment occurs only if no priority zero bugs exist, and validation that there are no active incidents takes place after deployment.

References:

<https://docs.microsoft.com/en-us/azure/devops/pipelines/release/deploy-using-approvals?view=azure-devops#configure-gate>

Question Set 3

QUESTION 146

You administer an Azure DevOps project that includes package feeds.

You need to ensure that developers can unlist and deprecate packages. The solution must use the principle of least privilege.

Which access level should you grant to the developers?

- A. Collaborator
- B. Contributor
- C. Owner

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Feeds have four levels of access: Owners, Contributors, Collaborators, and Readers. Owners can add any type of identity-individuals, teams, and groups-to any access level.

Permission	Reader	Collaborator	Contributor	Owner
List and restore/install packages	✓	✓	✓	✓
Save packages from upstream sources	✓		✓	✓
Push packages		✓	✓	
Unlist/deprecate packages		✓	✓	
Promote a package to a view			✓	✓
Delete/unpublish package				✓
Edit feed permissions				✓

Reference:

<https://docs.microsoft.com/en-us/azure/devops/artifacts/feeds/feed-permissions>

QUESTION 147

You plan to share packages that you wrote, tested, validated, and deployed by using Azure Artifacts.

You need to release multiple builds of each package by using a single feed. The solution must limit the release of packages that are in development.

What should you use?

- A. local symbols
- B. views
- C. global symbols

D. upstream sources

Answer: D

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Upstream sources enable you to manage all of your product's dependencies in a single feed. We recommend publishing all of the packages for a given product to that product's feed, and managing that product's dependencies from remote feeds in the same feed, via upstream sources. This setup has a few benefits:

Simplicity: your NuGet.config, .npmrc, or settings.xml contains exactly one feed (your feed).

Determinism: your feed resolves package requests in order, so rebuilding the same codebase at the same commit or changeset uses the same set of packages

Provenance: your feed knows the provenance of packages it saved via upstream sources, so you can verify that you're using the original package, not a custom or malicious copy published to your feed

Peace of mind: packages used via upstream sources are guaranteed to be saved in the feed on first use; if the upstream source is disabled/removed, or the remote feed goes down or deletes a package you depend on, you can continue to develop and build

References: <https://docs.microsoft.com/en-us/azure/devops/artifacts/concepts/upstream-sources?view=vsts>

QUESTION 148

Your company is concerned that when developers introduce open source libraries, it creates licensing compliance issues.

You need to add an automated process to the build pipeline to detect when common open source libraries are added to the code base.

What should you use?

- A. Microsoft Visual SourceSafe
- B. PDM
- C. WhiteSource
- D. OWASP ZAP

Answer: C

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

WhiteSource is the leader in continuous open source software security and compliance management. WhiteSource integrates into your build process, irrespective of your programming languages, build tools, or development environments. It works automatically, continuously, and silently in the background, checking the security, licensing, and quality of your open source components against WhiteSource constantly-updated definitive database of open source repositories.

Azure DevOps integration with WhiteSource Bolt will enable you to:

1. Detect and remedy vulnerable open source components.
2. Generate comprehensive open source inventory reports per project or build.
3. Enforce open source license compliance, including dependencies' licenses.

4. Identify outdated open source libraries with recommendations to update.

Note: Black duck would also be a good answer, but it is not an option here.

References: <https://www.azuredevopslabs.com/labs/vstsextend/WhiteSource/>

QUESTION 149

Your company is concerned that when developers introduce open source libraries, it creates licensing compliance issues.

You need to add an automated process to the build pipeline to detect when common open source libraries are added to the code base.

What should you use?

- A. Microsoft Visual SourceSafe
- B. Code Style
- C. Black Duck
- D. Jenkins
- E. SourceGea
- F. OWASP ZAP

Answer: C

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Secure and Manage Open Source Software

Black Duck helps organizations identify and mitigate open source security, license compliance and code-quality risks across application and container portfolios.

Black Duck Hub and its plugin for Team Foundation Server (TFS) allows you to automatically find and fix open source security vulnerabilities during the build process, so you can proactively manage risk. The integration allows you to receive alerts and fail

builds when any Black Duck Hub policy violations are met.

Note: WhiteSource would also be a good answer, but it is not an option here.

References:

<https://marketplace.visualstudio.com/items?itemName=black-duck-software.hub-tfs>

QUESTION 150

DRAG DROP

You need to find and isolate shared code. The shared code will be maintained in a series of packages.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions

Group the related components.

Assign ownership to each component group.

Create a dependency graph for the application.

Identify the most common language used.

Rewrite the components in the most common language.

Answer Area

A. Answer:

Actions

-
-
-
- Identify the most common language used.
- Rewrite the components in the most common language.

Answer Area

- Create a dependency graph for the application.
- Group the related components.
- Assign ownership to each component group.

Explanation

Explanation/Reference:
Explanation:

Step 1: Create a dependency graph for the application
By linking work items and other objects, you can track related work, dependencies, and changes made over time. All links are defined with a specific link type. For example, you can use Parent/Child links to link work items to support a hierarchical tree structure.
Whereas, the Commit and Branch link types support links between work items and commits and branches, respectively.

Step 2: Group the related components.
Packages enable you to share code across your organization: you can compose a large product, develop multiple products based on a common shared framework, or create and share reusable components and libraries.

Step 3: Assign ownership to each component graph

References:
<https://docs.microsoft.com/en-us/azure/devops/boards/queries/link-work-items-support-traceability?view=azure-devops&tabs=new-web-form>

<https://docs.microsoft.com/en-us/visualstudio/releasenotes/tfs2017-relnotes>

Answer:

Explanation/Reference:
Explanation

Explanation/Reference:
Explanation:

Step 1: Create a dependency graph for the application

By linking work items and other objects, you can track related work, dependencies, and changes made over time. All links are defined with a specific link type. For example, you can use Parent/Child links to link work items to support a hierarchical tree structure.

Whereas, the Commit and Branch link types support links between work items and commits and branches, respectively.

Step 2: Group the related components.

Packages enable you to share code across your organization: you can compose a large product, develop multiple products based on a common shared framework, or create and share reusable components and libraries.

Step 3: Assign ownership to each component graph

References:

<https://docs.microsoft.com/en-us/azure/devops/boards/queries/link-work-items-support-traceability?view=azure-devops&tabs=new-web-form>

<https://docs.microsoft.com/en-us/visualstudio/releasenotes/tfs2017-relnotes>

QUESTION 151

DRAG DROP

You are implementing a package management solution for a Node.js application by using Azure Artifacts.

You need to configure the development environment to connect to the package repository. The solution must minimize the likelihood that credentials will be leaked.

Which file should you use to configure each connection? To answer, drag the appropriate files to the correct connections. Each file may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Answer Area

Files

The .npmrc file in the project

The .npmrc file in the user's home folder

The Package.json file in the project

The Project.json file in the project

Feed registry information:

Credentials:

A. Answer:

Answer Area

Files
The Package.json file in the project
The Project.json file in the project

Feed registry information: The .npmrc file in the project

Credentials: The .npmrc file in the user's home folder

Explanation

Explanation/Reference:

Explanation:

All Azure Artifacts feeds require authentication, so you'll need to store credentials for the feed before you can install or publish packages. npm uses .npmrc configuration files to store feed URLs and credentials. Azure DevOps Services recommends using two .npmrc files.

Feed registry information: The .npmrc file in the project

One .npmrc should live at the root of your git repo adjacent to your project's package.json. It should contain a "registry" line for your feed and it should not contain credentials since it will be checked into git.

Credentials: The .npmrc file in the user's home folder

On your development machine, you will also have a .npmrc in \$home for Linux or Mac systems or \$env.HOME for win systems. This .npmrc should contain credentials for all of the registries that you need to connect to. The NPM client will look at your project's .npmrc, discover the registry, and fetch matching credentials from \$home/.npmrc or \$env.HOME/.npmrc.

References:

<https://docs.microsoft.com/en-us/azure/devops/artifacts/npm/npmrc?view=azure-devops&tabs=windows>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

All Azure Artifacts feeds require authentication, so you'll need to store credentials for the feed before you can install or publish packages. npm uses .npmrc configuration files to store feed URLs and credentials. Azure DevOps Services recommends using two .npmrc files.

Feed registry information: The .npmrc file in the project

One .npmrc should live at the root of your git repo adjacent to your project's package.json. It should contain a "registry" line for your feed and it should not contain credentials since it will be checked into git.

Credentials: The .npmrc file in the user's home folder

On your development machine, you will also have a .npmrc in \$home for Linux or Mac systems or \$env.HOME for win systems. This .npmrc should contain credentials for all of the registries that you need to connect to. The NPM client will look at your project's .npmrc, discover the registry, and fetch matching credentials from \$home/.npmrc or \$env.HOME/.npmrc.

References:

<https://docs.microsoft.com/en-us/azure/devops/artifacts/npm/npmrc?view=azure-devops&tabs=windows>

QUESTION 152

DRAG DROP

You are creating a NuGet package.

You plan to distribute the package to your development team privately.
You need to share the package and test that the package can be consumed.

Which four actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions

Answer Area

Create a new Azure Artifacts feed.

Configure a self-hosted agent.

Publish a package.

Install a package.

Connect to an Azure Artifacts feed.



A. Answer:

Actions

Install a package.

Answer Area

Configure a self-hosted agent.

Create a new Azure Artifacts feed.

Publish a package.

Connect to an Azure Artifacts feed.



Explanation

Explanation/Reference:

Explanation:

Step 1: Configure a self-hosted agent.

The build will run on a Microsoft hosted agent.

Step 2: Create a new Azure Artifacts feed

Microsoft offers an official extension for publishing and managing your private NuGet feeds.

Step 3: Publish the package.

Publish, pack and push the built project to your NuGet feed.

Step 4: Connect to an Azure Artifacts feed.

With the package now available, you can point Visual Studio to the feed, and download the newly published package

References:

<https://medium.com/@dan.cokely/creating-nuget-packages-in-azure-devops-with-azure-pipelines-and-yaml-d6fa30f0f15e>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Step 1: Configure a self-hosted agent.
The build will run on a Microsoft hosted agent.

Step 2: Create a new Azure Artifacts feed
Microsoft offers an official extension for publishing and managing your private NuGet feeds.

Step 3: Publish the package.
Publish, pack and push the built project to your NuGet feed.

Step 4: Connect to an Azure Artifacts feed.
With the package now available, you can point Visual Studio to the feed, and download the newly published package

References:
<https://medium.com/@dan.cokely/creating-nuget-packages-in-azure-devops-with-azure-pipelines-and-yaml-d6fa30f0f15e>

QUESTION 153

HOTSPOT

You have an Azure DevOps project that contains a build pipeline. The build pipeline uses approximately 50 open source libraries.

You need to ensure that the project can be scanned for known security vulnerabilities in the open source libraries.

What should you do? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Object to create:

- A build task
- A deployment task
- An artifacts repository

Service to use:

- WhiteSource Bolt
- Bamboo
- CMake
- Chef

A. Answer:

Answer Area

Object to create:

- A build task
- A deployment task
- An artifacts repository

Service to use:

- WhiteSource Bolt
- Bamboo
- CMake
- Chef

Explanation

Explanation/Reference:
Explanation:

Box 1: A Build task
Trigger a build

You have a Java code provisioned by the Azure DevOps demo generator. You will use WhiteSource Bolt extension to check the vulnerable components present in this code.

1. Go to Builds section under Pipelines tab, select the build definition WhiteSourceBolt and click on Queue to trigger a build.
2. To view the build in progress status, click on ellipsis and select View build results.

Box 2: WhiteSource Bolt

WhiteSource is the leader in continuous open source software security and compliance management. WhiteSource integrates into your build process, irrespective of your programming languages, build tools, or development environments. It works automatically, continuously, and silently in the background, checking the security, licensing, and quality of your open source components against WhiteSource constantly-updated definitive database of open source repositories.

References:

<https://www.azuredevopslabs.com/labs/vstsextend/whitesource/>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box 1: A Build task

Trigger a build

You have a Java code provisioned by the Azure DevOps demo generator. You will use WhiteSource Bolt extension to check the vulnerable components present in this code.

1. Go to Builds section under Pipelines tab, select the build definition WhiteSourceBolt and click on Queue to trigger a build.

2. To view the build in progress status, click on ellipsis and select View build results.

Box 2: WhiteSource Bolt

WhiteSource is the leader in continuous open source software security and compliance management. WhiteSource integrates into your build process, irrespective of your programming languages, build tools, or development environments. It works automatically, continuously, and silently in the background, checking the security, licensing, and quality of your open source components against WhiteSource constantly-updated definitive database of open source repositories.

References:

<https://www.azuredevopslabs.com/labs/vstsextend/whitesource/>

QUESTION 154

SIMULATION

You plan to store signed images in an Azure Container Registry instance named az4009940427acr1.

You need to modify the SKU for az4009940427acr1 to support the planned images. The solution must minimize costs.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

1. Open Microsoft Azure Portal, and select the Azure Container Registry instance named az4009940427acr1.

2. Under Policies, select Content Trust > Enabled > Save.

The screenshot shows the 'myregistry - Content Trust' settings page in the Azure portal. The left sidebar has sections for Services (Repositories, Webhooks, Replications, Tasks) and Policies (Content trust, which is highlighted with a red border). The main area shows a status message: 'When turned on, content trust enables you to push trusted images to the registry.' Below it is a 'Status' section with two buttons: 'Disabled' and 'Enabled'. The 'Enabled' button is highlighted with a red border. At the top right are 'Save' and 'Discard' buttons.

References:

<https://docs.microsoft.com/en-us/azure/container-registry/container-registry-content-trust>

Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

1. Open Microsoft Azure Portal, and select the Azure Container Registry instance named az4009940427acr1.

2. Under Policies, select Content Trust > Enabled > Save.

The screenshot shows the 'Content trust' settings for a container registry named 'myregistry'. The status is set to 'Enabled'. A note explains that enabling content trust allows pushing trusted images to the registry. The left sidebar lists other management options: Services (Repositories, Webhooks, Replications, Tasks), Policies (Content trust, highlighted with a red border), and Monitoring.

References:

<https://docs.microsoft.com/en-us/azure/container-registry/container-registry-content-trust>

QUESTION 155

You have an Azure DevOps project that contains a build pipeline. The build pipeline uses approximately 50 open source libraries.

You need to ensure that all the open source libraries comply with your company's licensing standards.

Which service should you use?

- A. NuGet
- B. Maven
- C. Black Duck
- D. Helm

Answer: C

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Secure and Manage Open Source Software

Black Duck helps organizations identify and mitigate open source security, license compliance and code-quality risks across application and container portfolios.

Black Duck Hub and its plugin for Team Foundation Server (TFS) allows you to automatically find and fix open source security vulnerabilities during the build process, so you can proactively manage risk. The integration allows you to receive alerts and fail

builds when any Black Duck Hub policy violations are met.

Note: WhiteSource would also be a good answer, but it is not an option here.

Reference:

<https://marketplace.visualstudio.com/items?itemName=black-duck-software.hub-tfs>

Question Set 1

QUESTION 156

SIMULATION

You plan to add a new web farm that will be published by using an IP address of 10.0.0.5.

You need to allow traffic from the web farm to an Azure Database for MySQL server named az400-11566895-mysql.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

Server-level firewall rules can be used to manage access to an Azure Database for MySQL Server from a specified IP address or a range of IP addresses.

Create a server-level firewall rule in the Azure portal

1. On the MySQL server page, under Settings heading, click Connection Security to open the Connection Security page for the Azure Database for MySQL.

The screenshot shows the 'mydemoserver - Connection security' page in the Azure portal. The left sidebar has a red box around the 'Connection security' item under the 'Settings' section. The main area shows the 'Firewall rules' section with a note: 'Connections from the IPs specified below provides access to all the databases in andrela-mysql.' Below it is a table for defining firewall rules, with columns for Rule name, Start IP address, and End IP address. A note says 'No firewall rules configured.' The 'VNET rules' section shows a table with columns: Rule name, Virtual network, Subnet, Address range, Endpoint status, Resource group, Subscription ID, and State. A note says 'No results.' The 'SSL settings' section shows a note: 'Enforcing SSL connections on your server may require additional configuration to your applications connecting to the server. Click here to learn more.' Below it is a switch for 'Enforce SSL connection' which is set to 'ENABLED'.

2. In the firewall rules for the Azure Database for MySQL, you can specify a single IP address or a range of addresses. If you want to limit the rule to a single IP address, type the same address in the Start IP and End IP fields. Opening the firewall enables administrators, users, and application to access any database on the MySQL server to which they have valid credentials.

Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation

Server-level firewall rules can be used to manage access to an Azure Database for MySQL Server from a specified IP address or a range of IP addresses.

Create a server-level firewall rule in the Azure portal

1. On the MySQL server page, under Settings heading, click Connection Security to open the Connection Security page for the Azure Database for MySQL.

The screenshot shows the 'mydemoserver - Connection security' page in the Azure portal. The left sidebar has a red box around the 'Connection security' tab. The main content area has a note: 'Connections from the IPs specified below provides access to all the databases in andrela-mysql.' There is an 'ON/OFF' switch for 'Allow access to Azure services'. A table for 'Firewall rules' is shown with columns: Rule name, Start IP address, and End IP address. Below it, a note says 'No firewall rules configured.' The 'VNET rules' section has a '+ Adding existing virtual network' and '+ Create new virtual network' button. The 'SSL settings' section shows 'Enforce SSL connection' set to 'ENABLED'. The left sidebar also lists 'Connection strings', 'Server parameters', 'Replication', 'Pricing tier', 'Properties', 'Locks', 'Export template', and 'Security' (which includes 'Advanced Threat Protection ...').

2. In the firewall rules for the Azure Database for MySQL, you can specify a single IP address or a range of addresses. If you want to limit the rule to a single IP address, type the same address in the Start IP and End IP fields. Opening the firewall enables administrators, users, and application to access any database on the MySQL server to which they have valid credentials.

The screenshot shows the 'mydemoserver - Connection security' page in the Azure portal. The left sidebar lists various settings like Overview, Activity log, and Connection strings. The main area is titled 'Firewall rules' and contains a table with columns: Rule name, Start IP address, and End IP address. A new rule is being added, with the 'Rule name' field empty and the 'Start IP address' and 'End IP address' both set to '123.123.123.123'. The 'End IP address' field is highlighted with a red box. Below the table, there's a section for VNET rules and an SSL settings section where 'Enforce SSL connection' is set to 'ENABLED'.

Rule name	Start IP address	End IP address
ClientIPAddress_2019-9-4_13-47-46	123.123.123.123	123.123.123.123
RangeOfAddresses	123.123.123.0	123.123.123.255

3. Click Save on the toolbar to save this server-level firewall rule. Wait for the confirmation that the update to the firewall rules is successful.

Reference:

<https://docs.microsoft.com/en-us/azure/mysql/howto-manage-firewall-using-portal#create-a-server-level-firewall-rule-in-the-azure-portal>

QUESTION 157

DRAG DROP

You plan to use Azure Kubernetes Service (AKS) to host containers deployed from images hosted in a Docker Trusted Registry.

You need to recommend a solution for provisioning and connecting to AKS. The solution must ensure that AKS is RBAC-enabled and uses a custom service principal.

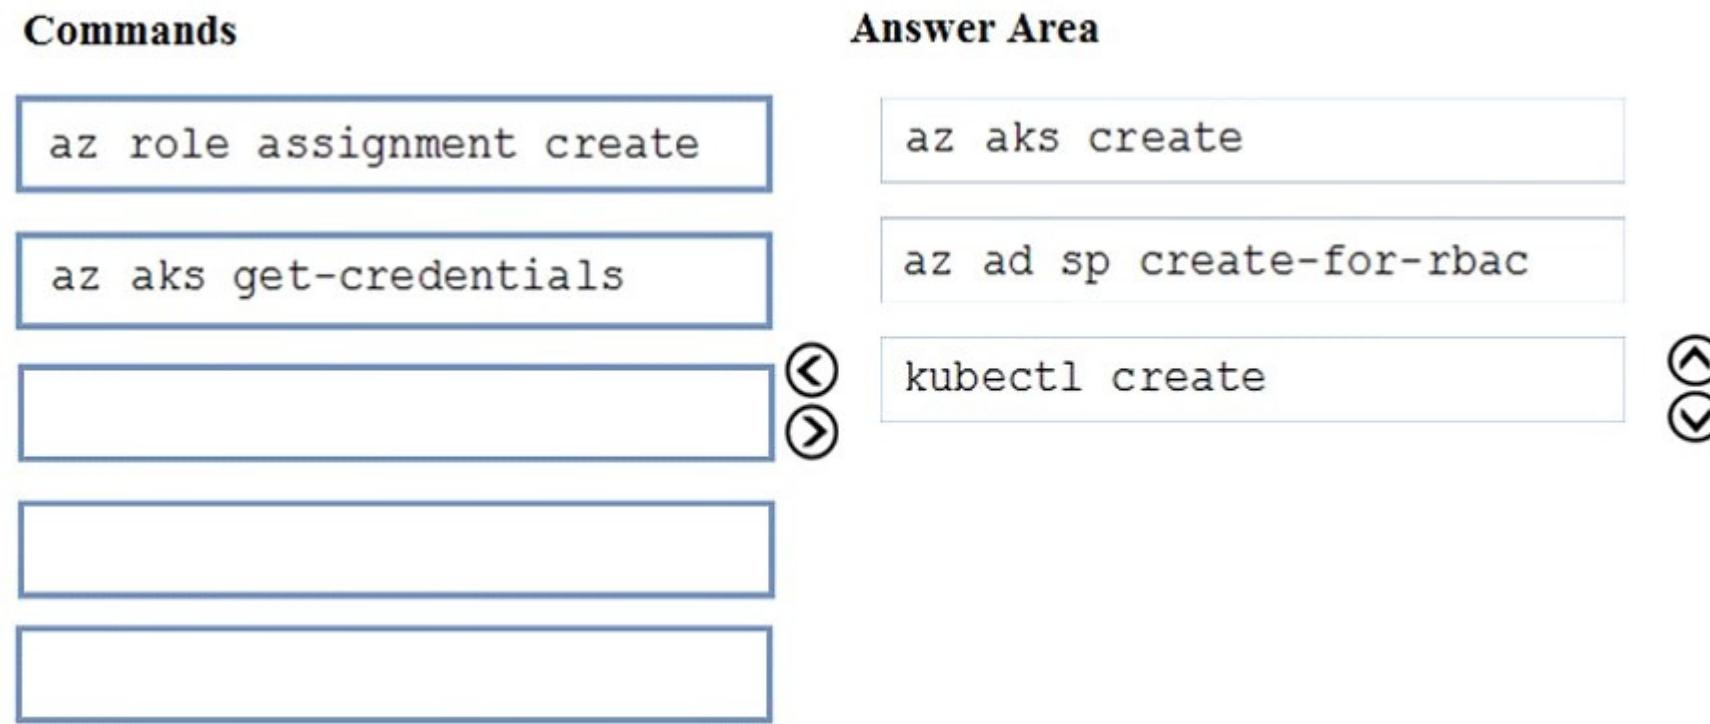
Which three commands should you recommend be run in sequence? To answer, move the appropriate commands from the list of commands to the answer area and arrange them in the correct order.

Select and Place:

Commands	Answer Area
az role assignment create	
az aks get-credentials	
az aks create	◀ ▶
az ad sp create-for-rbac	
kubectl create	



A. Answer:



Explanation

Explanation/Reference:
Explanation:

Step 1 : az acr create

An Azure Container Registry (ACR) can also be created using the new Azure CLI.

```
az acr create  
--name <REGISTRY_NAME>  
--resource-group <RESOURCE_GROUP_NAME>  
--sku Basic
```

Step 2: az ad sp create-for-rbac

Once the ACR has been provisioned, you can either enable administrative access (which is okay for testing) or you create a Service Principal (sp) which will provide a client_id and a client_secret.

```
az ad sp create-for-rbac  
--scopes /subscriptions/<SUBSCRIPTION_ID>/resourcegroups/<RG_NAME>/providers/Microsoft.ContainerRegistry/registries/<REGISTRY_NAME>  
--role Contributor  
--name <SERVICE_PRINCIPAL_NAME>
```

Step 3: kubectl create

Create a new Kubernetes Secret.

```
kubectl create secret docker-registry <SECRET_NAME>  
--docker-server <REGISTRY_NAME>.azurecr.io  
--docker-email <YOUR_MAIL>  
--docker-username=<SERVICE_PRINCIPAL_ID>  
--docker-password <YOUR_PASSWORD>
```

References:

<https://thorsten-hans.com/how-to-use-private-azure-container-registry-with-kubernetes>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation

Step 1 : az acr create

An Azure Container Registry (ACR) can also be created using the new Azure CLI.

```
az acr create  
--name <REGISTRY_NAME>  
--resource-group <RESOURCE_GROUP_NAME>  
--sku Basic
```

Step 2: az ad sp create-for-rbac

Once the ACR has been provisioned, you can either enable administrative access (which is okay for testing) or you create a Service Principal (sp) which will provide a client_id and a client_secret.

```
az ad sp create-for-rbac  
--scopes /subscriptions/<SUBSCRIPTION_ID>/resourcegroups/<RG_NAME>/providers/Microsoft.ContainerRegistry/registries/<REGISTRY_NAME>  
--role Contributor  
--name <SERVICE_PRINCIPAL_NAME>
```

Step 3: kubectl create

Create a new Kubernetes Secret.

```
kubectl create secret docker-registry <SECRET_NAME>  
--docker-server <REGISTRY_NAME>.azurecr.io  
--docker-email <YOUR_MAIL>  
--docker-username=<SERVICE_PRINCIPAL_ID>  
--docker-password <YOUR_PASSWORD>
```

References:

<https://thorsten-hans.com/how-to-use-private-azure-container-registry-with-kubernetes>

QUESTION 158

Your company has a project in Azure DevOps for a new application. The application will be deployed to several Azure virtual machines that run Windows Server 2016.

You need to recommend a deployment strategy for the virtual machines. The strategy must meet the following requirements:

Ensure that the virtual machines maintain a consistent configuration.

Minimize administrative effort to configure the virtual machines.

What should you include in the recommendation?

- A. Azure Resource Manager templates and the PowerShell Desired State Configuration (DSC) extension for Windows
- B. Deployment YAML and Azure pipeline deployment groups
- C. Azure Resource Manager templates and the Custom Script Extension for Windows

D. Deployment YAML and Azure pipeline stage templates

Answer: C

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

The Custom Script Extension downloads and executes scripts on Azure virtual machines. This extension is useful for post deployment configuration, software installation, or any other configuration or management tasks. Scripts can be downloaded from Azure storage or GitHub, or provided to the Azure portal at extension run time. The Custom Script Extension integrates with Azure Resource Manager templates, and can be run using the Azure CLI, PowerShell, Azure portal, or the Azure Virtual Machine REST API.

Incorrect Answers:

B: YAML doesn't work with Azure pipeline deployment groups.

References:

<https://docs.microsoft.com/en-us/azure/virtual-machines/extensions/custom-script-windows>

QUESTION 159

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You manage a project in Azure DevOps.

You need to prevent the configuration of the project from changing over time.

Solution: Add a code coverage step to the build pipelines.

Does this meet the goal?

A. Yes

B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Instead implement Continuous Assurance for the project.

Reference:

<https://azsk.azurewebsites.net/04-Continous-Assurance/Readme.html>

QUESTION 160

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You manage a project in Azure DevOps.

You need to prevent the configuration of the project from changing over time.

Solution: Implement Continuous Integration for the project.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Instead implement Continuous Assurance for the project.

Reference:

<https://azsk.azurewebsites.net/04-Continous-Assurance/Readme.html>

QUESTION 161

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You manage a project in Azure DevOps.

You need to prevent the configuration of the project from changing over time.

Solution: Implement Continuous Assurance for the project.

Does this meet the goal?

- A. Yes
- B. No

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

The basic idea behind Continuous Assurance (CA) is to setup the ability to check for "drift" from what is considered a secure snapshot of a system. Support for Continuous Assurance lets us treat security truly as a 'state' as opposed to a 'point in time' achievement. This is particularly important in today's context when 'continuous change' has become a norm.

There can be two types of drift:

Drift involving 'baseline' configuration: This involves settings that have a fixed number of possible states (often pre-defined/statically determined ones). For instance, a SQL DB can have TDE encryption turned ON or OFF...or a Storage Account may have

auditing turned ON however the log retention period may be less than 365 days.

Drift involving 'stateful' configuration: There are settings which cannot be constrained within a finite set of well-known states. For instance, the IP addresses configured to have access to a SQL DB can be any (arbitrary) set of IP addresses. In such scenarios,

usually human judgment is initially required to determine whether a particular configuration should be considered 'secure' or not. However, once that is done, it is important to ensure that there is no "stateful drift" from the attested configuration. (E.g., if, in a troubleshooting session, someone adds the IP address of a developer machine to the list, the Continuous Assurance feature should be able to identify the drift and generate notifications/alerts or even trigger 'auto-remediation' depending on the severity of the change).

Reference:
<https://azsk.azurewebsites.net/04-Continous-Assurance/Readme.html>

QUESTION 162

Your company has a release pipeline in an Azure DevOps project.

You plan to deploy to an Azure Kubernetes Services (AKS) cluster by using the Helm package and deploy task.

You need to install a service in the AKS namespace for the planned deployment.

Which service should you install?

- A. Azure Container Registry
- B. Chart
- C. Kubectl
- D. Tiller

Answer: D

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Before you can deploy Helm in an RBAC-enabled AKS cluster, you need a service account and role binding for the Tiller service.

Incorrect Answers:

C: Kubectl is a command line interface for running commands against Kubernetes clusters.

Reference:

<https://docs.microsoft.com/en-us/azure/aks/kubernetes-helm>

QUESTION 163

Your company develops an app for iOS. All users of the app have devices that are members of a private distribution group in Microsoft Visual Studio App Center.

You plan to distribute a new release of the app.

You need to identify which certificate file you require to distribute the new release from App Center.

Which file type should you upload to App Center?

- A. .cer

- B. .pfx
- C. .p12
- D. .pvk

Answer: C

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

A successful IOS device build will produce an ipa file. In order to install the build on a device, it needs to be signed with a valid provisioning profile and certificate. To sign the builds produced from a branch, enable code signing in the configuration pane and upload a provisioning profile (.mobileprovision) and a valid certificate (.p12), along with the password for the certificate.

References:

<https://docs.microsoft.com/en-us/appcenter/build/xamarin/ios/>

QUESTION 164

SIMULATION

You need to create a virtual machine template in an Azure DevTest Labs environment named az400-9940427-dtl1. The template must be based on Windows Server 2016 Datacenter. Virtual machines created from the template must include the selenium tool and the Google Chrome browser.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

1. Open Microsoft Azure Portal

2. Select All Services, and then select DevTest Labs in the DEVOPS section.

The screenshot shows the Microsoft Azure portal's 'All services' page. On the left, there's a sidebar with various navigation items like Home, Dashboard, and DevTest Labs. The 'All services' item is highlighted with a red box and has a yellow circle with the number '1' above it. In the main list, under the 'DEVOPS' category, the 'DevTest Labs' service is also highlighted with a red box and has a yellow circle with the number '2' to its right. Additionally, the 'DevTest Labs' entry in the list itself has a yellow circle with the number '3' to its left, pointing to it from the sidebar.

3. From the list of labs, select the az400-9940427-dtl1 lab

4. On the home page for your lab, select + Add on the toolbar.

5. Select the Windows Server 2016 Datacenter base image for the VM.

6. Select automation options at the bottom of the page above the Submit button.

7. You see the Azure Resource Manager template for creating the virtual machine.

8. The JSON segment in the resources section has the definition for the image type you selected earlier.

Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

1. Open Microsoft Azure Portal

2. Select All Services, and then select DevTest Labs in the DEVOPS section.

The screenshot shows the Microsoft Azure portal interface. On the left, there's a sidebar with various navigation options like Home, Dashboard, and All services. The 'All services' option is highlighted with a red box and has a yellow callout with the number 1 pointing to it. In the main content area, there are several sections: SECURITY (7), DEVOPS (5), and MIGRATE (6). The DEVOPS section is expanded, showing DevOps Projects, Application Insights, Lab Accounts, API Management services, and DevTest Labs. 'DevTest Labs' is highlighted with a yellow box and has a red star icon with the number 2. In the bottom right corner of the 'DevTest Labs' card, there's another yellow box with the number 3. The top of the screen features a search bar and various global navigation icons.

3. From the list of labs, select the az400-9940427-dtl1 lab

4. On the home page for your lab, select + Add on the toolbar.

5. Select the Windows Server 2016 Datacenter base image for the VM.

6. Select automation options at the bottom of the page above the Submit button.
7. You see the Azure Resource Manager template for creating the virtual machine.
8. The JSON segment in the resources section has the definition for the image type you selected earlier.

References:

<https://docs.microsoft.com/bs-cyrl-ba/azure//lab-services/devtest-lab-vm-powershell>

QUESTION 165

SIMULATION

You need to prepare a network security group (NSG) named az400-9940427-nsg1 to host an Azure DevOps pipeline agent. The solution must allow only the required outbound port for Azure DevOps and deny all other inbound and outbound access to the Internet.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

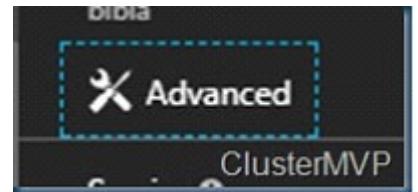
Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

1. Open Microsoft Azure Portal and Log into your Azure account.
2. Select network security group (NSG) named az400-9940427-nsg1
3. Select Settings, Outbound security rules, and click Add
4. Click Advanced



5. Change the following settings:
Destination Port range: 8080

Protocol. TCP

Action: Allow

Note: By default, Azure DevOps Server uses TCP Port 8080.

References:

<https://robertsmit.wordpress.com/2017/09/11/step-by-step-azure-network-security-groups-nsg-security-center-azure-nsg-network/>

<https://docs.microsoft.com/en-us/azure/devops/server/architecture/required-ports?view=azure-devops>

Answer:

Explanation/Reference:

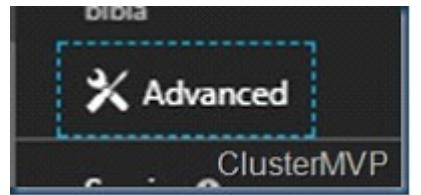
Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

1. Open Microsoft Azure Portal and Log into your Azure account.
2. Select network security group (NSG) named az400-9940427-nsg1
3. Select Settings, Outbound security rules, and click Add
4. Click Advanced



5. Change the following settings:
Destination Port range: 8080

Protocol: TCP

Action: Allow

Note: By default, Azure DevOps Server uses TCP Port 8080.

References:

<https://robertsmitt.wordpress.com/2017/09/11/step-by-step-azure-network-security-groups-nsg-security-center-azure-nsg-network/>

<https://docs.microsoft.com/en-us/azure/devops/server/architecture/required-ports?view=azure-devops>

QUESTION 166

SIMULATION

You plan to deploy a template named D:\Deploy.json to a resource group named Deploy-1d9940427.

You need to modify the template to meet the following requirements, and then to deploy the template:

The address space must be reduced to support only 256 total IP addresses.

The subnet address space must be reduced to support only 64 total IP addresses.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

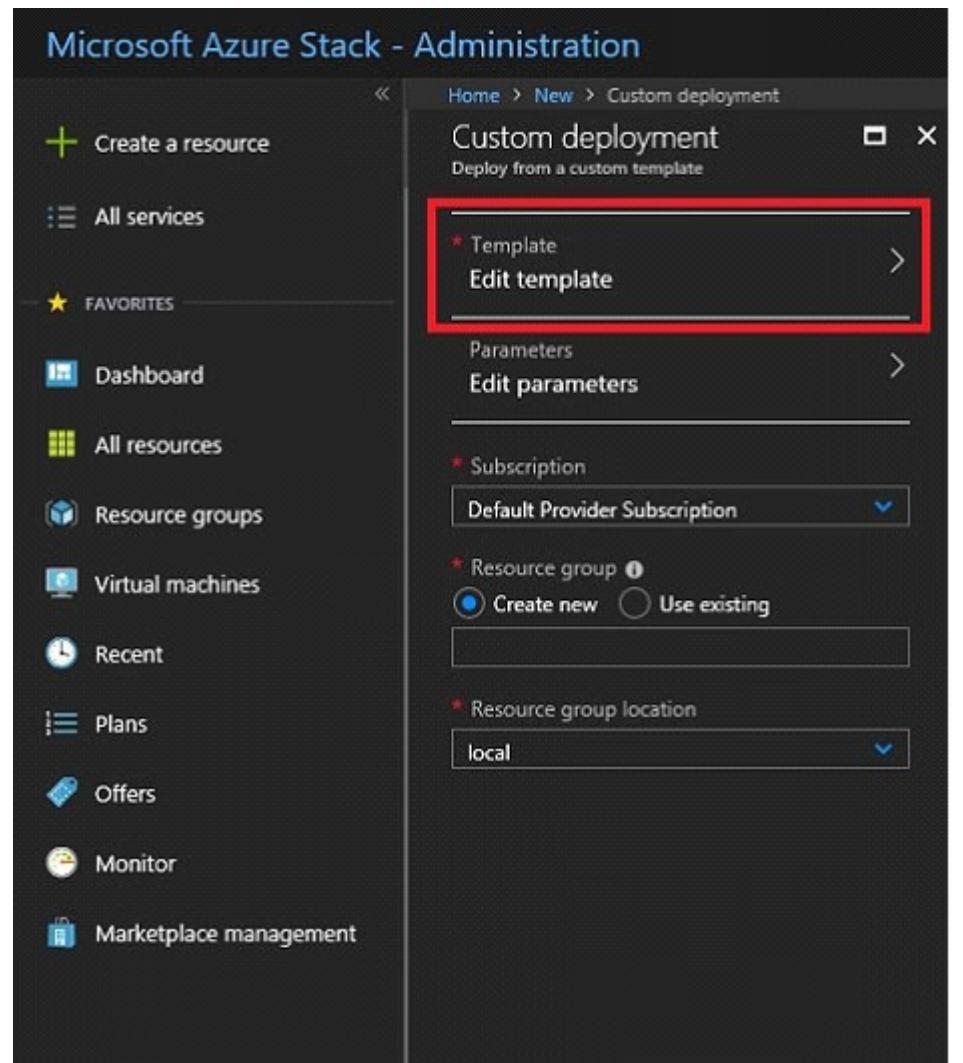
Explanation/Reference:

Explanation:

1. Sign in to the portal,
2. Choose template Deploy-1od9940427
3. Select Edit template, and then paste your JSON template code into the code window.
4. Change the AddressPrefixes to 10.0.0.0/24 in order to support only 256 total IP addresses.
addressSpace": {"addressPrefixes": ["10.0.0.0/24"]},
5. Change the firstSubnet addressprefix to 10.0.0.0/26 to support only 64 total IP addresses.

```
"subnets": [  
{  
    "name": "firstSubnet",  
    "properties": {  
        "addressPrefix": "10.0.0.0/24"  
    }  
}]
```

6. Select Save.



Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation

1. Sign in to the portal,

2. Choose template Deploy-lod9940427

3. Select Edit template, and then paste your JSON template code into the code window.

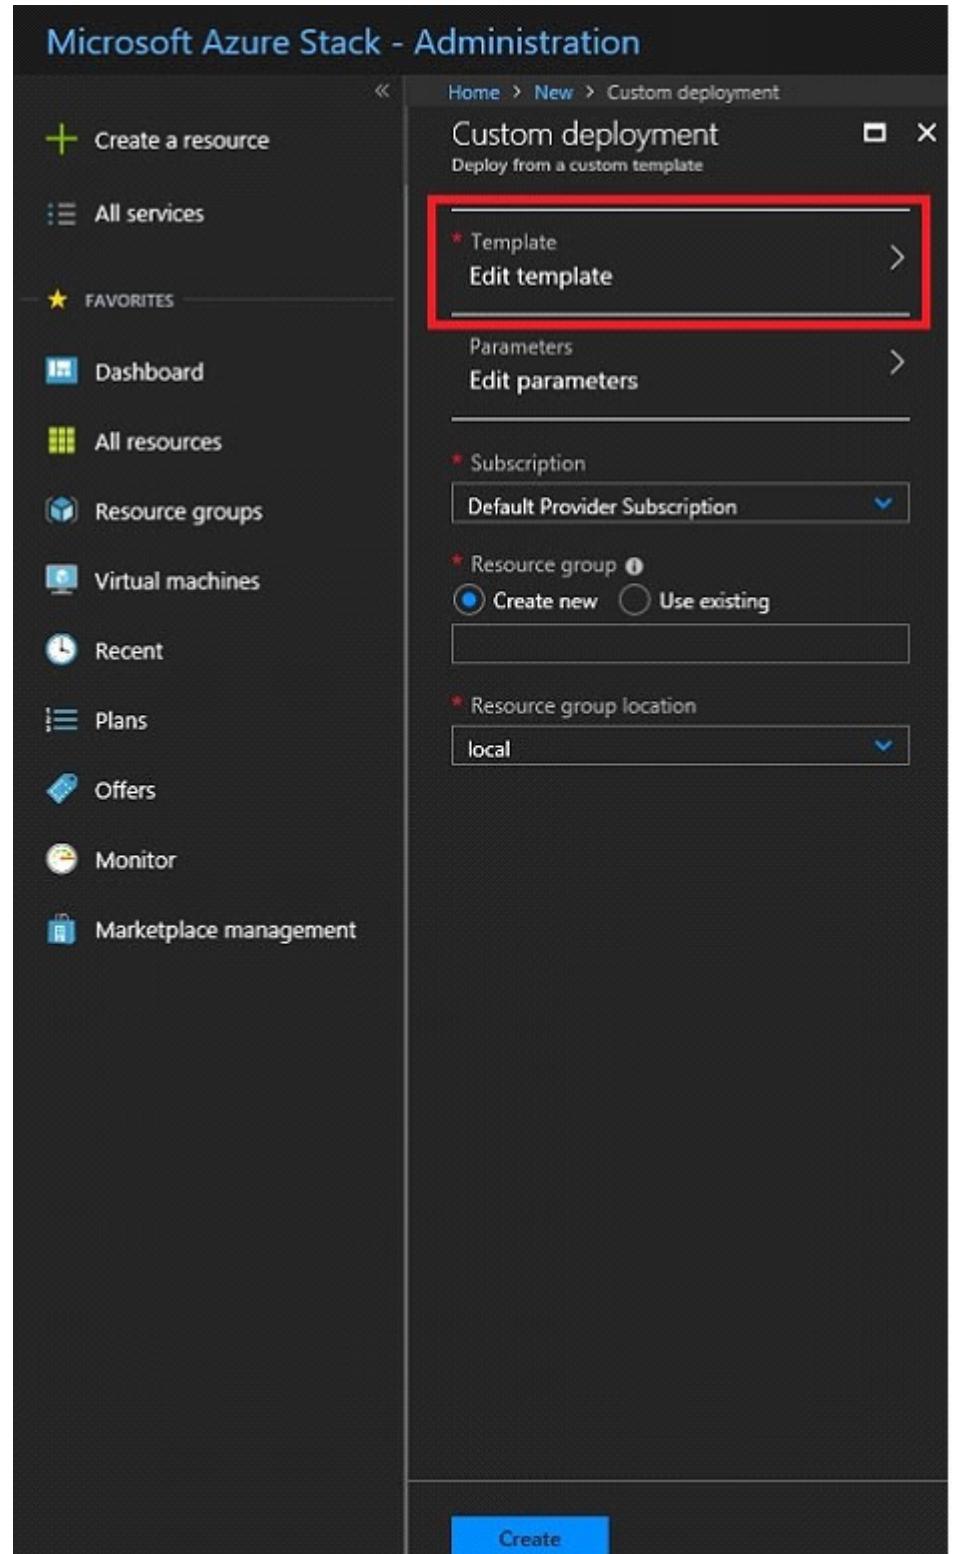
4. Change the ASddressPrefixes to 10.0.0.0/24 in order to support only 256 total IP addresses.

addressSpace": {"addressPrefixes": ["10.0.0.0/24"]},

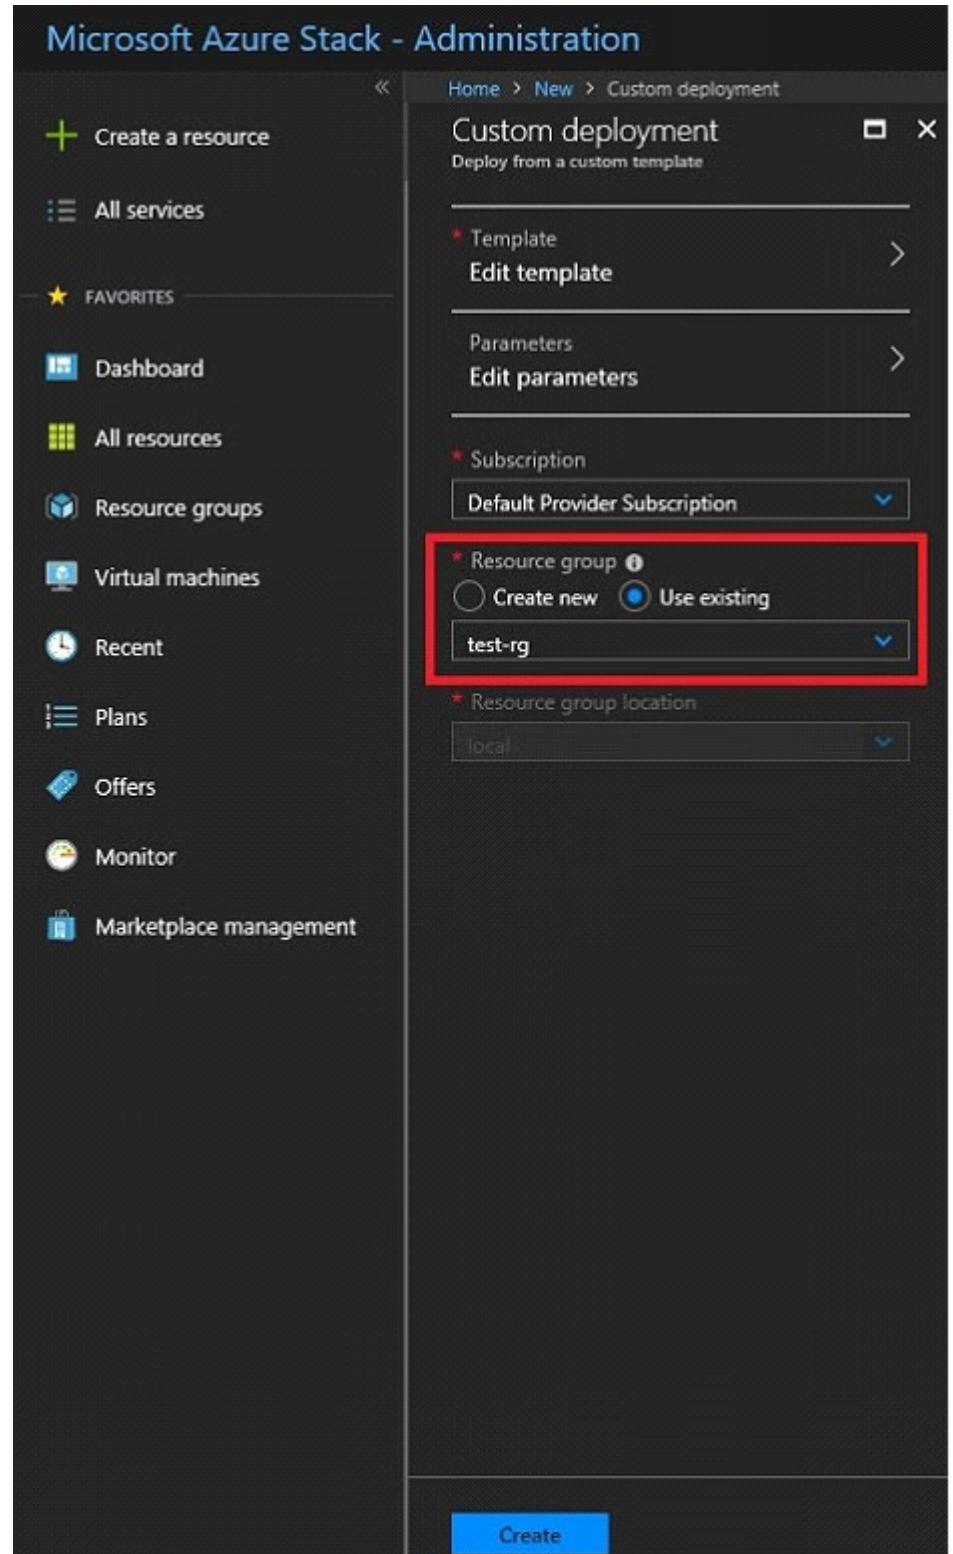
5. Change the firstSubnet addressprefix to 10.0.0.0/26 to support only 64 total IP addresses.

```
"subnets": [  
{  
  "name": "firstSubnet",  
  "properties": {  
    "addressPrefix": "10.0.0.0/24"  
  }  
}]
```

6. Select Save.



7. Select Edit parameters, provide values for the parameters that are shown, and then select OK.
8. Select Subscription. Choose the subscription you want to use, and then select OK.
9. Select Resource group. Choose an existing resource group or create a new one, and then select OK.



10. Select Create. A new tile on the dashboard tracks the progress of your template deployment.

References:

<https://docs.microsoft.com/en-us/azure-stack/user/azure-stack-deploy-template-portal?view=azs-1908>

<https://docs.microsoft.com/en-us/azure/architecture/building-blocks/extending-templates/update-resource>

QUESTION 167**SIMULATION**

You need to configure an Azure web app named az400-9940427-main to contain an environmental variable named "MAX_ITEMS". The environmental variable must have a value of 50. To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.
Explanation

Explanation/Reference:
Explanation:

1. In the Azure portal, navigate to the az400-9940427-main app's management page. In the app's left menu, click Configuration > Application settings.

The screenshot shows the Azure portal configuration blade for an App Service named 'my-core-app'. The left sidebar lists various configuration sections: Security, Deployment, Quickstart, Deployment slots, Deployment Center, Settings, and Configuration. The 'Configuration' section is selected and highlighted with a red box. The main content area is titled 'Application settings' and contains a note about encrypted data and environment variables. It includes buttons for 'New application setting', 'Show values', 'Advanced edit', and 'Filter'. A table below shows no application settings. The 'Connection strings' section also shows no entries.

2. Click New Application settings

3. Enter the following:
Name: MAX_ITEMS

Value: 50

References:
<https://docs.microsoft.com/en-us/azure/app-service/configure-common>

Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation

1. In the Azure portal, navigate to the az400-9940427-main app's management page. In the app's left menu, click Configuration > Application settings.

The screenshot shows the Azure portal configuration interface for an App Service named "my-core-app". The left sidebar lists various configuration sections like Security, Deployment, Quickstart, Deployment slots, Deployment Center, Settings, and Configuration (which is highlighted with a red box). The main content area has tabs for Application settings, General settings, Default documents, and Path mappings. The Application settings tab is active, displaying a message about encrypted transmission and environment variables. It includes buttons for New application setting, Show values, Advanced edit, and Filter. A table shows no application settings. The Connection strings tab also shows no connection strings.

2. Click New Application settings

3. Enter the following:

Name: MAX_ITEMS

Value: 50

References:
<https://docs.microsoft.com/en-us/azure/app-service/configure-common>

QUESTION 168

DRAG DROP

You have an Azure Kubernetes Service (AKS) implementation that is RBAC-enabled.

You plan to use Azure Container Instances as a hosted development environment to run containers in the AKS implementation.

You need to configure Azure Container Instances as a hosted environment for running the containers in AKS.

Which three actions should you perform in sequence? To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.

Select and Place:

Actions	Answer Area
Run helm init.	0
Run az aks install-connector.	0
Create a YAML file.	0
Run az role assignment create	
Run kubectl apply.	

A. Answer:

Actions	Answer Area
Run az aks install-connector.	Create a YAML file. Run kubectl apply. Run helm init.
	Run az role assignment create

Explanation

Explanation/Reference:

Explanation:

Step 1: Create a YAML file.

If your AKS cluster is RBAC-enabled, you must create a service account and role binding for use with Tiller. To create a service account and role binding, create a file named rbac-virtual-kubelet.yaml

Step 2: Run kubectl apply.

Apply the service account and binding with kubectl apply and specify your rbac-virtual-kubelet.yaml file.

Step 3: Run helm init.

Configure Helm to use the tiller service account:

```
helm init --service-account tiller
```

You can now continue to installing the Virtual Kubelet into your AKS cluster.

References: <https://docs.microsoft.com/en-us/azure/aks/virtual-kubelet>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Step 1: Create a YAML file.

If your AKS cluster is RBAC-enabled, you must create a service account and role binding for use with Tiller. To create a service account and role binding, create a file named rbac-virtual-kubelet.yaml

Step 2: Run kubectl apply.

Apply the service account and binding with kubectl apply and specify your rbac-virtual-kubelet.yaml file.

Step 3: Run helm init.

Configure Helm to use the tiller service account:

```
helm init --service-account tiller
```

You can now continue to installing the Virtual Kubelet into your AKS cluster.

References: <https://docs.microsoft.com/en-us/azure/aks/virtual-kubelet>

QUESTION 169

You have an application that consists of several Azure App Service web apps and Azure functions.

You need to assess the security of the web apps and the functions.

Which Azure feature can you use to provide a recommendation for the security of the application?

- A. Security & Compliance in Azure Log Analytics
- B. Resource health in Azure Service Health
- C. Smart Detection in Azure Application Insights
- D. Compute & apps in Azure Security Center

Answer: D

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Monitor compute and app services: Compute & apps include the App Services tab, which App services: list of your App service environments and current security state of each.

Recommendations

This section has a set of recommendations for each VM and computer, web and worker roles, Azure App Service Web Apps, and Azure App Service Environment that Security Center monitors. The first column lists the recommendation. The second column shows the total number of resources that are affected by that recommendation. The third column shows the severity of the issue.

Incorrect Answers:

C: Smart Detection automatically warns you of potential performance problems, not security problems in your web application.

References:

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/proactive-diagnostics>

QUESTION 170

You have a private distribution group that contains provisioned and unprovisioned devices.

You need to distribute a new iOS application to the distribution group by using Microsoft Visual Studio App Center.

What should you do?

- A. Request the Apple ID associated with the user of each device.
- B. Register the devices on the Apple Developer portal.
- C. Create an active subscription in App Center Test.
- D. Add the device owner to the organization in App Center.

Answer: B

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:
When releasing an iOS app signed with an ad-hoc or development provisioning profile, you must obtain tester's device IDs (UDIDs), and add them to the provisioning profile before compiling a release. When you enable the distribution group's Automatically manage devices setting, App Center automates the before mentioned operations and removes the constraint for you to perform any manual tasks. As part of automating the workflow, you must provide the user name and password for your Apple ID and your production certificate in a .p12 format.

App Center starts the automated tasks when you distribute a new release or one of your testers registers a new device. First, all devices from the target distribution group will be registered, using your Apple ID, in your developer portal and all provisioning profiles used in the app will be generated with both new and existing device ID. Afterward, the newly generated provisioning profiles are downloaded to App Center servers.

References:
<https://docs.microsoft.com/en-us/appcenter/distribution/groups>

QUESTION 171

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

After you answer a question in this section, you will NOT be able to return to it. As a result, these questions will not appear in the review screen.

You manage a project in Azure DevOps.

You need to prevent the configuration of the project from changing over time.

Solution: Perform a Subscription Health scan when packages are created.

Does this meet the goal?

- A. Yes
- B. No

Answer: B

Explanation/Reference:
Explanation

Explanation/Reference:

Explanation:

Instead implement Continuous Assurance for the project.

Note: The Subscription Security health check features in AzSK contains a set of scripts that examines a subscription and flags off security issues, misconfigurations or obsolete artifacts/settings which can put your subscription at higher risk.

Reference:
<https://azsk.azurewebsites.net/04-Continous-Assurance/Readme.html>

QUESTION 172

You are developing an iOS application by using Azure DevOps.

You need to test the application manually on 10 devices without releasing the application to the public.

Which two actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Create a Microsoft Intune device compliance policy.
- B. Deploy a certificate from an internal certification authority (CA) to each device.
- C. Register the application in the iTunes store.
- D. Onboard the devices into Microsoft Intune.
- E. Distribute a new release of the application.
- F. Register the IDs of the devices in the Apple Developer portal.

Answer: BF

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

B: Follow these steps to register the devices:

Select the Register devices button.

A dialog prompts for your username and password used in the Apple Developer portal.

Once you sign in with your Apple username and password, App Center adds the unprovisioned devices to both your Apple developer account and the releases provisioning profile.

Optionally you can upload a .p12 file to re-sign the app and distribute it to the newly added devices. Read more on how to generate a .p12 file.

F: Registering a device means making it part of the list of devices on the Apple Developer portal that can then be included in a provisioning profile.

Incorrect Answers:

C: Only register the application in the iTunes store when it is ready for public release.

References:

<https://docs.microsoft.com/en-us/appcenter/distribution/auto-provisioning>

QUESTION 173

You have a private distribution group that contains provisioned and unprovisioned devices.

You need to distribute a new iOS application to the distribution group by using Microsoft Visual Studio App Center.

What should you do?

- A. Select Register devices and sign my app.
- B. Generate a new .p12 file for each device.
- C. Create an active subscription in App Center Test.
- D. Add the device owner to the collaborators group.

Answer: A

Explanation/Reference:

Explanation

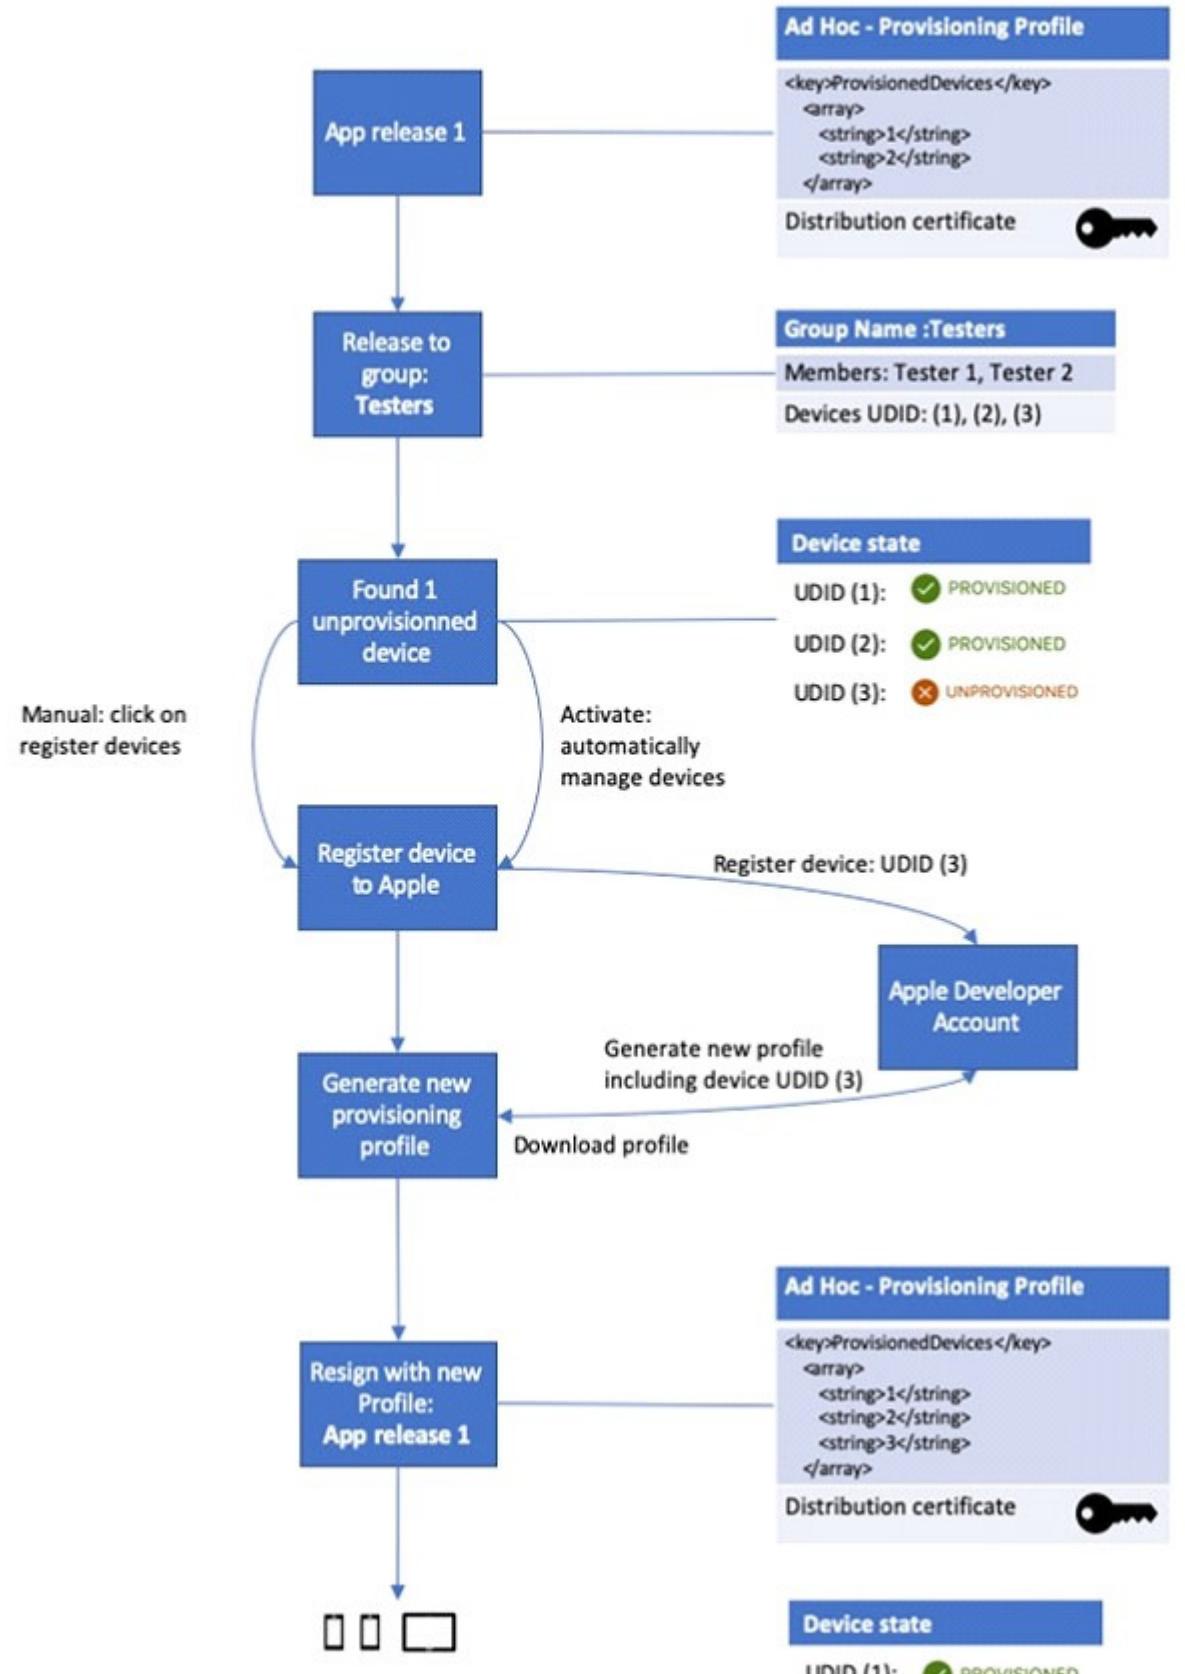
Explanation/Reference:

Explanation:

The following diagram displays the entire app re-signing flow in App Center.



App Center
install.appcenter.ms/apps



Incorrect Answers:

B: Only one .p12 file for the app, not one for each device.

Reference:

<https://docs.microsoft.com/hu-hu/appcenter/distribution/auto-provisioning>

QUESTION 174

SIMULATION

You plan to deploy a website that will be hosted in two Azure regions.

You need to create an Azure Traffic Manager profile named az40011566895n1-tm in a resource group named RG1lod11566895. The solution must ensure that users will always connect to a copy of the website that is in the same country.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

1. Go to the Azure portal, navigate to Traffic Manager profiles and click on the Add button to create a routing profile.

The screenshot shows the 'Traffic Manager profiles' page in the Azure portal. At the top, there's a header with a menu icon, the title 'Traffic Manager profiles', and the Microsoft logo. Below the header is a toolbar with a '+' icon for 'Add', a 'Columns' icon, and a 'Refresh' icon. A message 'Subscriptions: All 4 selected' is displayed above a search bar containing 'Filter by name...'. To the right of the search bar is a button labeled 'All subscriptions'. At the bottom left, it says '22 items'.

2. In the Create Traffic Manager profile, enter, or select these settings:

Name: az40011566895n1-tm

Routing method: Geographic

Resource group: RG1lod11566895

The screenshot shows the 'Create Traffic Manager profile' dialog. It has several input fields:

- * Name: samplegeoprofile.trafficmanager.net
- Routing method: Geographic
- * Subscription: A dropdown menu showing a single option.
- * Resource group: Create new (radio button selected). Option: geoprofilerg
- * Resource group location: West US

Note: Traffic Manager profiles can be configured to use the Geographic routing method so that users are directed to specific endpoints (Azure, External or Nested) based on which geographic location their DNS query originates from. This empowers Traffic

Answer:

Explanation/Reference:

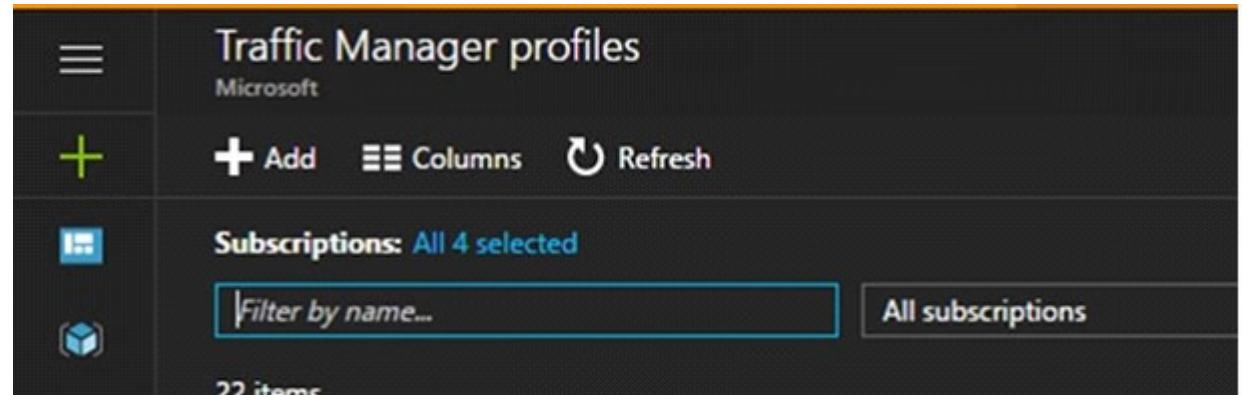
Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

1. Go to the Azure portal, navigate to Traffic Manager profiles and click on the Add button to create a routing profile.

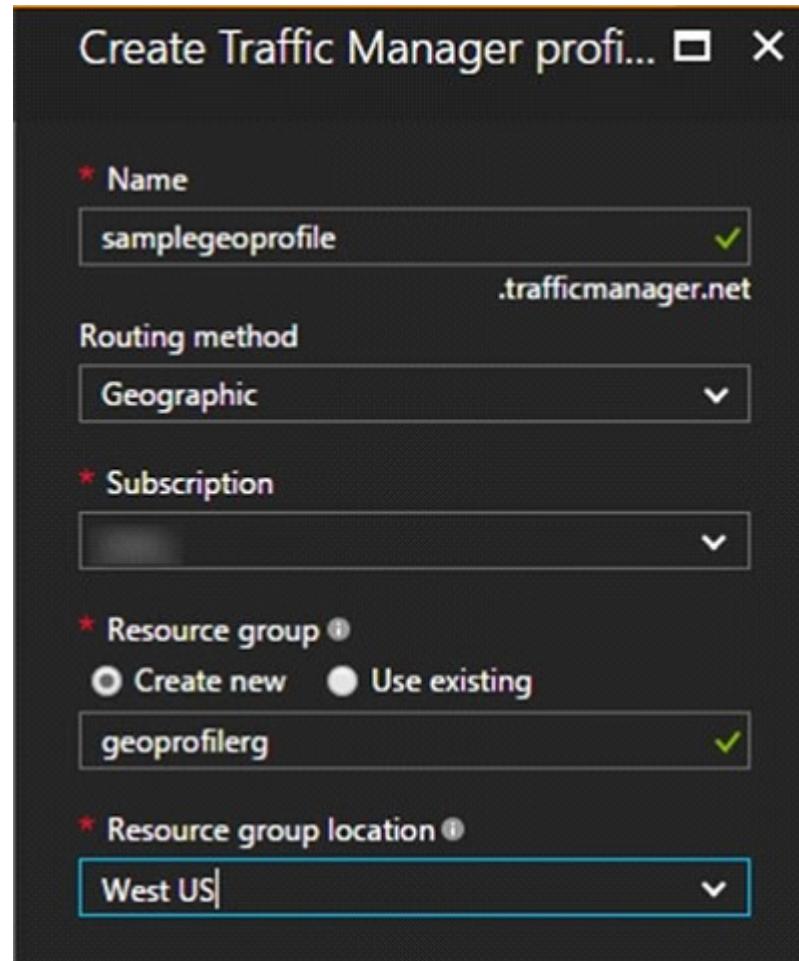


2. In the Create Traffic Manager profile, enter, or select these settings:

Name: az40011566895n1-tm

Routing method: Geographic

Resource group: RG1lod11566895



Note: Traffic Manager profiles can be configured to use the Geographic routing method so that users are directed to specific endpoints (Azure, External or Nested) based on which geographic location their DNS query originates from. This empowers Traffic Manager customers to enable scenarios where knowing a user's geographic region and routing them based on that is important.

Reference:

<https://azure.microsoft.com/en-us/blog/announcing-the-general-availability-of-geographic-routing-capability-in-azure-traffic-manager/>
Question Set 1

QUESTION 175

You have a multi-tier application. The front end of the application is hosted in Azure App Service.

You need to identify the average load times of the application pages.

What should you use?

- A. Azure Application Insights
- B. the activity log of the App Service
- C. the diagnostics logs of the App Service
- D. Azure Advisor

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Application Insights will tell you about any performance issues and exceptions, and help you find and diagnose the root causes.

Application Insights can monitor both Java and ASP.NET web applications and services, WCF services. They can be hosted on-premises, on virtual machines, or as Microsoft Azure websites.

On the client side, Application Insights can take telemetry from web pages and a wide variety of devices including iOS, Android, and Windows Store apps.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/web-monitor-performance>

QUESTION 176

SIMULATION

You need to create an instance of Azure Application Insights named az400-9940427-main and configure the instance to receive telemetry data from an Azure web app named az400-9940427-main.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

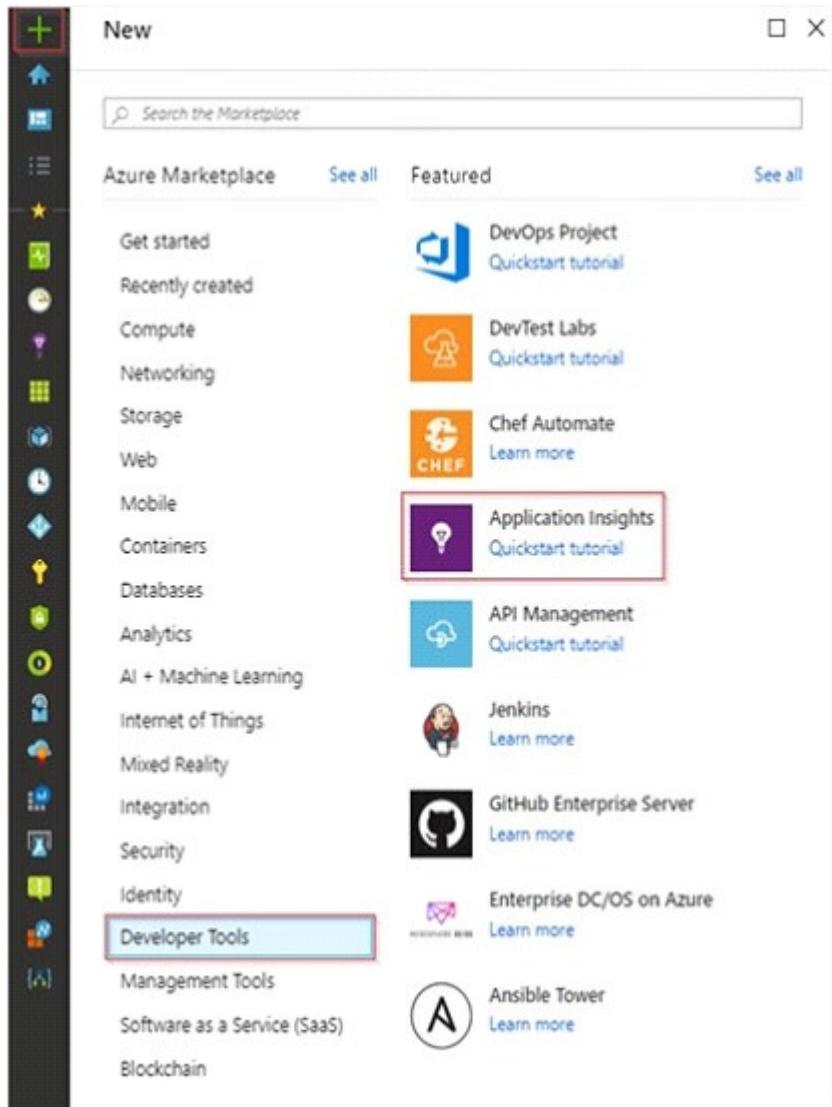
Explanation/Reference:

Explanation:

Step 1: Create an instance of Azure Application Insights

1. Open Microsoft Azure Portal

2. Log into your Azure account, Select Create a resource > Developer tools > Application Insights.

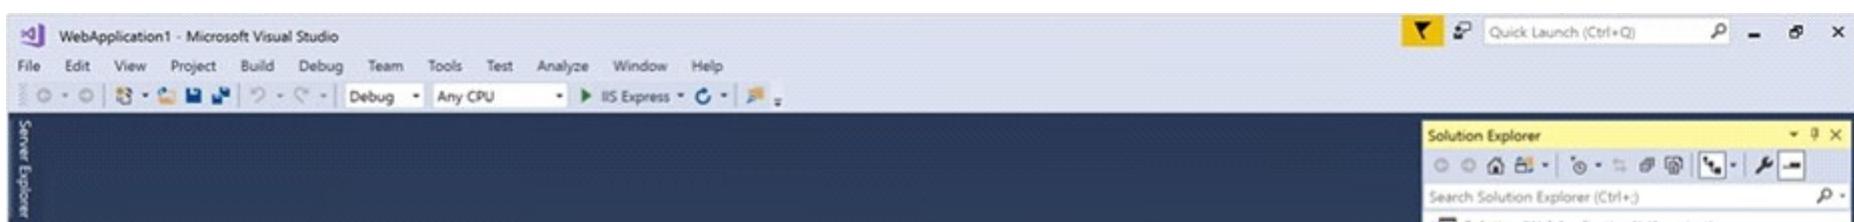


3. Enter the following settings, and then select Review + create.

Name: az400-9940427-main

Step 2: Configure App Insights SDK

4. Open your ASP.NET Core Web App project in Visual Studio > Right-click on the AppName in the Solution Explorer > Select Add > Application Insights Telemetry.



Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

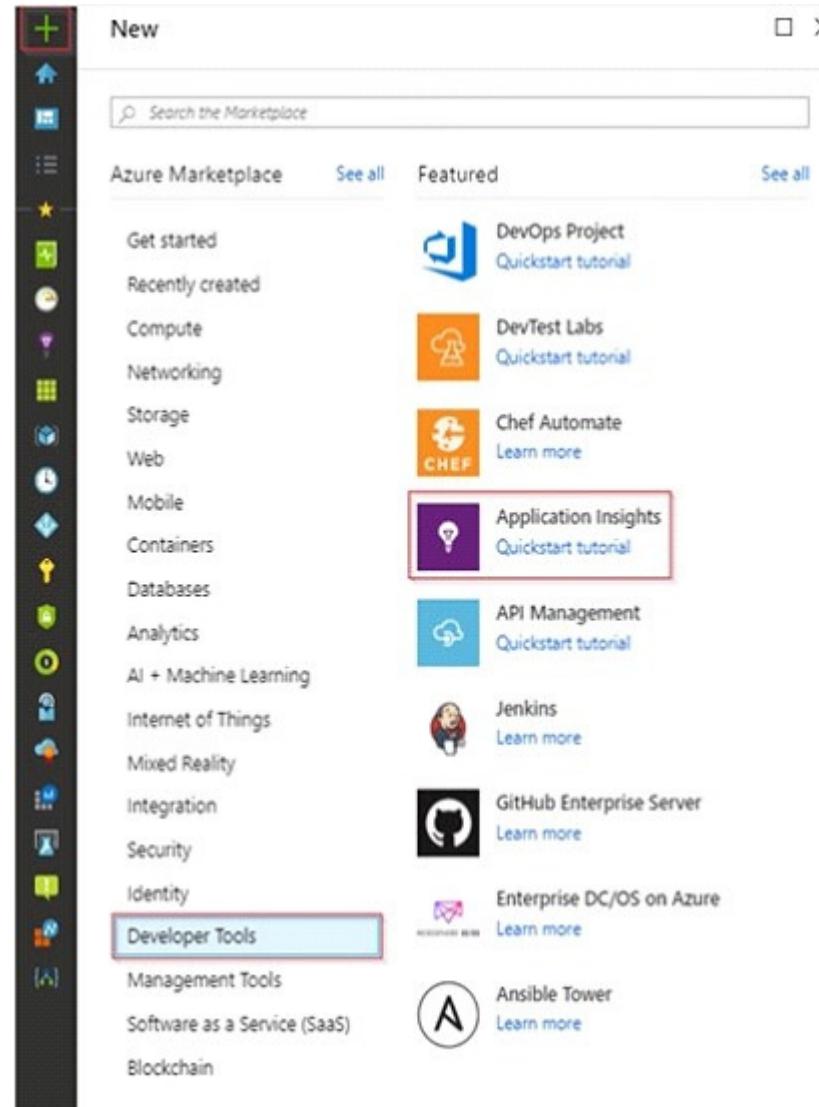
Explanation/Reference:

Explanation:

Step 1: Create an instance of Azure Application Insights

1. Open Microsoft Azure Portal

2. Log into your Azure account, Select Create a resource > Developer tools > Application Insights.

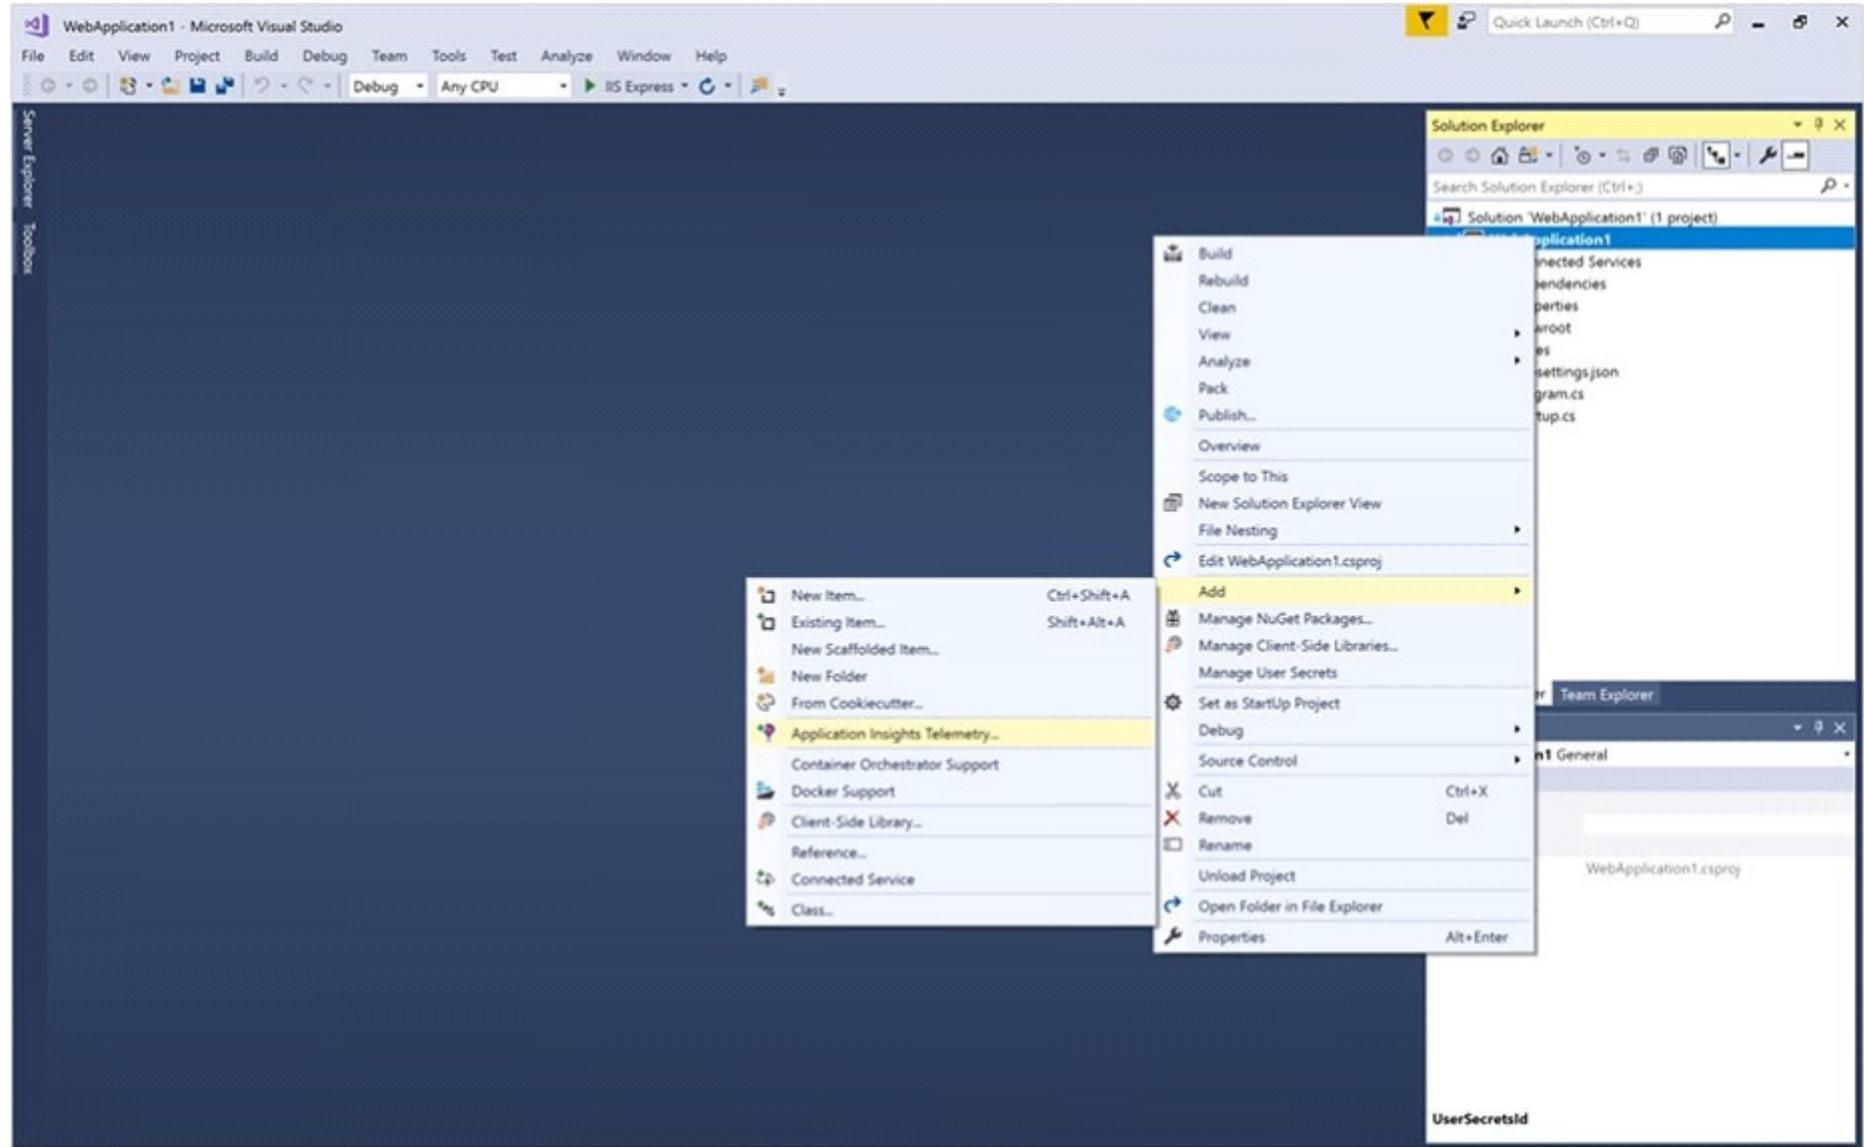


3. Enter the following settings, and then select Review + create.

Name: az400-9940427-main

Step 2: Configure App Insights SDK

4. Open your ASP.NET Core Web App project in Visual Studio > Right-click on the AppName in the Solution Explorer > Select Add > Application Insights Telemetry.



5. Click the Get Started button

6. Select your account and subscription > Select the Existing resource you created in the Azure portal > Click Register.

References:

<https://docs.microsoft.com/bs-latn-ba/azure/azure-monitor/learn/dotnetcore-quick-start?view=vs-2017>

QUESTION 177

Your company uses ServiceNow for incident management.

You develop an application that runs on Azure.

The company needs to generate a ticket in ServiceNow when the application fails to authenticate.

Which Azure Log Analytics solution should you use?

- A. Application Insights Connector

- B. Automation & Control
- C. IT Service Management Connector (ITSM)
- D. Insight & Analytics

Answer: C

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

The IT Service Management Connector (ITSMC) allows you to connect Azure and a supported IT Service Management (ITSM) product/service.

ITSMC supports connections with the following ITSM tools:

ServiceNow

System Center Service Manager

Provance

Cherwell

With ITSMC, you can

Create work items in ITSM tool, based on your Azure alerts (metric alerts, Activity Log alerts and Log Analytics alerts).

Optionally, you can sync your incident and change request data from your ITSM tool to an Azure Log Analytics workspace.

References: <https://docs.microsoft.com/en-us/azure/azure-monitor/platform/itsmc-overview>

QUESTION 178

HOTSPOT

Your company is building a new web application.

You plan to collect feedback from pilot users on the features being delivered.

All the pilot users have a corporate computer that has Google Chrome and the Microsoft Test & Feedback extension installed. The pilot users will test the application by using Chrome.

You need to identify which access levels are required to ensure that developers can request and gather feedback from the pilot users. The solution must use the principle of least privilege.

Which access levels in Azure DevOps should you identify? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

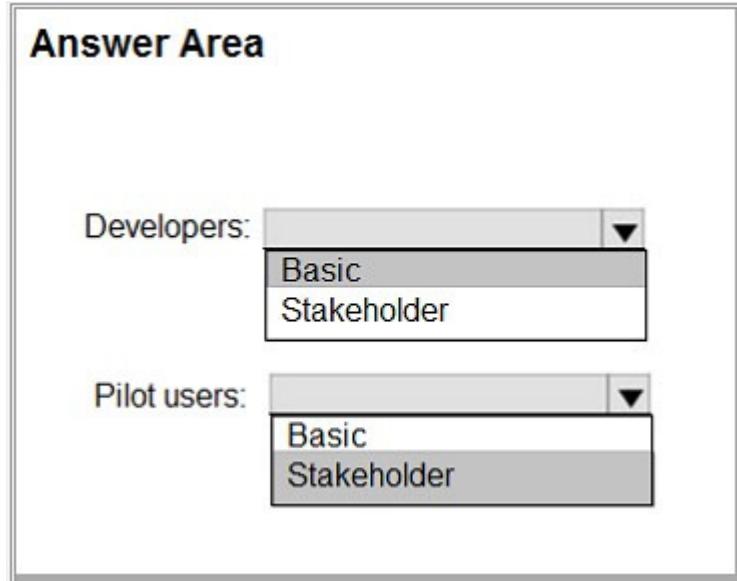
Developers:

- Basic
- Stakeholder

Pilot users:

- Basic
- Stakeholder

A. Answer:



Explanation

Explanation/Reference:
Explanation:

Box 1: Basic

Assign Basic to users with a TFS CAL, with a Visual Studio Professional subscription, and to users for whom you are paying for Azure Boards & Repos in an organization.

Box 2: Stakeholder

Assign Stakeholders to users with no license or subscriptions who need access to a limited set of features.

Note:

You assign users or groups of users to one of the following access levels:

Basic: provides access to most features

VS Enterprise: provides access to premium features

Stakeholders: provides partial access, can be assigned to unlimited users for free

References: <https://docs.microsoft.com/en-us/azure/devops/organizations/security/access-levels?view=vsts>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box 1: Basic

Assign Basic to users with a TFS CAL, with a Visual Studio Professional subscription, and to users for whom you are paying for Azure Boards & Repos in an organization.

Box 2: Stakeholder

Assign Stakeholders to users with no license or subscriptions who need access to a limited set of features.

Note:

You assign users or groups of users to one of the following access levels:

Basic: provides access to most features

VS Enterprise: provides access to premium features

Stakeholders: provides partial access, can be assigned to unlimited users for free

References: <https://docs.microsoft.com/en-us/azure/devops/organizations/security/access-levels?view=vsts>

QUESTION 179

You use Azure SQL Database Intelligent Insights and Azure Application Insights for monitoring.

You need to write ad-hoc queries against the monitoring data.

Which query language should you use?

- A. Azure Log Analytics
- B. PL/pgSQL
- C. PL/SQL
- D. Transact-SQL

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Data analysis in Azure SQL Analytics is based on Log Analytics language for your custom querying and reporting.

References: <https://docs.microsoft.com/en-us/azure/azure-monitor/insights/azure-sql>

QUESTION 180

You have a multi-tier application that has an Azure Web Apps front end and an Azure SQL Database back end.

You need to recommend a solution to capture and store telemetry data. The solution must meet the following requirements:

Support using ad-hoc queries to identify baselines.

Trigger alerts when metrics in the baseline are exceeded.

Store application and database metrics in a central location.

What should you include in the recommendation?

- A. Azure Event Hubs
- B. Azure SQL Database Intelligent Insights
- C. Azure Application Insights
- D. Azure Log Analytics

Answer: D

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Azure Platform as a Service (PaaS) resources, like Azure SQL and Web Sites (Web Apps), can emit performance metrics data natively to Log Analytics.

The Premium plan will retain up to 12 months of data, giving you an excellent baseline ability.

There are two options available in the Azure portal for analyzing data stored in Log analytics and for creating queries for ad hoc analysis.

Incorrect Answers:

B: Intelligent Insights analyzes database performance by comparing the database workload from the last hour with the past seven-day baseline workload. However, we need handle application metrics as well.

References: <https://docs.microsoft.com/en-us/azure/azure-monitor/platform/collect-azurepass-posh>

QUESTION 181

Your company creates a web application.

You need to recommend a solution that automatically sends to Microsoft Teams a daily summary of the exceptions that occur in the application.

Which two Azure services should you recommend? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Azure Logic Apps
- B. Azure Pipelines
- C. Microsoft Visual Studio App Center
- D. Azure DevOps Project
- E. Azure Application Insights

Answer: AE

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

E: Exceptions in your live web app are reported by Application Insights.

Note: Periodical reports help keep a team informed on how their business critical services are doing. Developers, DevOps/SRE teams, and their managers can be productive with automated reports reliably delivering insights without requiring everyone to sign in
the portal. Such reports can also help identify gradual increases in latencies, load or failure rates that may not trigger any alert rules.

A: You can programmatically query Application Insights data to generate custom reports on a schedule. The following options can help you get started quickly:

Automate reports with Microsoft Flow

Automate reports with Logic Apps

References:

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/asp-net-exceptions>

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/automate-custom-reports>

QUESTION 182**DRAG DROP**

Your company wants to use Azure Application Insights to understand how user behaviors affect an application.

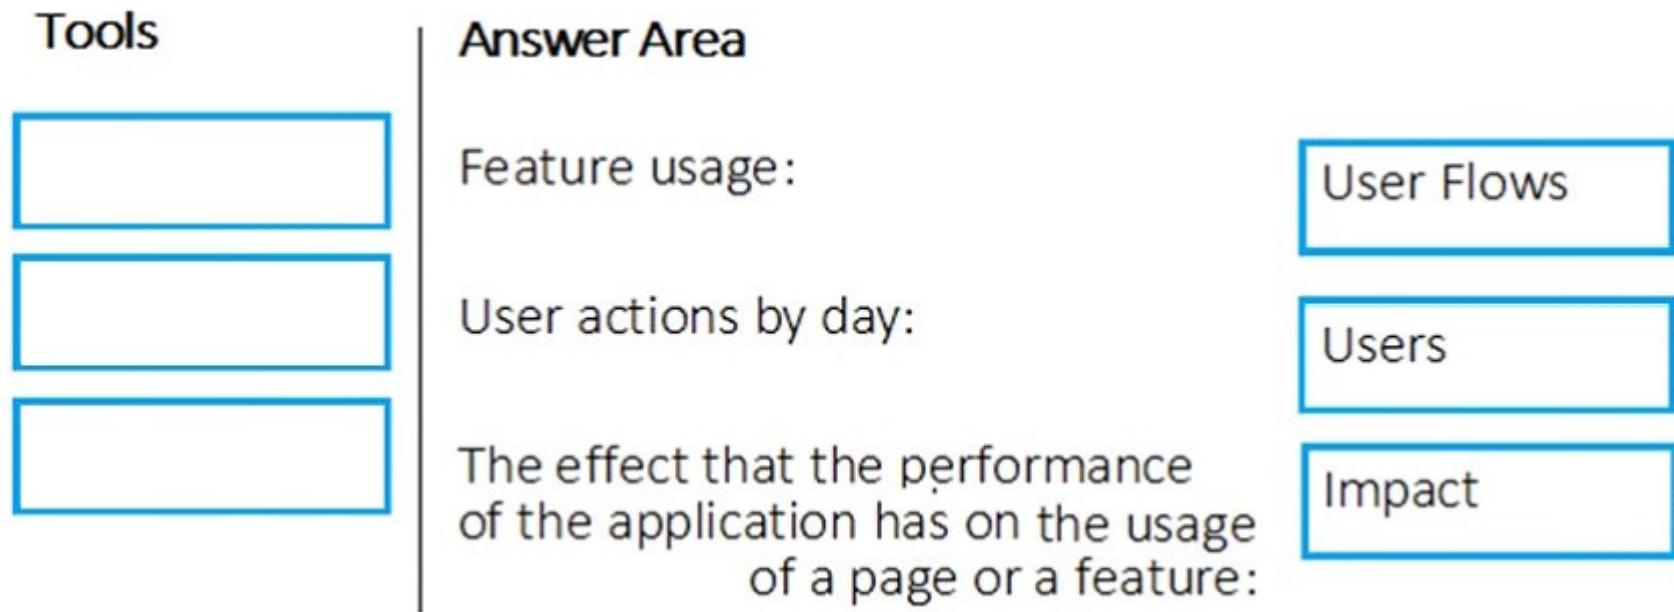
Which Application Insights tool should you use to analyze each behavior? To answer, drag the appropriate tools to the correct behaviors. Each tool may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Tools	Answer Area
Impact	Feature usage: <input type="text"/>
User Flows	User actions by day: <input type="text"/>
Users	The effect that the performance of the application has on the usage of a page or a feature: <input type="text"/>

A. Answer:



Explanation

Explanation/Reference:
Explanation:

Box 1: User Flows

The User Flows tool visualizes how users navigate between the pages and features of your site. It's great for answering questions like:

How do users navigate away from a page on your site?

What do users click on a page on your site?

Where are the places that users churn most from your site?

Are there places where users repeat the same action over and over?

Box 2: Users

Box 3: Impact

Reference:

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/usage-flows>

Answer:

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

Box 1: User Flows

The User Flows tool visualizes how users navigate between the pages and features of your site. It's great for answering questions like:

How do users navigate away from a page on your site?

What do users click on a page on your site?

Where are the places that users churn most from your site?

Are there places where users repeat the same action over and over?

Box 2: Users
Box 3: Impact

Reference:
<https://docs.microsoft.com/en-us/azure/azure-monitor/app/usage-flows>

QUESTION 183

Your company is building a mobile app that targets Android and iOS devices.

Your team uses Azure DevOps to manage all work items and release cycles.

You need to recommend a solution to perform the following tasks:

Collect crash reports for issue analysis.

Distribute beta releases to your testers.

Get user feedback on the functionality of new apps.

What should you include in the recommendation?

- A. the Microsoft Test & Feedback extension
- B. Microsoft Visual Studio App Center integration
- C. Azure Application Insights widgets
- D. Jenkins integration

Answer: A

Explanation/Reference:

Explanation

Explanation/Reference:

Explanation:

The "Exploratory Testing" extension is now "Test & Feedback" and is now Generally Available.

Anyone can now test web apps and give feedback, all directly from the browser on any platform: Windows, Mac, or Linux. Available for Google Chrome and Mozilla Firefox (required version 50.0 or above) currently. Support for Microsoft Edge is in the pipeline and will be enabled once Edge moves to a Chromium-compatible web platform.

References:

<https://marketplace.visualstudio.com/items?itemName=ms.vss-exploratorytesting-web>

QUESTION 184

You have an Azure DevOps project named Project1 and an Azure subscription named Sub1. Sub1 contains an Azure virtual machine scale set named VMSS1. VMSS1 hosts a web application named WebApp1. WebApp1 uses stateful sessions.

The WebApp1 installation is managed by using the Custom Script extension. The script resides in an Azure Storage account named sa1.

You plan to make a minor change to a UI element of WebApp1 and to gather user feedback about the change.

You need to implement limited user testing for the new version of WebApp1 on VMSS1.

Which three actions should you perform? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- A. Modify the load balancer settings of VMSS1.
- B. Redeploy VMSS1.
- C. Upload a custom script file to sa1.
- D. Modify the Custom Script extension settings of VMSS1.
- E. Update the configuration of a virtual machine in VMSS1.

Answer: BCD

Explanation/Reference:

Explanation

Explanation/Reference:

QUESTION 185

SIMULATION

You need to create a notification if the peak average response time of an Azure web app named az400-9940427-main is more than five seconds when evaluated during a five-minute period. The notification must trigger the "<https://contoso.com/notify>" webhook.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

1. Open Microsoft Azure Portal

2. Log into your Azure account and go to App Service and look under Monitoring then you will see Alert.

3. Select Add an alert rule

4. Configure the alert rule as per below and click Ok.

Source: Alert on Metrics

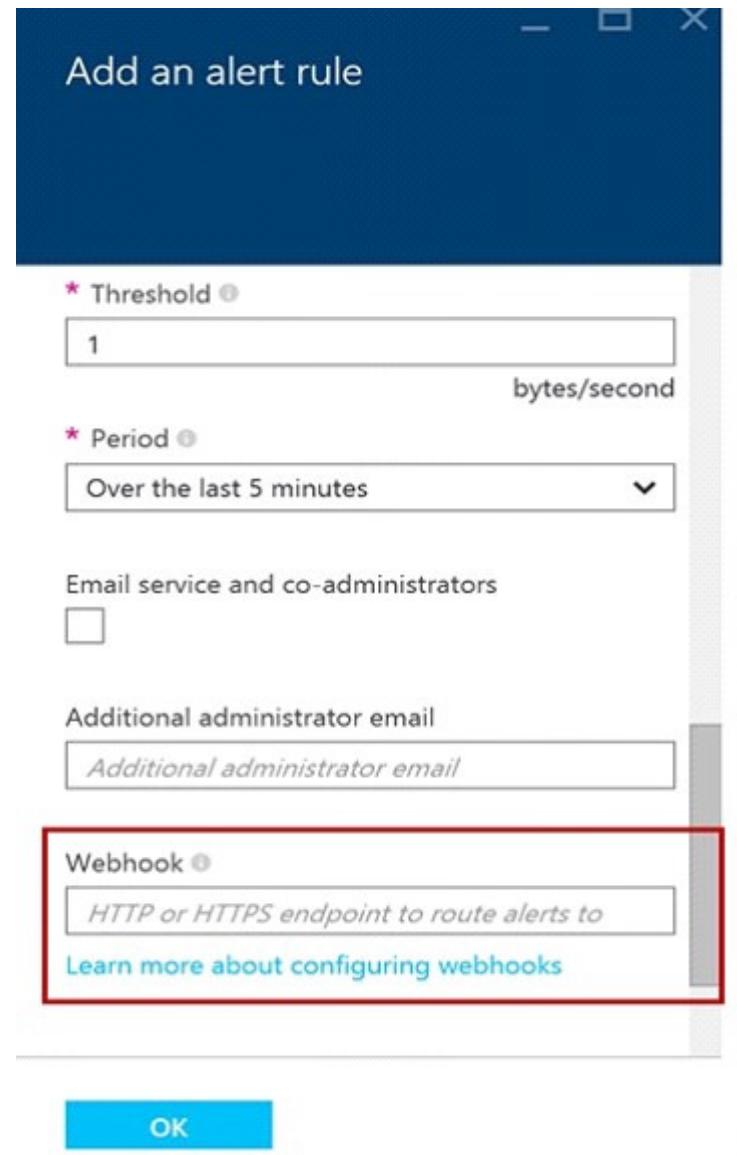
Resource Group: az400-9940427-main

Resource: az400-9940427-main

Threshold: 5

Period: Over the last 5 minutes

Webhook: <https://contoso.com/notify>



Reference:

<https://azure.microsoft.com/es-es/blog/webhooks-for-azure-alerts/>

Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

1. Open Microsoft Azure Portal

2. Log into your Azure account and go to App Service and look under Monitoring then you will see Alert.

3. Select Add an alert rule

4. Configure the alert rule as per below and click Ok.

Source: Alert on Metrics

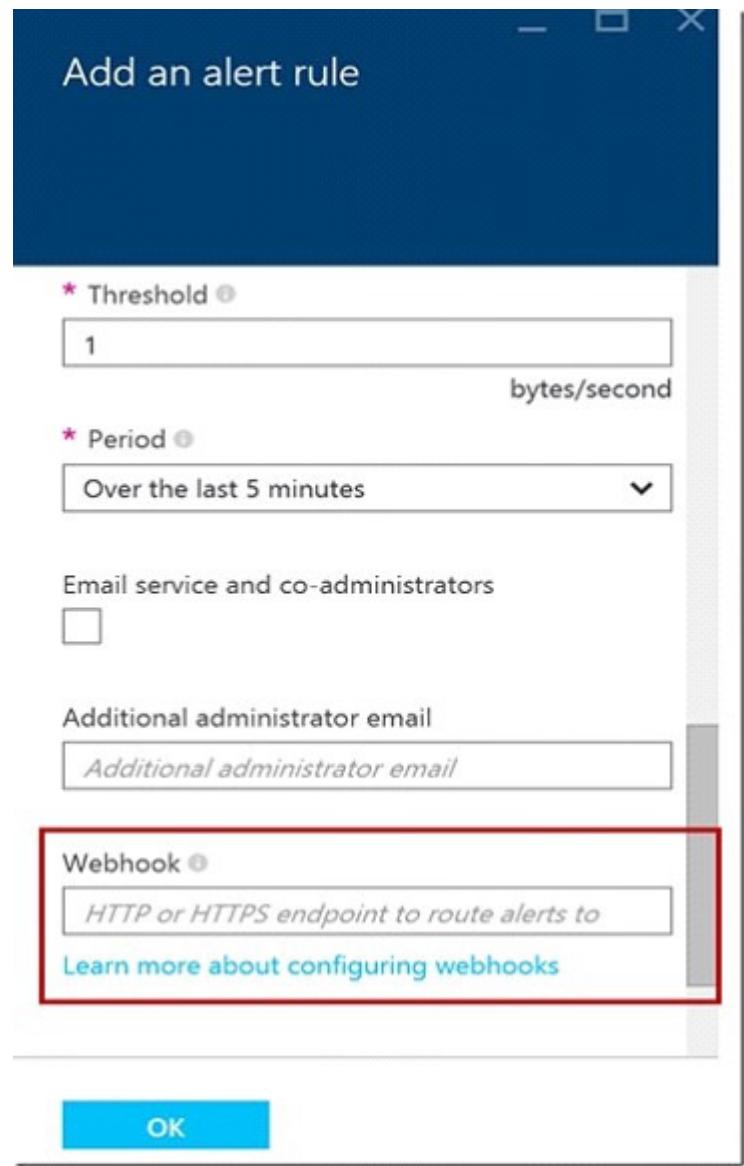
Resource Group: az400-9940427-main

Resource: az400-9940427-main

Threshold: 5

Period: Over the last 5 minutes

Webhook: <https://contoso.com/notify>



Reference:
<https://azure.microsoft.com/es-es/blog/webhooks-for-azure-alerts/>

QUESTION 186

SIMULATION

You need to create and configure an Azure Storage account named az400lod11566895stor in a resource group named RG1lod11566895 to store the boot diagnostics for a virtual machine named VM1.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

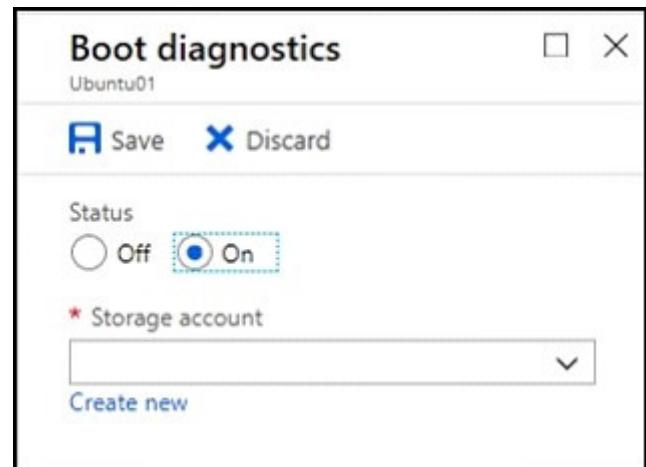
Step 1: To create a general-purpose v2 storage account in the Azure portal, follow these steps:

1. On the Azure portal menu, select All services. In the list of resources, type Storage Accounts. As you begin typing, the list filters based on your input. Select Storage Accounts.
2. On the Storage Accounts window that appears, choose Add.
3. Select the subscription in which to create the storage account.
4. Under the Resource group field, select RG1lod11566895
5. Next, enter a name for your storage account named: az400lod11566895stor
6. Select Create.

Step 2: Enable boot diagnostics on existing virtual machine

To enable Boot diagnostics on an existing virtual machine, follow these steps:

1. Sign in to the Azure portal, and then select the virtual machine VM1.
2. In the Support + troubleshooting section, select Boot diagnostics, then select the Settings tab.
3. In Boot diagnostics settings, change the status to On, and from the Storage account drop-down list, select the storage account az400lod11566895stor.
4. Save the change.



You must restart the virtual machine for the change to take effect.

Reference:

<https://docs.microsoft.com/en-us/azure/storage/common/storage-account-create>

<https://docs.microsoft.com/en-us/azure/virtual-machines/troubleshooting/boot-diagnostics>

Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

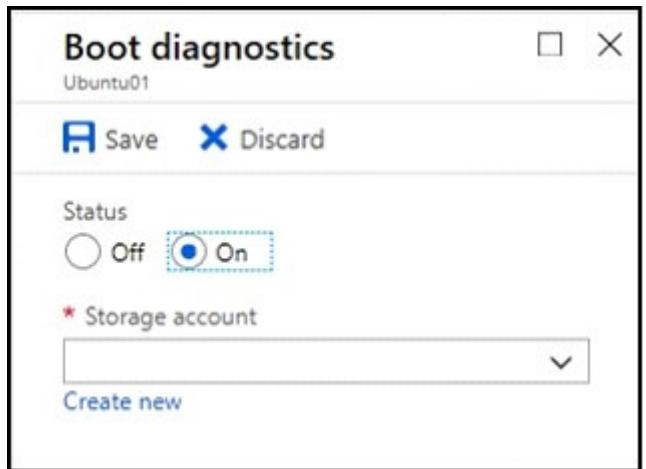
Step 1: To create a general-purpose v2 storage account in the Azure portal, follow these steps:

1. On the Azure portal menu, select All services. In the list of resources, type Storage Accounts. As you begin typing, the list filters based on your input. Select Storage Accounts.
2. On the Storage Accounts window that appears, choose Add.
3. Select the subscription in which to create the storage account.
4. Under the Resource group field, select RG1lod11566895
5. Next, enter a name for your storage account named: az400lod11566895stor
6. Select Create.

Step 2: Enable boot diagnostics on existing virtual machine

To enable Boot diagnostics on an existing virtual machine, follow these steps:

1. Sign in to the Azure portal, and then select the virtual machine VM1.
2. In the Support + troubleshooting section, select Boot diagnostics, then select the Settings tab.
3. In Boot diagnostics settings, change the status to On, and from the Storage account drop-down list, select the storage account az400lod11566895stor.
4. Save the change.



You must restart the virtual machine for the change to take effect.

Reference:

<https://docs.microsoft.com/en-us/azure/storage/common/storage-account-create>

<https://docs.microsoft.com/en-us/azure/virtual-machines/troubleshooting/boot-diagnostics>

QUESTION 187

SIMULATION

You have a web app that connects to an Azure SQL Database named db1.

You need to configure db1 to send Query Store runtime statistics to Azure Log Analytics.

To complete this task, sign in to the Microsoft Azure portal.

A. Answer:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

To enable streaming of diagnostic telemetry for a single or a pooled database, follow these steps:

1. Go to Azure SQL database resource.
2. Select Diagnostics settings.
3. Select Turn on diagnostics if no previous settings exist, or select Edit setting to edit a previous setting. You can create up to three parallel connections to stream diagnostic telemetry.
4. Select Add diagnostic setting to configure parallel streaming of diagnostics data to multiple resources.

The screenshot shows the 'CRM Database - Diagnostic settings' page in the Azure portal. The left sidebar contains navigation links: Connection strings, Sync to other databases, Add Azure Search, Properties, Locks, Automation script, Security (Advanced Threat Protection, Auditing, Dynamic Data Masking, Transparent data encryption), Intelligent Performance (Performance overview, Performance recommendations, Query Performance Insight, Automatic tuning), Monitoring (Alerts (Classic), Metrics, Diagnostic settings), Support + troubleshooting, and Resource health. The 'Diagnostic settings' link is highlighted with a red box. The main content area includes fields for Subscription (Workload Insight dev/test subscription), Resource group (widemo), and Resource type (0 selected). A button labeled 'Turn on diagnostics' is visible, also highlighted with a red box. The breadcrumb navigation shows 'Workload Insight dev/test subscription > widemo > CRM Database'.

Answer:

Explanation/Reference:

Answer: See solution below.

Explanation

Explanation/Reference:

Explanation:

To enable streaming of diagnostic telemetry for a single or a pooled database, follow these steps:

1. Go to Azure SQL database resource.
2. Select Diagnostics settings.
3. Select Turn on diagnostics if no previous settings exist, or select Edit setting to edit a previous setting. You can create up to three parallel connections to stream diagnostic telemetry.
4. Select Add diagnostic setting to configure parallel streaming of diagnostics data to multiple resources.

CRM Database - Diagnostic settings

SQL database

Search (Ctrl+ /)

Refresh

* Subscription: Workload Insight dev/test subscription

Resource group: widemo

Resource type: 0 selected

Workload Insight dev/test subscription > widemo > CRM Database

Turn on diagnostics to collect the following data.

Connection strings

Sync to other databases

Add Azure Search

Properties

Locks

Automation script

Security

- Advanced Threat Protection
- Auditing
- Dynamic Data Masking
- Transparent data encryption

Intelligent Performance

- Performance overview
- Performance recommendations
- Query Performance Insight
- Automatic tuning

Monitoring

- Alerts (Classic)
- Metrics
- Diagnostic settings

Support + troubleshooting

Resource health

The screenshot shows the 'Diagnostic settings' page for a 'CRM Database' in the Azure portal. The left sidebar has a navigation menu with several sections: 'Connection strings', 'Sync to other databases', 'Add Azure Search', 'Properties', 'Locks', 'Automation script', 'Security' (with sub-options like 'Advanced Threat Protection'), 'Intelligent Performance' (with sub-options like 'Performance overview'), 'Monitoring' (with sub-options like 'Alerts (Classic)', 'Metrics', and 'Diagnostic settings'), 'Support + troubleshooting', and 'Resource health'. The 'Diagnostic settings' link is highlighted with a red box. The main content area shows 'Subscription' set to 'Workload Insight dev/test subscription', 'Resource group' set to 'widemo', and 'Resource type' set to '0 selected'. A note says 'Turn on diagnostics to collect the following data.' Below this, there's a breadcrumb trail: 'Workload Insight dev/test subscription > widemo > CRM Database'.

5. Enter a setting name for your own reference.
6. Select a destination resource for the streaming diagnostics data: Archive to storage account, Stream to an event hub, or Send to Log Analytics.
7. For the standard, event-based monitoring experience, select the following check boxes for database diagnostics log telemetry: QueryStoreRuntimeStatistics

Diagnostics settings

Save Discard Delete

* Name
 ✓

Archive to a storage account

Stream to an event hub

Send to Log Analytics

Subscription

▼

Log Analytics Workspace

▼

LOG

- SQLInsights
- AutomaticTuning
- QueryStoreRuntimeStatistics
- QueryStoreWaitStatistics
- Errors
- DatabaseWaitStatistics
- Timeouts
- Blocks
- Deadlocks

METRIC

8. For an advanced, one-minute-based monitoring experience, select the check box for Basic metrics.

9. Select Save.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-sql/database/metrics-diagnostic-telemetry-logging-streaming-export-configure>